

MP4

April 12, 2024

1 (Optional) Colab Setup

If you aren't using Colab, you can delete the following code cell. This is just to help students with mounting to Google Drive to access the other .py files and downloading the data, which is a little trickier on Colab than on your local machine using Jupyter.

```
[1]: # you will be prompted with a window asking to grant permissions
from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

```
[ ]: # fill in the path in your Google Drive in the string below. Note: do not ↵
      ↵escape slashes or spaces
import os
datadir = "/content/assignment4"
if not os.path.exists(datadir):
    !ln -s "/content/drive/My Drive/CS444/assignment4" $datadir # TODO: Fill your ↵
      ↵Assignment 4 path
os.chdir(datadir)
!pwd
```

2 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
[ ]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
```

```

plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[ ]: from gan.train import train
```

```
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
print(device)
```

cuda:0

3 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

3.0.1 GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log (1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
[ ]: from gan.losses import discriminator_loss, generator_loss
```

3.0.2 Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
[ ]: from gan.losses import ls_discriminator_loss, ls_generator_loss
```

4 GAN model architecture

TODO: Implement the `Discriminator` and `Generator` networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in_channels=1024, out_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=512, out_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=256, out_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[ ]: from gan.models import Discriminator, Generator
```

5 Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

TODO: Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
[ ]: batch_size = 64
imsize = 64
cat_root = './cats'

cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imsize)),
    transforms.RandomCrop(imsize),
]))

cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)
```

5.0.1 Visualize dataset

```
[ ]: from gan.utils import show_images

try:
    imgs = next(iter(cat_loader_train))[0].numpy().squeeze()
except:
```

```
imgs = cat_loader_train.__iter__().next()[0].numpy().squeeze()  
  
show_images(imgs, color=True)
```



6 Training

TODO: Fill in the training loop in `gan/train.py`.

```
[ ]: NOISE_DIM = 100  
      NUM_EPOCHS = 50  
      learning_rate = 0.001
```

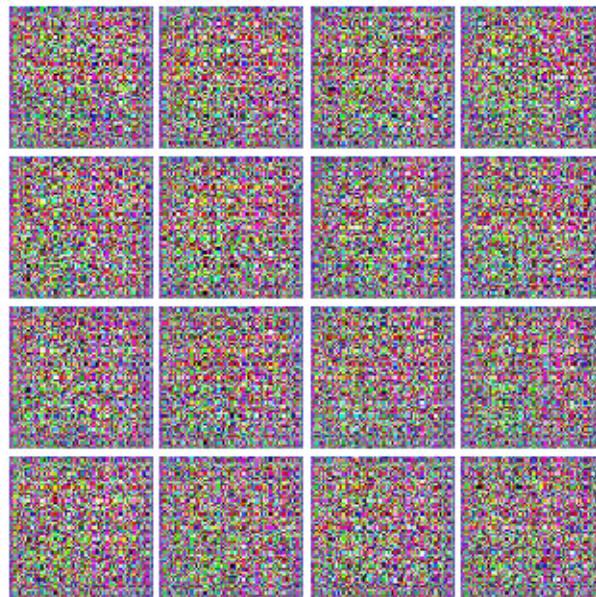
6.0.1 Train GAN

```
[ ]: D = Discriminator().to(device)
      G = Generator(noise_dim=NOISE_DIM).to(device)

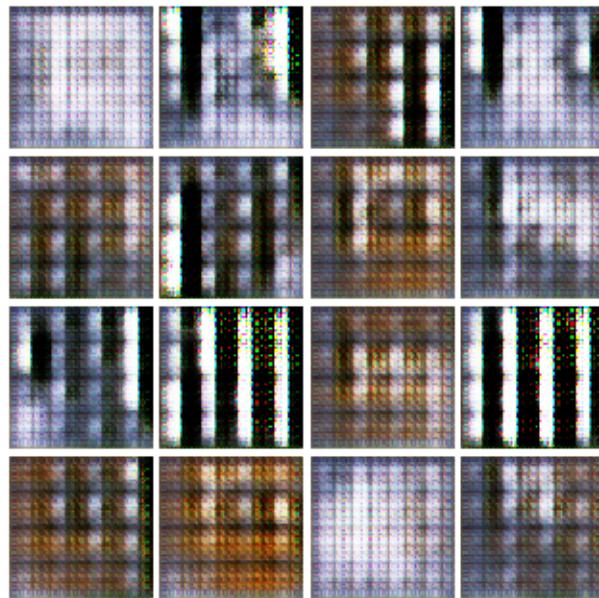
[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[ ]: # original gan
      train(D, G, D_optimizer, G_optimizer, discriminator_loss,
            generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
            batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

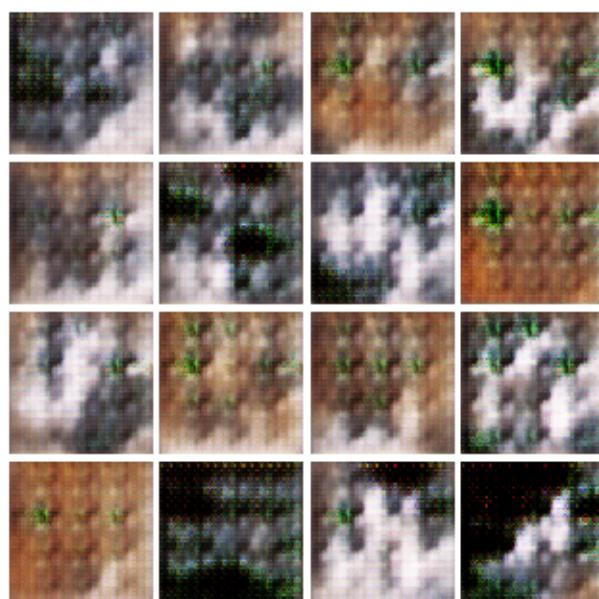
EPOCH: 1
Iter: 0, D: 1.434, G:6.76



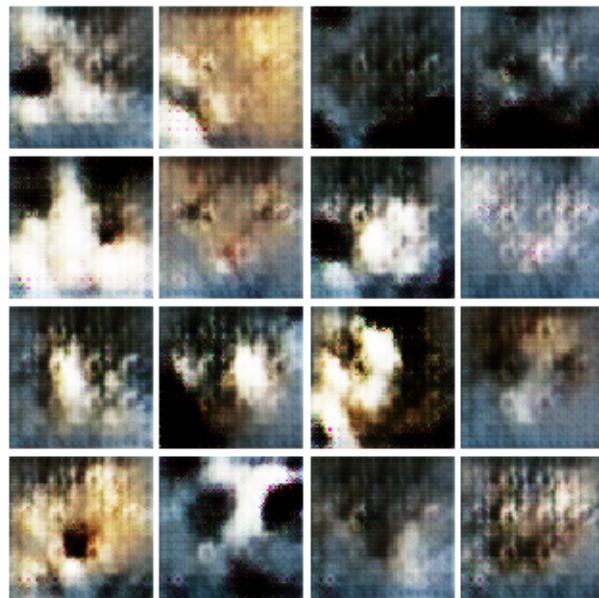
EPOCH: 2
Iter: 250, D: 0.4139, G:2.922



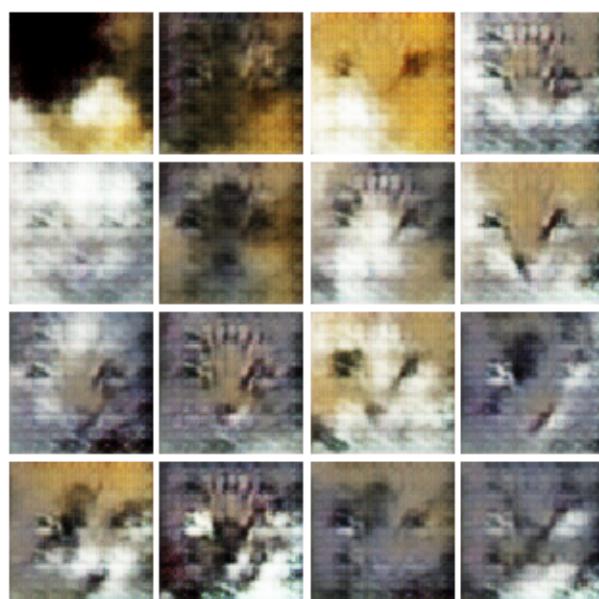
EPOCH: 3
Iter: 500, D: 1.321, G:0.954



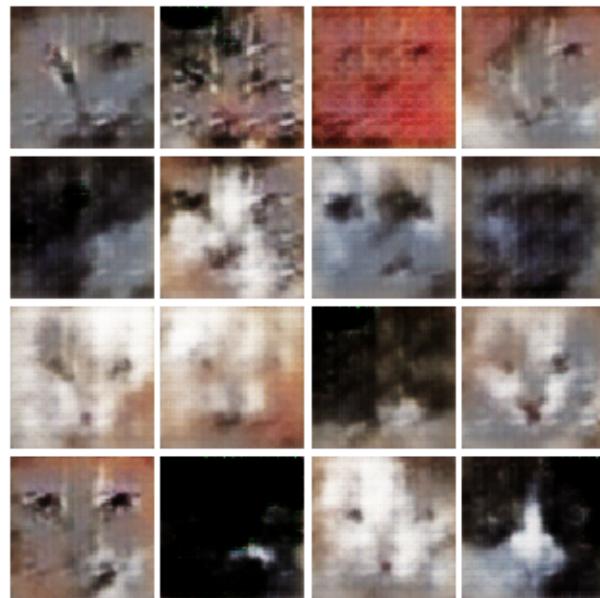
EPOCH: 4
Iter: 750, D: 1.81, G:2.555



EPOCH: 5
Iter: 1000, D: 2.105, G:1.181



EPOCH: 6
Iter: 1250, D: 0.7123, G:1.995



EPOCH: 7
Iter: 1500, D: 1.116, G:2.779



EPOCH: 8
Iter: 1750, D: 1.118, G:1.565



EPOCH: 9
Iter: 2000, D: 0.5186, G:1.965



EPOCH: 10
Iter: 2250, D: 1.07, G:5.21



EPOCH: 11

Iter: 2500, D: 0.4391, G:3.272



EPOCH: 12

Iter: 2750, D: 0.6128, G:3.432



EPOCH: 13

Iter: 3000, D: 0.8975, G:1.864



EPOCH: 14

Iter: 3250, D: 0.4588, G:3.39



EPOCH: 15

Iter: 3500, D: 0.8431, G: 6.011



EPOCH: 16

Iter: 3750, D: 0.6837, G: 2.197



EPOCH: 17
Iter: 4000, D: 0.76, G:3.597



EPOCH: 18
Iter: 4250, D: 2.298, G:0.9215



EPOCH: 19

Iter: 4500, D: 0.4436, G: 4.173



EPOCH: 20

Iter: 4750, D: 0.6505, G: 2.678



EPOCH: 21
Iter: 5000, D: 2.109, G:1.606



EPOCH: 22
Iter: 5250, D: 0.5301, G:5.821



EPOCH: 23

Iter: 5500, D: 0.3504, G:3.502



EPOCH: 24

Iter: 5750, D: 0.3673, G:2.862



EPOCH: 25
Iter: 6000, D: 0.469, G:4.492



EPOCH: 26
Iter: 6250, D: 0.3664, G:3.605



EPOCH: 27
Iter: 6500, D: 1.53, G:1.404



EPOCH: 28
Iter: 6750, D: 0.2472, G:2.748



EPOCH: 29

Iter: 7000, D: 1.059, G:6.211



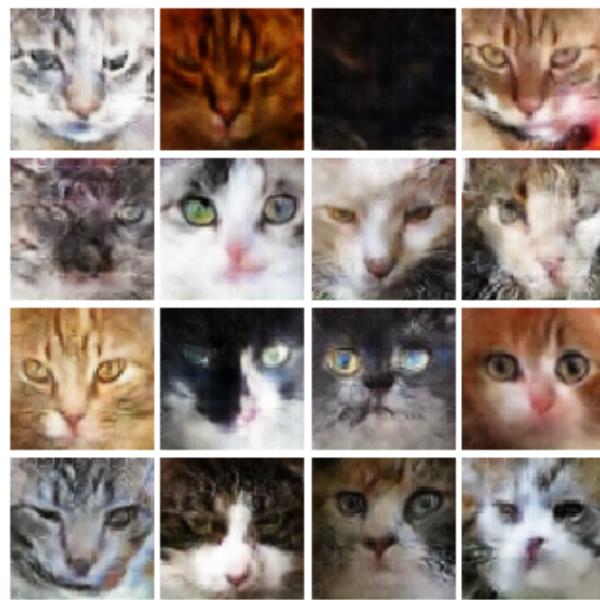
EPOCH: 30

Iter: 7250, D: 1.354, G:8.999



EPOCH: 31

Iter: 7500, D: 0.3022, G:4.046



EPOCH: 32

Iter: 7750, D: 0.7092, G:5.825



EPOCH: 33
Iter: 8000, D: 1.464, G:1.967



EPOCH: 34
Iter: 8250, D: 0.1711, G:4.863



EPOCH: 35
Iter: 8500, D: 0.2343, G:4.1



EPOCH: 36
Iter: 8750, D: 0.322, G:4.124



EPOCH: 37

Iter: 9000, D: 0.4345, G: 6.169



EPOCH: 38

Iter: 9250, D: 0.6034, G: 5.286



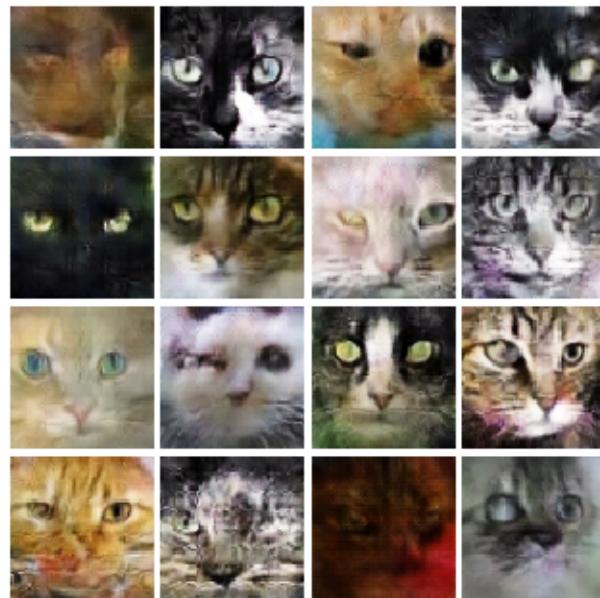
EPOCH: 39

Iter: 9500, D: 0.9227, G:8.206



EPOCH: 40

Iter: 9750, D: 0.5856, G:3.554



EPOCH: 41

Iter: 10000, D: 0.3183, G: 4.089



EPOCH: 42

Iter: 10250, D: 0.3083, G: 3.383



EPOCH: 43

Iter: 10500, D: 0.1295, G: 4.307



EPOCH: 44

Iter: 10750, D: 0.1736, G: 4.867



EPOCH: 45

Iter: 11000, D: 0.5499, G: 6.571



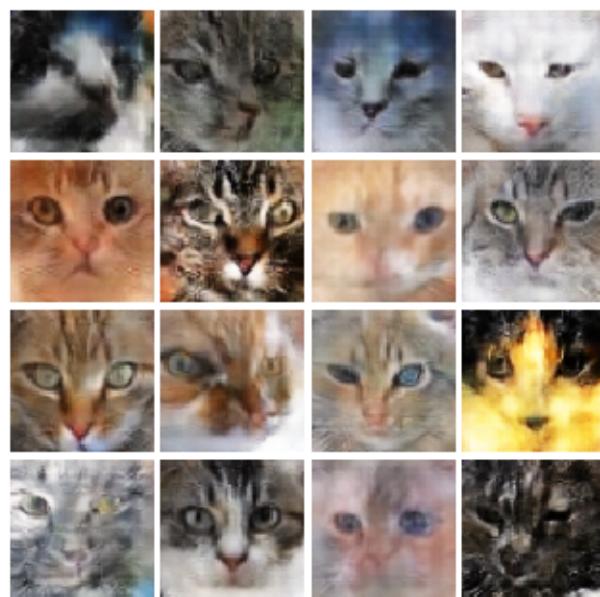
EPOCH: 46

Iter: 11250, D: 0.2223, G: 6.768



EPOCH: 47

Iter: 11500, D: 0.2271, G:5.016



EPOCH: 48

Iter: 11750, D: 0.5456, G:3.43



EPOCH: 49

Iter: 12000, D: 0.2531, G:3.519



EPOCH: 50

Iter: 12250, D: 5.078, G:2.906



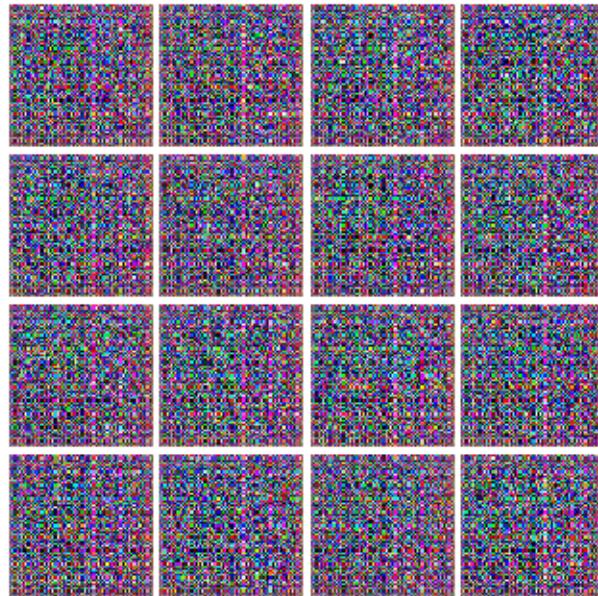
6.0.2 Train LS-GAN

```
[ ]: D = Discriminator().to(device)
      G = Generator(noise_dim=NOISE_DIM).to(device)

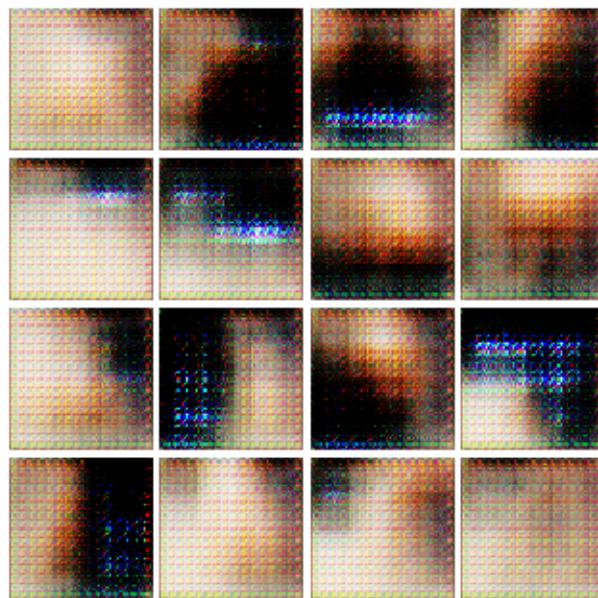
[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[ ]: # ls-gan
      train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
            ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
            batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

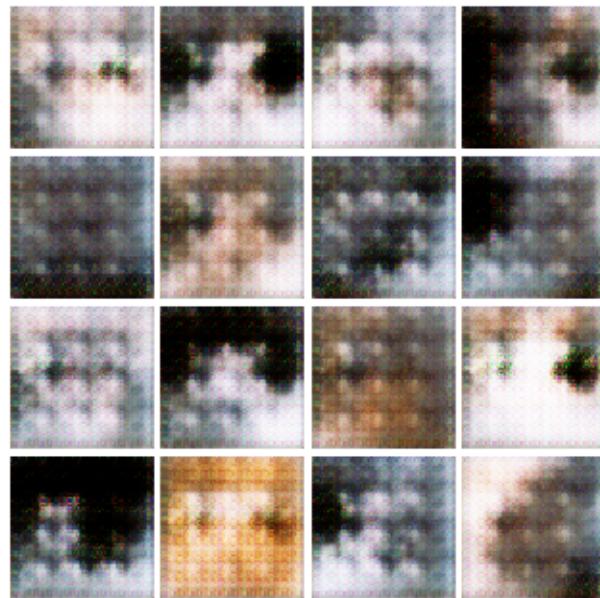
EPOCH: 1
Iter: 0, D: 0.4937, G:32.85



EPOCH: 2
Iter: 250, D: 0.2076, G:0.2198

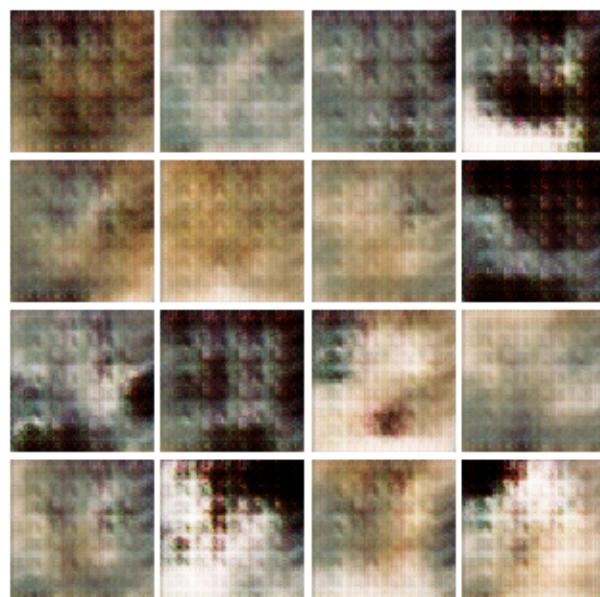


EPOCH: 3
Iter: 500, D: 0.7306, G:0.06443



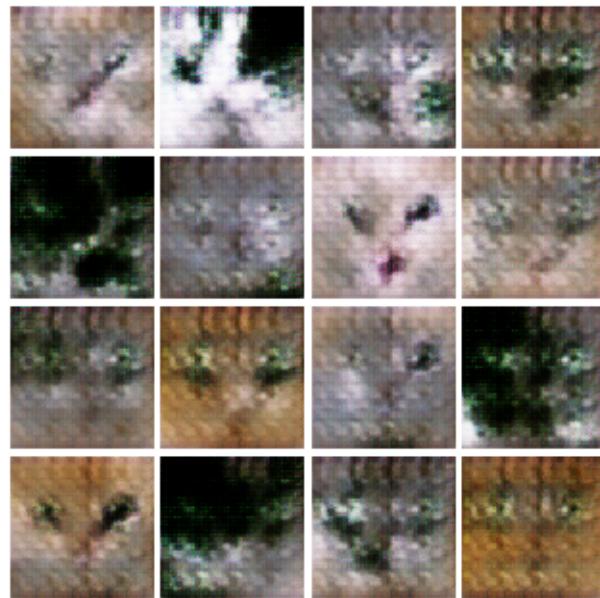
EPOCH: 4

Iter: 750, D: 0.1666, G:0.2884

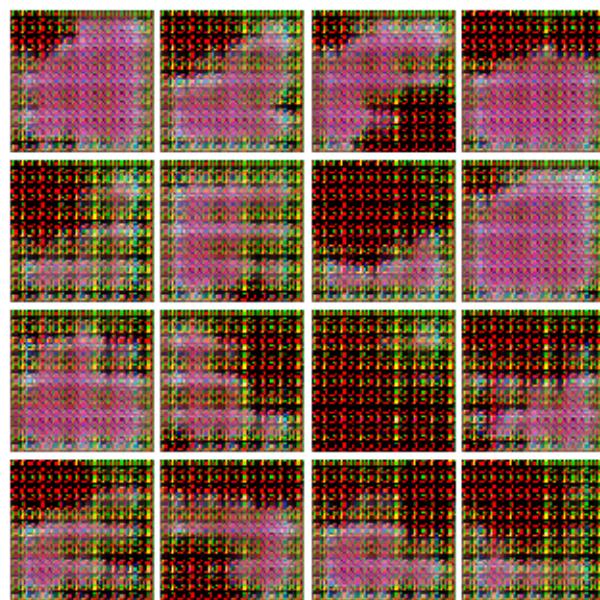


EPOCH: 5

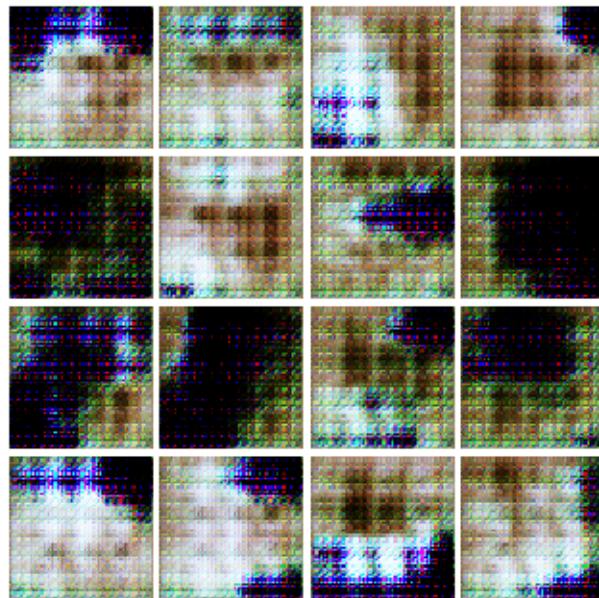
Iter: 1000, D: 0.2997, G:0.6397



EPOCH: 6
Iter: 1250, D: 0.1226, G: 0.2547

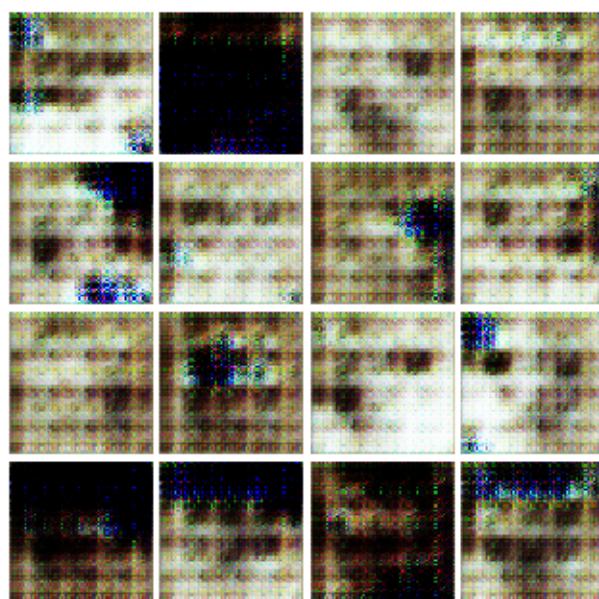


EPOCH: 7
Iter: 1500, D: 0.2253, G: 0.1375



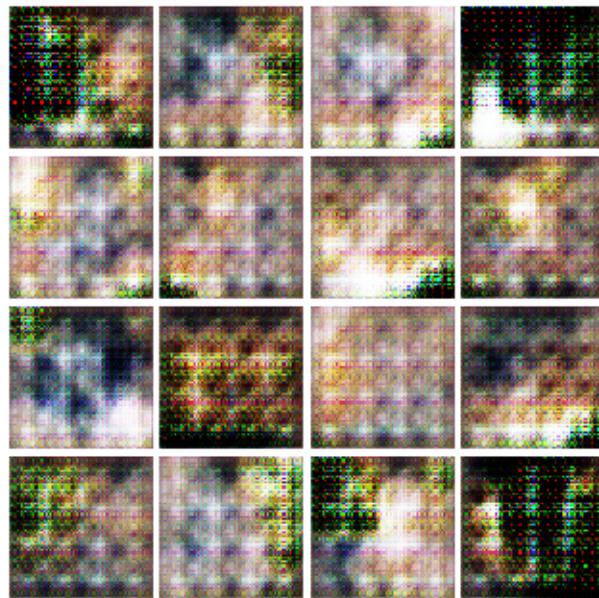
EPOCH: 8

Iter: 1750, D: 0.2253, G: 0.1999



EPOCH: 9

Iter: 2000, D: 0.2585, G: 0.1423



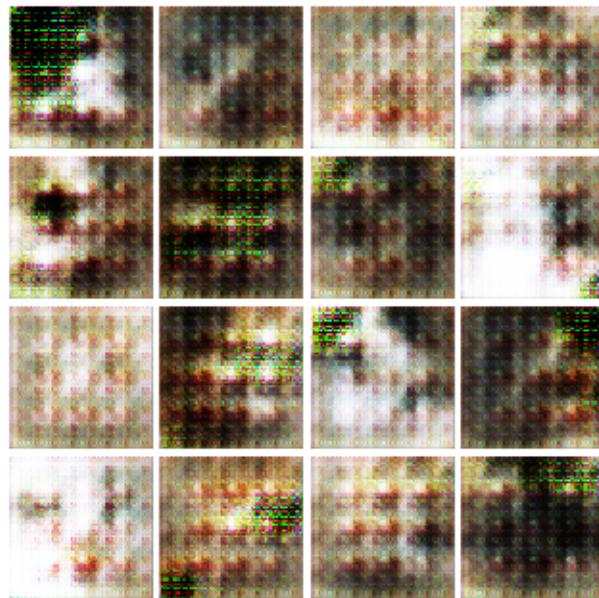
EPOCH: 10

Iter: 2250, D: 0.2479, G: 0.1602



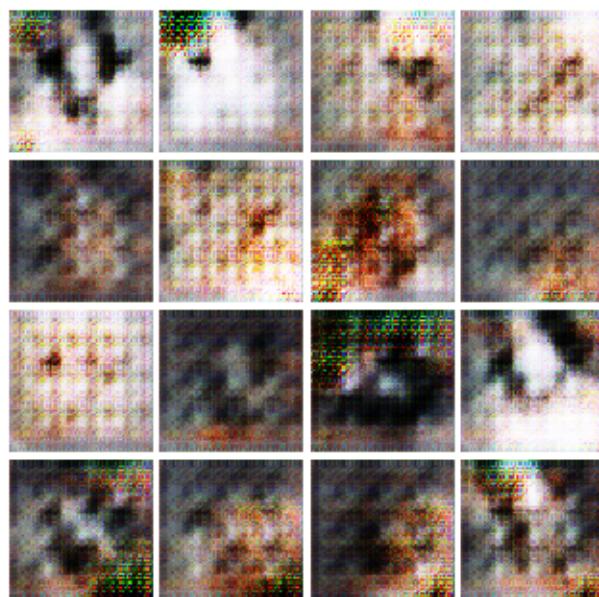
EPOCH: 11

Iter: 2500, D: 0.2509, G: 0.1554



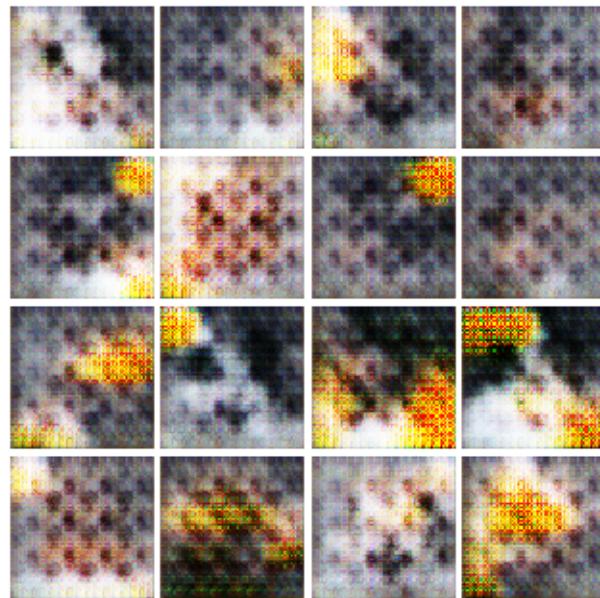
EPOCH: 12

Iter: 2750, D: 0.2521, G: 0.1607



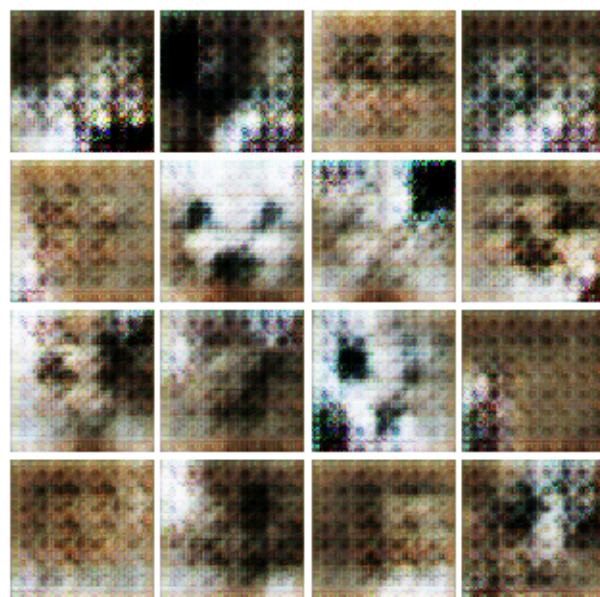
EPOCH: 13

Iter: 3000, D: 0.2594, G: 0.1666



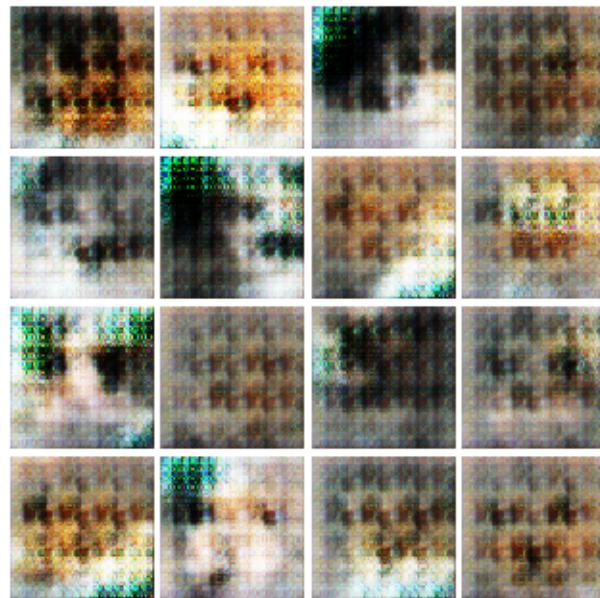
EPOCH: 14

Iter: 3250, D: 0.2588, G: 0.1337



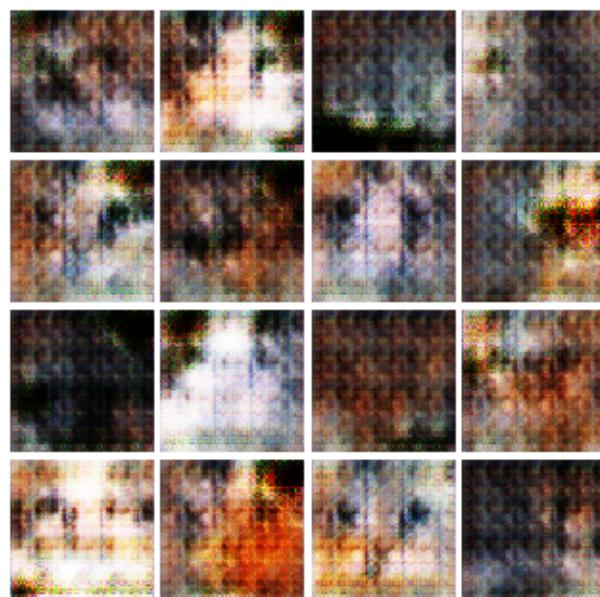
EPOCH: 15

Iter: 3500, D: 0.2274, G: 0.1446



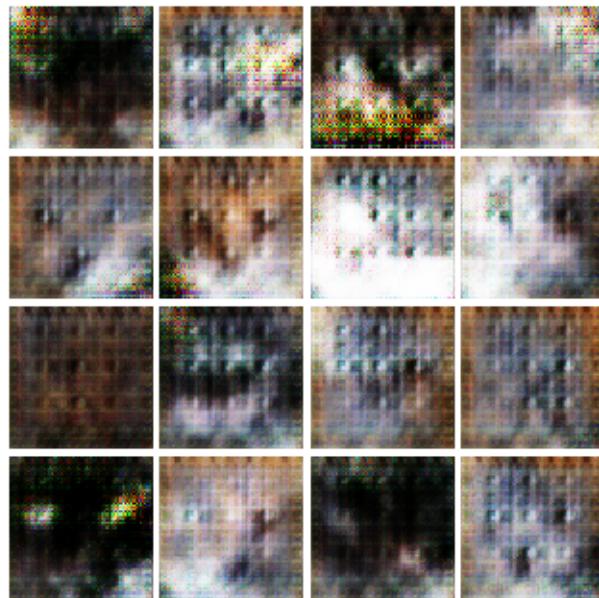
EPOCH: 16

Iter: 3750, D: 0.2629, G: 0.1613



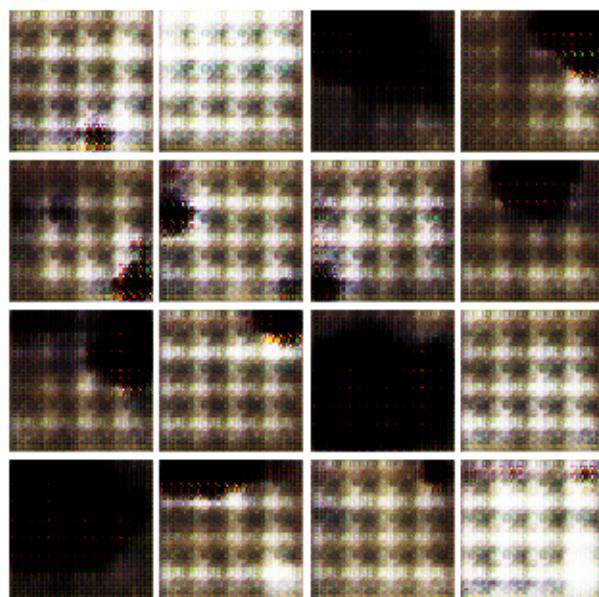
EPOCH: 17

Iter: 4000, D: 0.2655, G: 0.1462



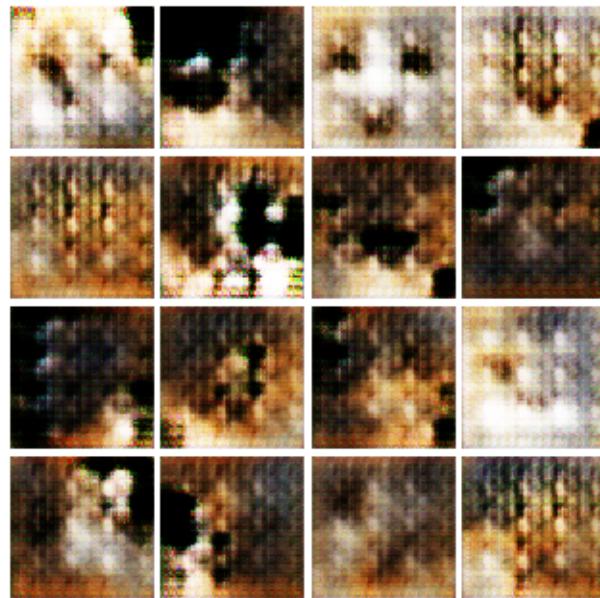
EPOCH: 18

Iter: 4250, D: 0.3798, G: 0.4034



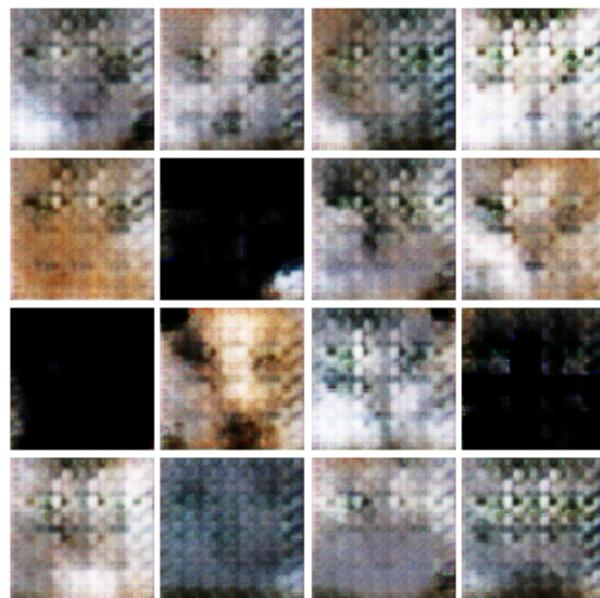
EPOCH: 19

Iter: 4500, D: 0.2048, G: 0.1754



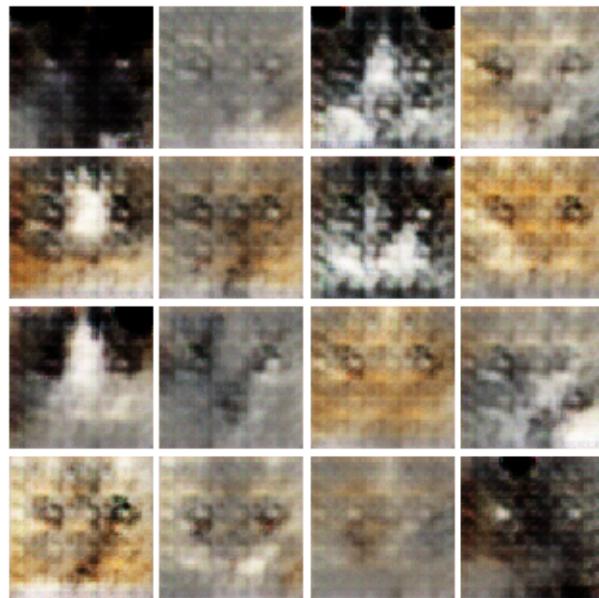
EPOCH: 20

Iter: 4750, D: 0.1935, G: 0.3828



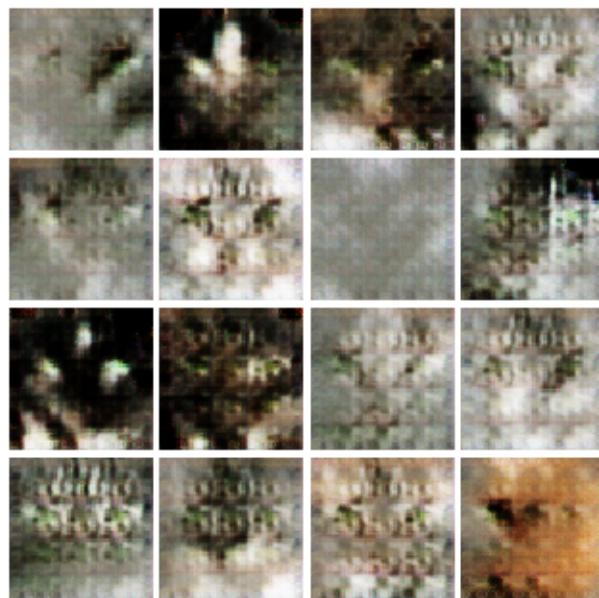
EPOCH: 21

Iter: 5000, D: 0.1796, G: 0.1603



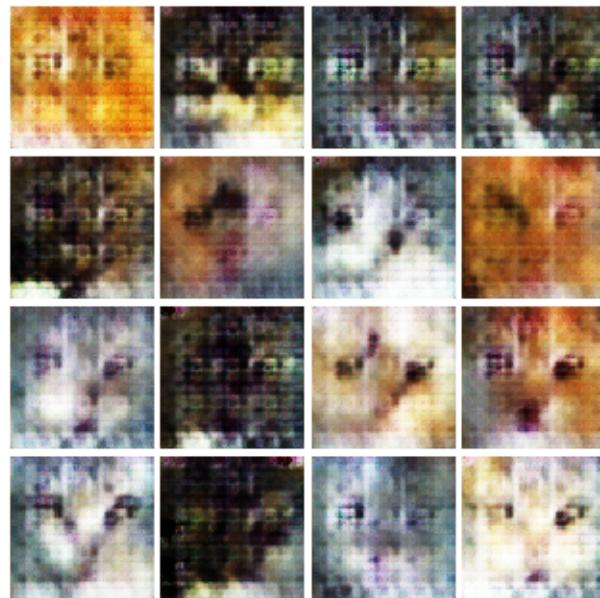
EPOCH: 22

Iter: 5250, D: 0.2654, G: 0.5572



EPOCH: 23

Iter: 5500, D: 0.27, G: 0.1135



EPOCH: 24

Iter: 5750, D: 0.2412, G: 0.2736



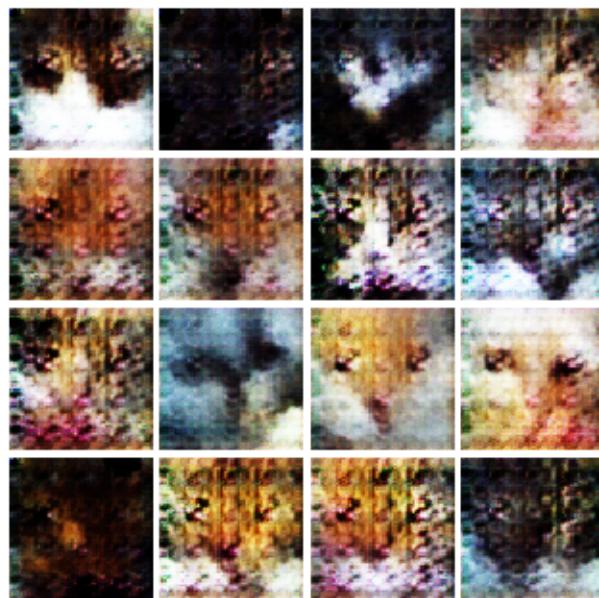
EPOCH: 25

Iter: 6000, D: 0.2285, G: 0.3909



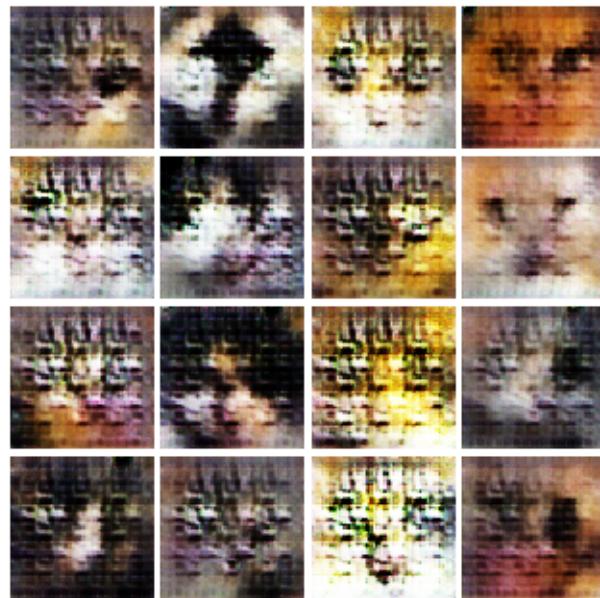
EPOCH: 26

Iter: 6250, D: 0.2662, G: 0.2258



EPOCH: 27

Iter: 6500, D: 0.4462, G: 0.3337



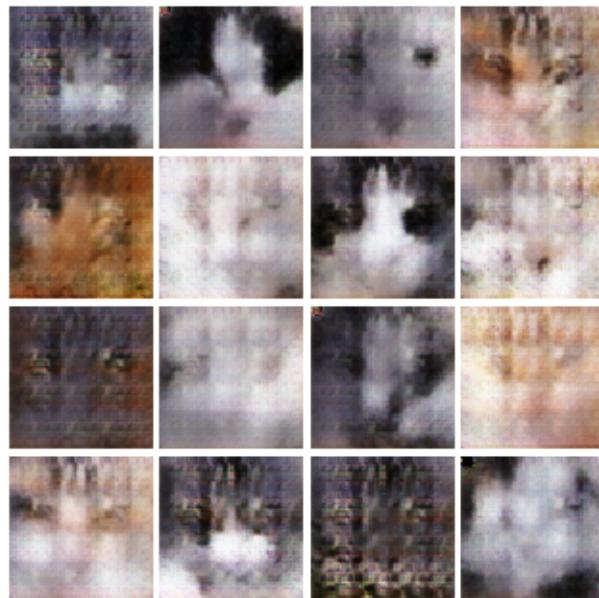
EPOCH: 28

Iter: 6750, D: 0.2506, G: 0.1854



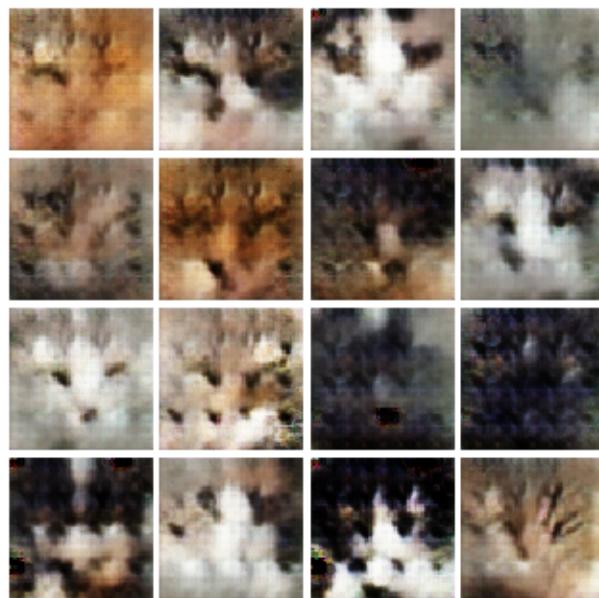
EPOCH: 29

Iter: 7000, D: 0.2969, G: 0.2533



EPOCH: 30

Iter: 7250, D: 0.2189, G: 0.1839



EPOCH: 31

Iter: 7500, D: 0.4495, G: 0.2258



EPOCH: 32

Iter: 7750, D: 0.239, G: 0.1689



EPOCH: 33

Iter: 8000, D: 0.1462, G: 0.241



EPOCH: 34

Iter: 8250, D: 0.179, G: 0.1969



EPOCH: 35

Iter: 8500, D: 0.1802, G: 0.4034



EPOCH: 36

Iter: 8750, D: 0.251, G:0.1944



EPOCH: 37

Iter: 9000, D: 0.221, G:0.6671



EPOCH: 38

Iter: 9250, D: 0.2752, G: 0.1956



EPOCH: 39

Iter: 9500, D: 0.05867, G: 0.2958



EPOCH: 40

Iter: 9750, D: 0.1881, G: 0.2339



EPOCH: 41

Iter: 10000, D: 0.1298, G: 0.2789



EPOCH: 42

Iter: 10250, D: 0.1615, G:0.271



EPOCH: 43

Iter: 10500, D: 0.1891, G:0.4147



EPOCH: 44

Iter: 10750, D: 0.2499, G: 0.1897



EPOCH: 45

Iter: 11000, D: 0.2359, G: 0.4049



EPOCH: 46

Iter: 11250, D: 0.1909, G: 0.2361



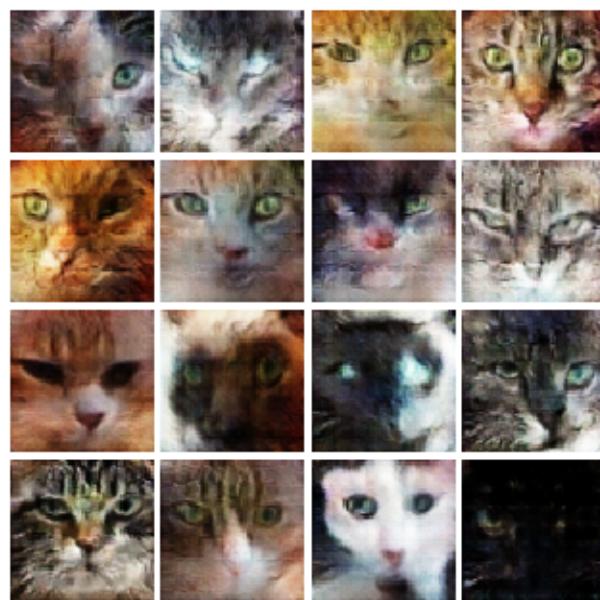
EPOCH: 47

Iter: 11500, D: 0.2203, G: 0.162



EPOCH: 48

Iter: 11750, D: 0.2374, G: 0.2051



EPOCH: 49

Iter: 12000, D: 0.1891, G: 0.2963



EPOCH: 50
Iter: 12250, D: 0.2157, G: 0.2054



7 Extra Credit 1: WGAN Loss

7.0.1 Train WGAN

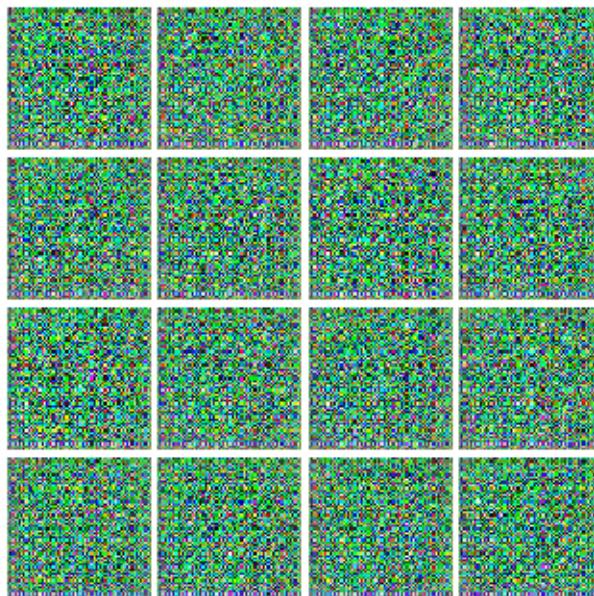
```
[ ]: D = Discriminator().to(device)
      G = Generator(noise_dim=NOISE_DIM).to(device)

[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

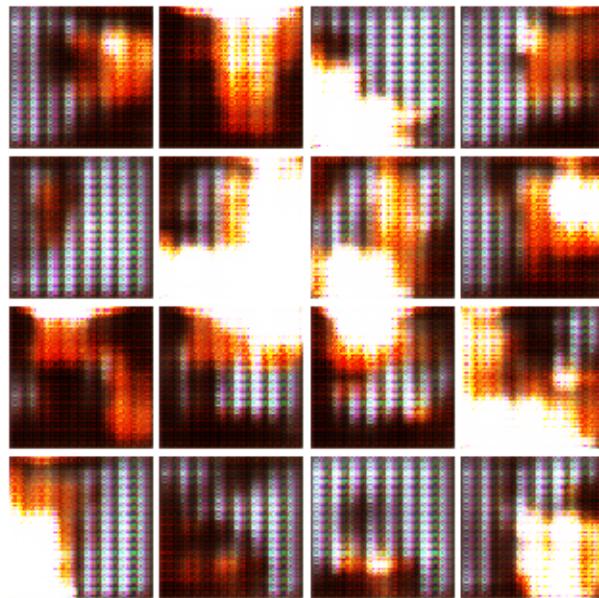
[ ]: from gan.losses import w_discriminator_loss, w_generator_loss

[ ]: # w-gan
      train(D, G, D_optimizer, G_optimizer, w_discriminator_loss,
            w_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
            batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 0.0248, G:6.091

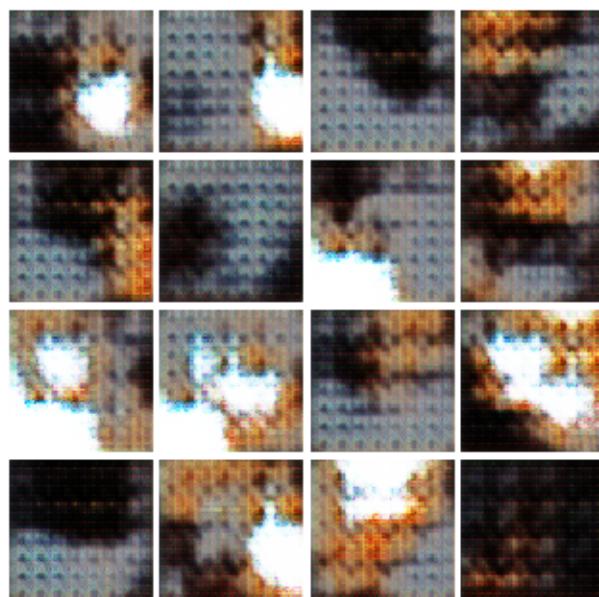


EPOCH: 2
Iter: 250, D: -1.148e+03, G:728.5



EPOCH: 3

Iter: 500, D: -4.819e+03, G:3.041e+03



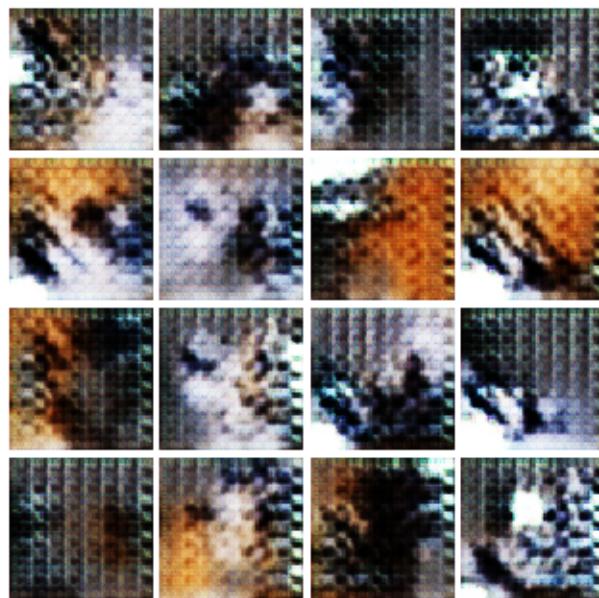
EPOCH: 4

Iter: 750, D: -5.589e+03, G:5.14e+03



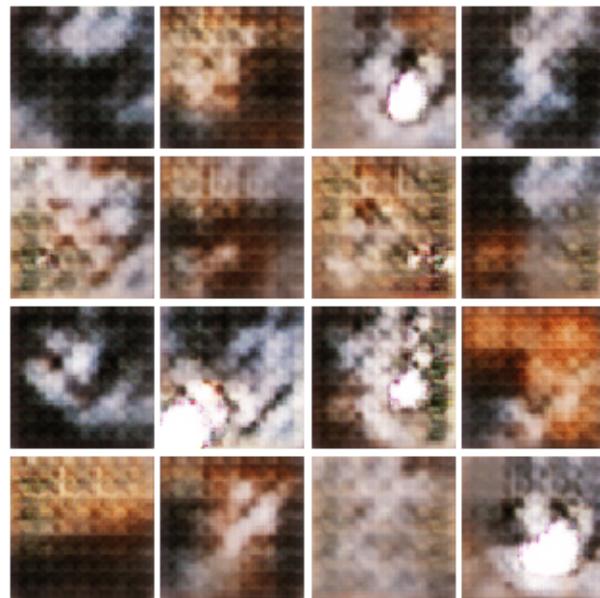
EPOCH: 5

Iter: 1000, D: -927.4, G:-7.76e+03



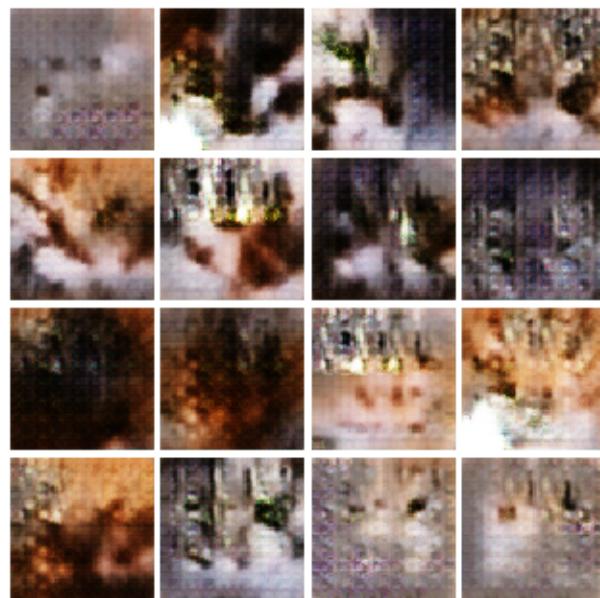
EPOCH: 6

Iter: 1250, D: -3.558e+03, G:2.24e+03



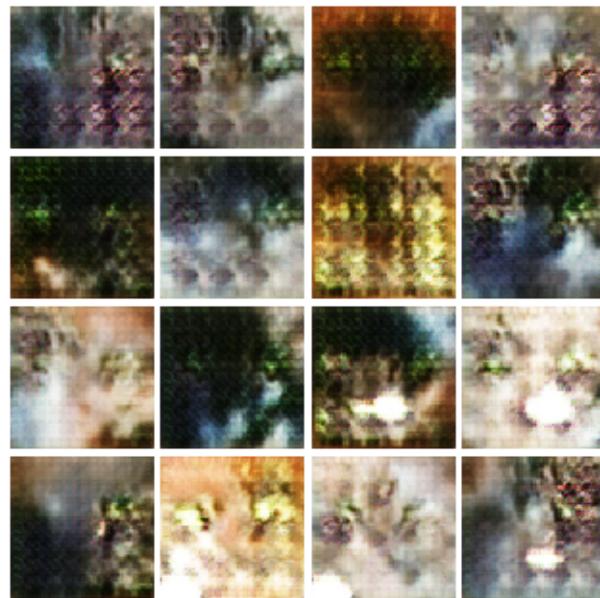
EPOCH: 7

Iter: 1500, D: -2.549e+04, G:1.537e+04



EPOCH: 8

Iter: 1750, D: -7.862e+03, G:-1.492e+03



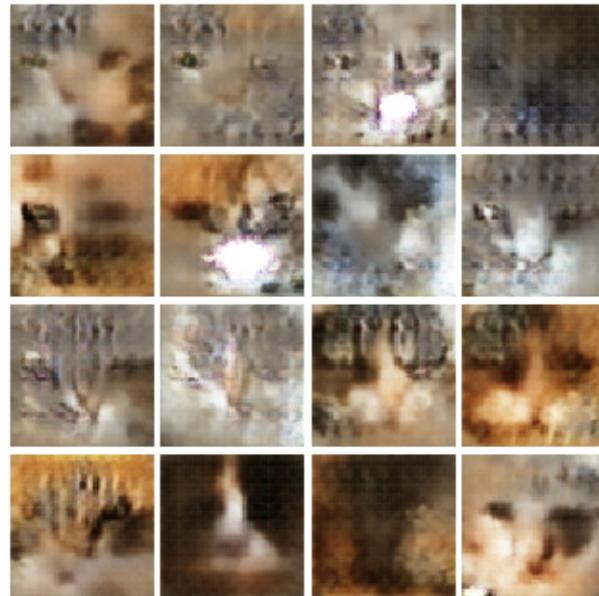
EPOCH: 9

Iter: 2000, D: -8.144e+03, G:-414.2



EPOCH: 10

Iter: 2250, D: -3.693e+04, G:2.31e+04



EPOCH: 11

Iter: 2500, D: -4.357e+04, G:4.149e+04



EPOCH: 12

Iter: 2750, D: -6.45e+04, G:5.011e+04



EPOCH: 13

Iter: 3000, D: -1.128e+04, G:3.343e+03



EPOCH: 14

Iter: 3250, D: -7.664e+04, G:5.96e+04



EPOCH: 15

Iter: 3500, D: -1.028e+05, G:8.698e+04



EPOCH: 16

Iter: 3750, D: -8.201e+04, G:8.154e+04



EPOCH: 17

Iter: 4000, D: -7.555e+04, G:5.451e+04



EPOCH: 18

Iter: 4250, D: -1.335e+05, G:4.106e+04



EPOCH: 19

Iter: 4500, D: -9.836e+04, G:1.176e+05



EPOCH: 20

Iter: 4750, D: -2.165e+05, G:-7.68e+04



EPOCH: 21

Iter: 5000, D: -4.391e+05, G:2.159e+05



EPOCH: 22

Iter: 5250, D: -9.723e+04, G:1.565e+05



EPOCH: 23

Iter: 5500, D: -1.618e+05, G:1.882e+05



EPOCH: 24

Iter: 5750, D: -2.526e+04, G:-1.874e+05



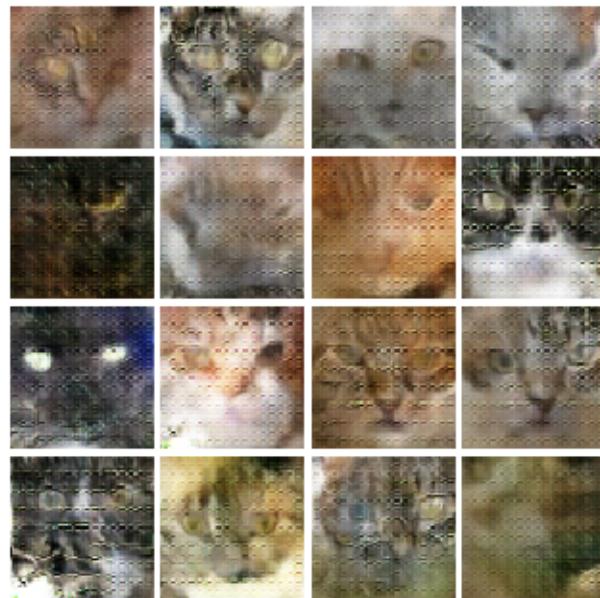
EPOCH: 25

Iter: 6000, D: -1.642e+05, G:2.041e+05



EPOCH: 26

Iter: 6250, D: -6.685e+05, G:3.294e+05



EPOCH: 27

Iter: 6500, D: -7.455e+05, G:3.499e+05



EPOCH: 28

Iter: 6750, D: -3.096e+05, G:2.773e+05



EPOCH: 29

Iter: 7000, D: -1.332e+05, G:3.551e+05



EPOCH: 30

Iter: 7250, D: -3.532e+05, G:1.486e+05



EPOCH: 31

Iter: 7500, D: -1.689e+05, G:2.7e+05



EPOCH: 32

Iter: 7750, D: -1.01e+06, G:4.815e+05



EPOCH: 33

Iter: 8000, D: -1.105e+06, G:5.479e+05



EPOCH: 34

Iter: 8250, D: -4.883e+05, G:3.625e+05



EPOCH: 35

Iter: 8500, D: -2.942e+05, G:4.262e+05



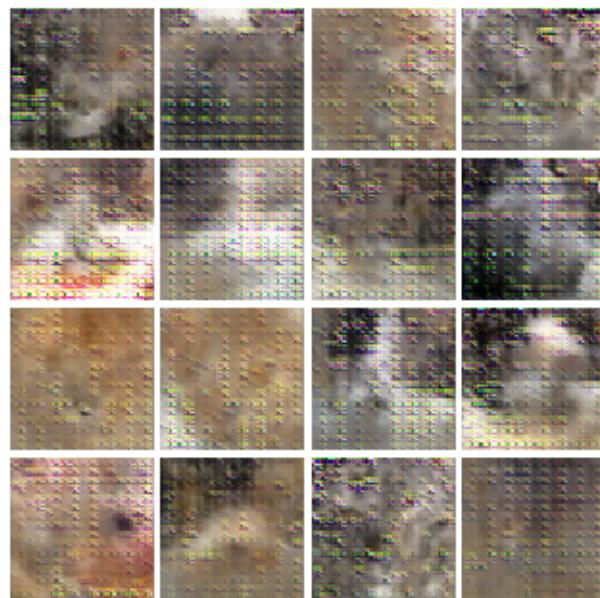
EPOCH: 36

Iter: 8750, D: -1.251e+06, G:6.217e+05



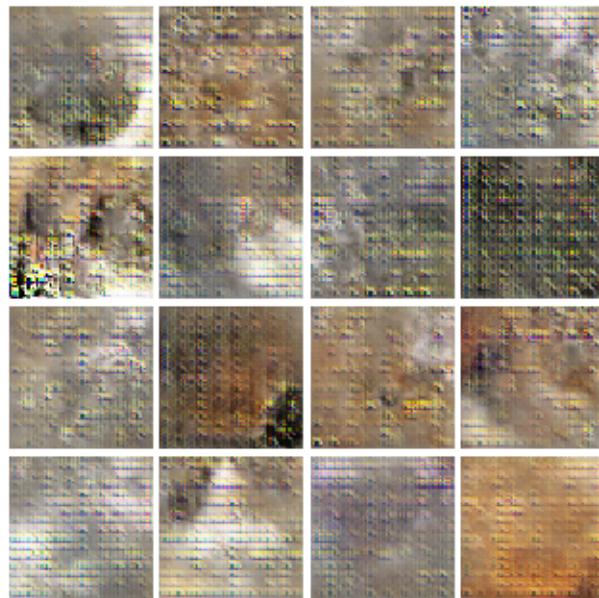
EPOCH: 37

Iter: 9000, D: -1.354e+06, G:6.386e+05



EPOCH: 38

Iter: 9250, D: -1.457e+06, G:6.874e+05



EPOCH: 39

Iter: 9500, D: -1.533e+06, G:7.188e+05



EPOCH: 40

Iter: 9750, D: -1.632e+06, G:7.716e+05



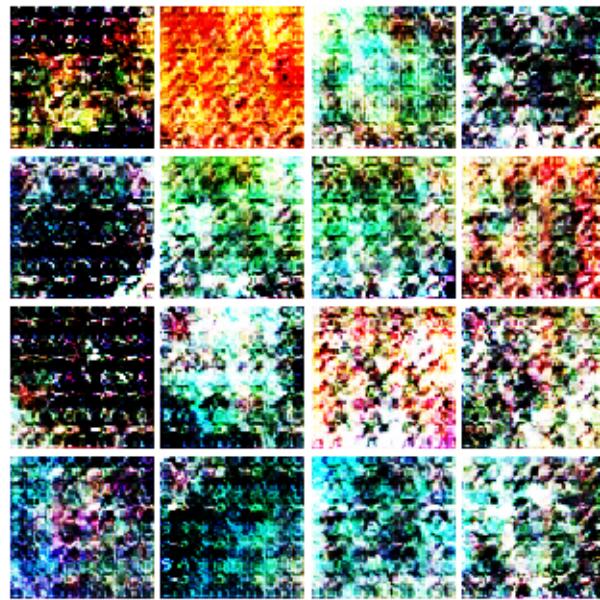
EPOCH: 41

Iter: 10000, D: -1.717e+06, G:8.128e+05



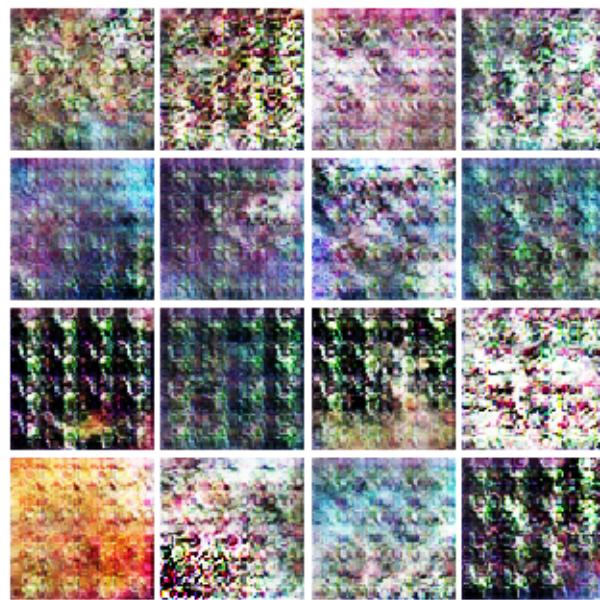
EPOCH: 42

Iter: 10250, D: -1.802e+06, G:8.5e+05



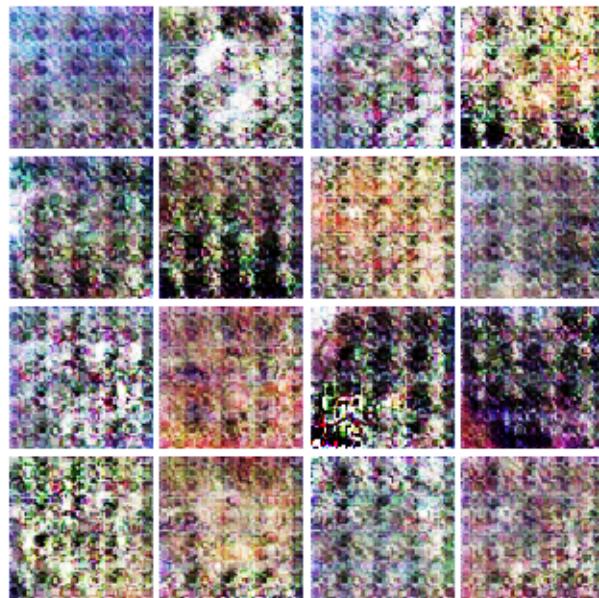
EPOCH: 43

Iter: 10500, D: -1.901e+06, G:8.912e+05



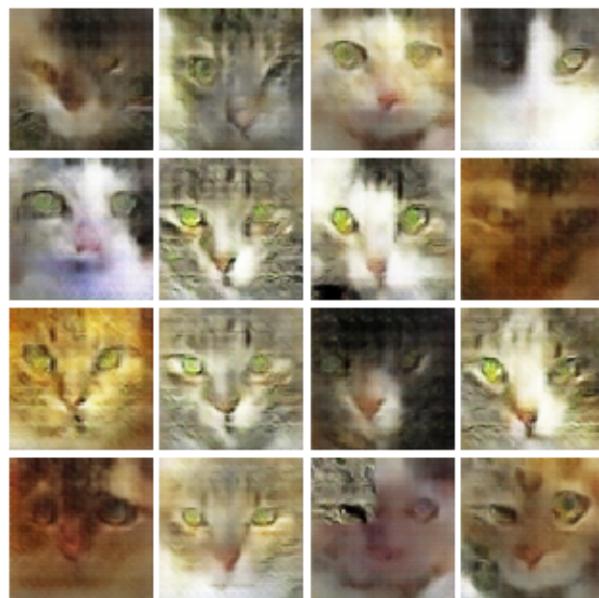
EPOCH: 44

Iter: 10750, D: -1.992e+06, G:9.374e+05



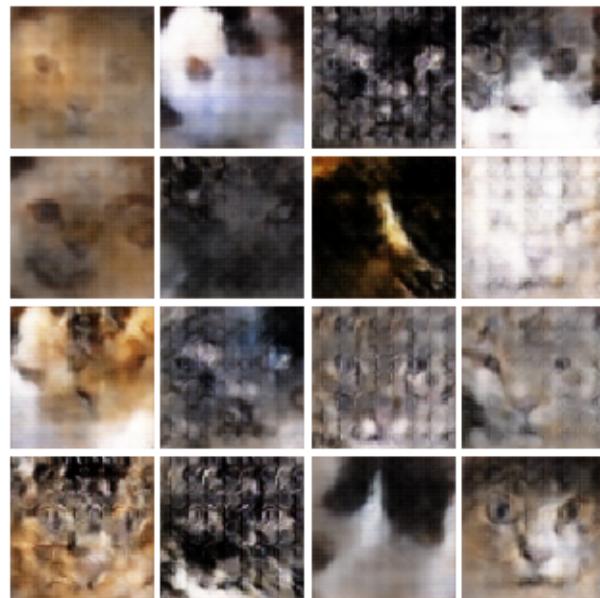
EPOCH: 45

Iter: 11000, D: -7.083e+04, G:-5.667e+05



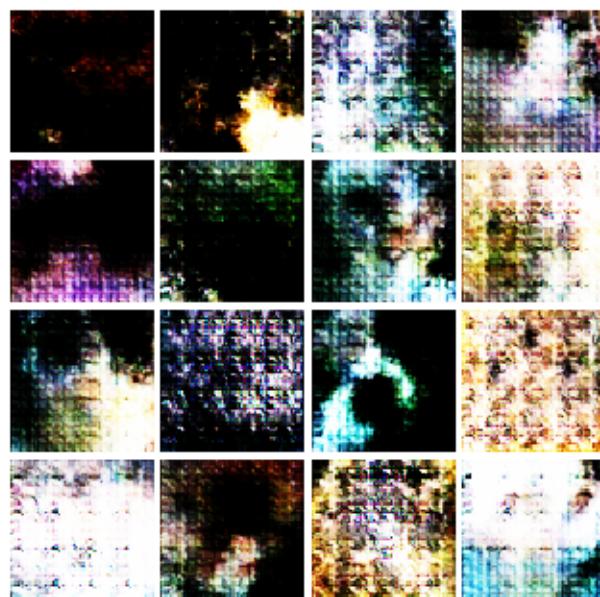
EPOCH: 46

Iter: 11250, D: -2.11e+06, G:1.011e+06



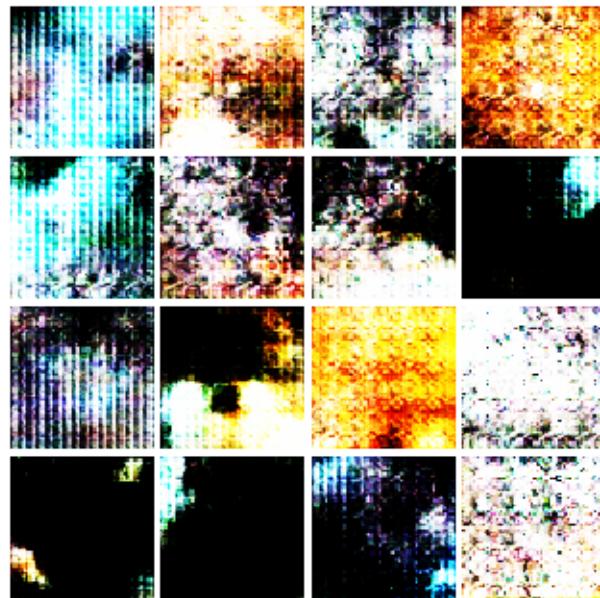
EPOCH: 47

Iter: 11500, D: -2.211e+06, G:1.041e+06



EPOCH: 48

Iter: 11750, D: -2.303e+06, G:1.081e+06



EPOCH: 49

Iter: 12000, D: -2.156e+05, G:1.669e+04



EPOCH: 50

Iter: 12250, D: 9.797e+04, G:-6.357e+05



8 Extra Credit 2: Spectral Normalization

```
[ ]: D = Discriminator().to(device)
      G = Generator(noise_dim=NOISE_DIM).to(device)

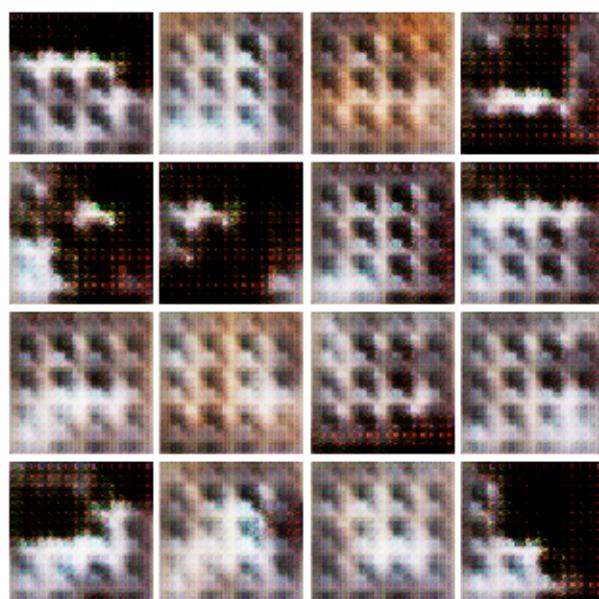
[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[ ]: train(D, G, D_optimizer, G_optimizer, discriminator_loss,
           generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
           batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

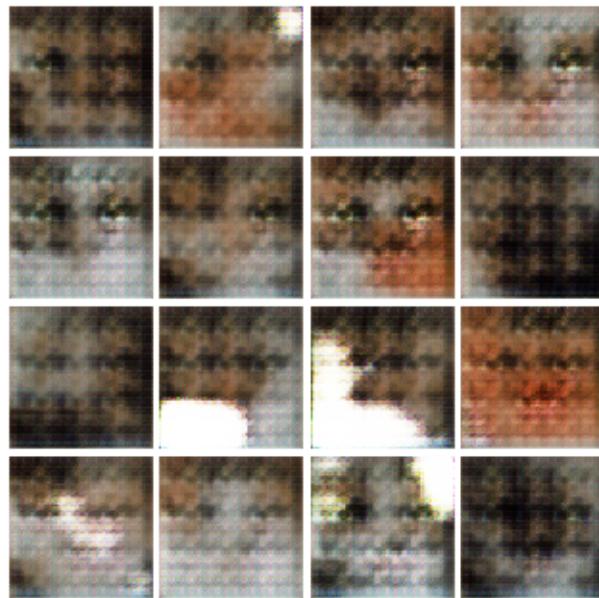
EPOCH: 1
Iter: 0, D: 1.404, G:4.587



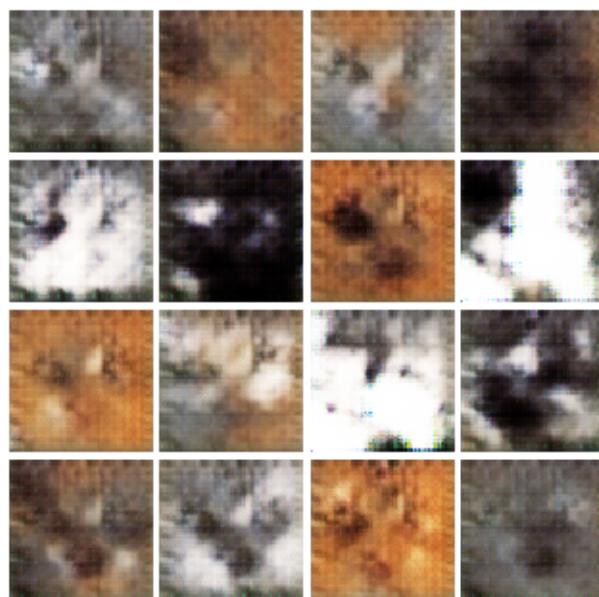
EPOCH: 2
Iter: 250, D: 1.468, G:3.173



EPOCH: 3
Iter: 500, D: 1.285, G:2.781



EPOCH: 4
Iter: 750, D: 1.686, G:5.162



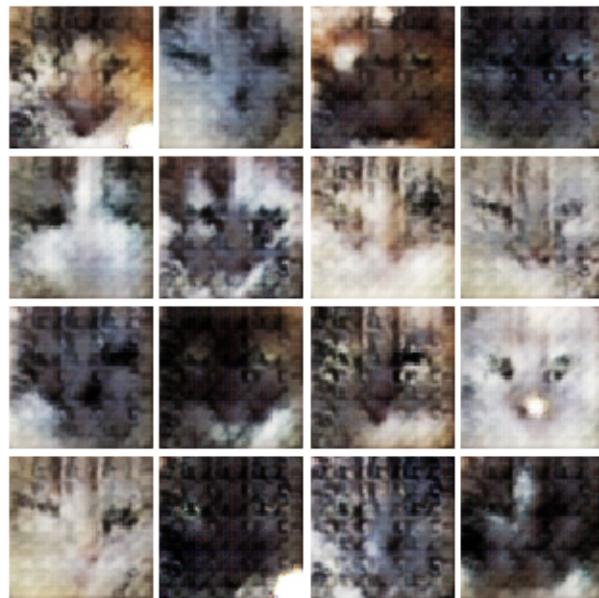
EPOCH: 5
Iter: 1000, D: 1.196, G:1.126



EPOCH: 6
Iter: 1250, D: 1.032, G: 1.983



EPOCH: 7
Iter: 1500, D: 1.266, G: 1.015



EPOCH: 8
Iter: 1750, D: 1.547, G:1.839

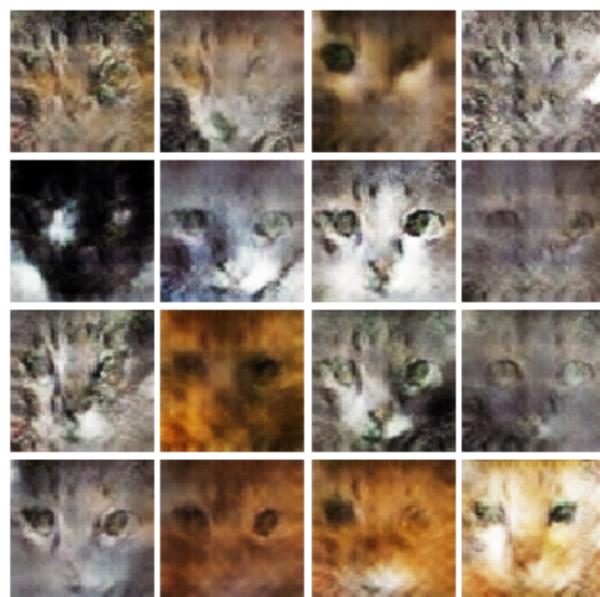


EPOCH: 9
Iter: 2000, D: 0.9348, G:1.46



EPOCH: 10

Iter: 2250, D: 1.001, G:1.955



EPOCH: 11

Iter: 2500, D: 0.738, G:2.569



EPOCH: 12
Iter: 2750, D: 1.719, G:1.755



EPOCH: 13
Iter: 3000, D: 0.4123, G:2.399



EPOCH: 14

Iter: 3250, D: 0.7801, G:3.544



EPOCH: 15

Iter: 3500, D: 0.6437, G:2.187



EPOCH: 16

Iter: 3750, D: 0.5209, G:3.502



EPOCH: 17

Iter: 4000, D: 0.3493, G:3.831



EPOCH: 18

Iter: 4250, D: 0.3992, G: 4.406



EPOCH: 19

Iter: 4500, D: 0.4835, G: 3.697



EPOCH: 20
Iter: 4750, D: 0.6429, G:2.772



EPOCH: 21
Iter: 5000, D: 0.6385, G:6.887



EPOCH: 22
Iter: 5250, D: 0.246, G:3.861



EPOCH: 23
Iter: 5500, D: 0.3747, G:3.706



EPOCH: 24

Iter: 5750, D: 0.308, G:3.607



EPOCH: 25

Iter: 6000, D: 0.1213, G:4.931



EPOCH: 26
Iter: 6250, D: 0.505, G:8.478



EPOCH: 27
Iter: 6500, D: 0.3444, G:3.813



EPOCH: 28

Iter: 6750, D: 0.3579, G:3.171



EPOCH: 29

Iter: 7000, D: 1.789, G:1.996



EPOCH: 30

Iter: 7250, D: 0.2094, G: 4.823



EPOCH: 31

Iter: 7500, D: 0.2543, G: 4.393



EPOCH: 32
Iter: 7750, D: 1.084, G:7.711



EPOCH: 33
Iter: 8000, D: 0.1085, G:5.36



EPOCH: 34

Iter: 8250, D: 0.2684, G:3.894



EPOCH: 35

Iter: 8500, D: 0.5873, G:1.972



EPOCH: 36

Iter: 8750, D: 0.2957, G:1.519



EPOCH: 37

Iter: 9000, D: 1.058, G:11.42



EPOCH: 38

Iter: 9250, D: 0.4308, G: 4.428



EPOCH: 39

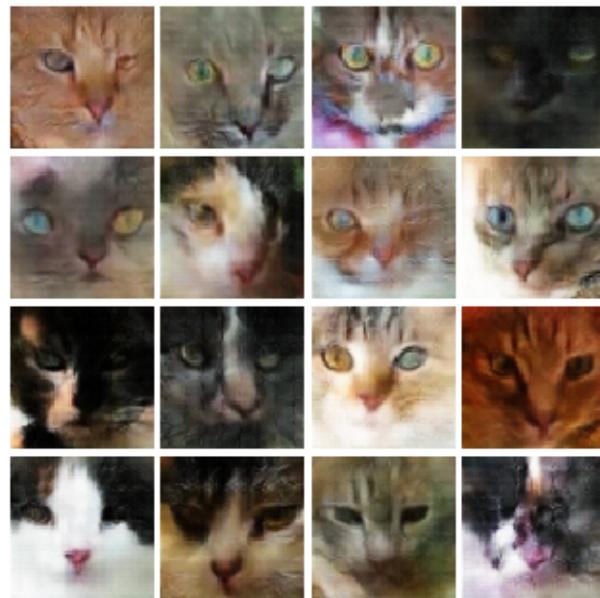
Iter: 9500, D: 0.3184, G: 5.904



EPOCH: 40
Iter: 9750, D: 0.3222, G:6.536



EPOCH: 41
Iter: 10000, D: 0.3402, G:5.17



EPOCH: 42

Iter: 10250, D: 0.1742, G: 5.425



EPOCH: 43

Iter: 10500, D: 0.148, G: 5.973



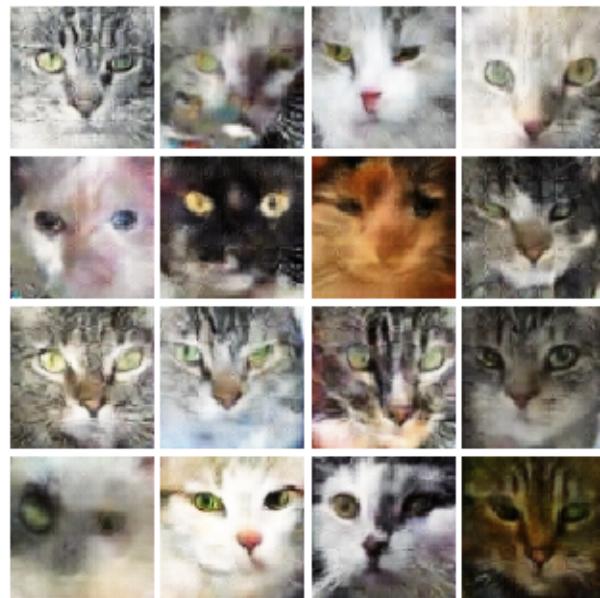
EPOCH: 44

Iter: 10750, D: 0.1794, G:6.006



EPOCH: 45

Iter: 11000, D: 0.3815, G:4.111



EPOCH: 46

Iter: 11250, D: 0.3606, G:4.85



EPOCH: 47

Iter: 11500, D: 0.2876, G:4.475



EPOCH: 48

Iter: 11750, D: 0.1511, G:5.33



EPOCH: 49

Iter: 12000, D: 0.2331, G:5.303



EPOCH: 50

Iter: 12250, D: 0.4051, G: 4.777



9 Extra Credit 3: Another Dataset

```
[ ]: ### Load Dog Dataset ###
batch_size = 64
imsize = 64
dog_root = './dog'

dog_train = ImageFolder(root=dog_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imszie)),
    transforms.RandomCrop(imszie),
]))
dog_loader_train = DataLoader(dog_train, batch_size=batch_size, drop_last=True)
```

```
[ ]: ### Visualize ###
from gan.utils import show_images

try:
    imgs = next(iter(dog_loader_train))[0].numpy().squeeze()
except:
    imgs = dog_loader_train.__iter__().next()[0].numpy().squeeze()

show_images(imgs, color=True)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



9.0.1 Training

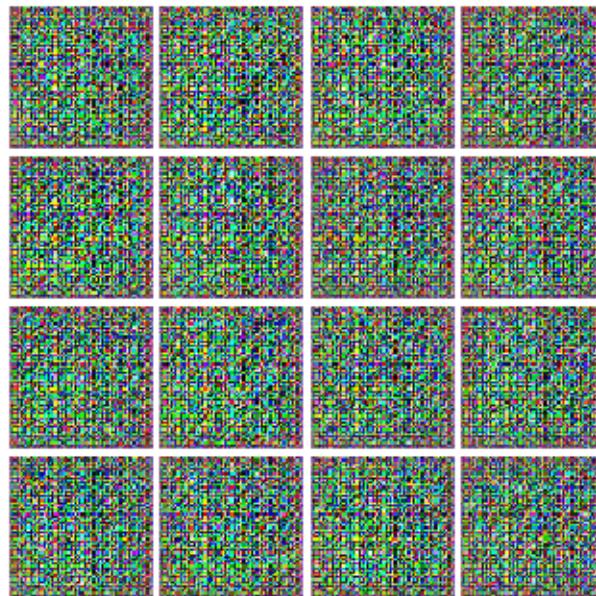
```
[ ]: NOISE_DIM = 100  
      NUM_EPOCHS = 200  
      learning_rate = 0.001
```

```
[ ]: D = Discriminator().to(device)  
      G = Generator(noise_dim=NOISE_DIM).to(device)
```

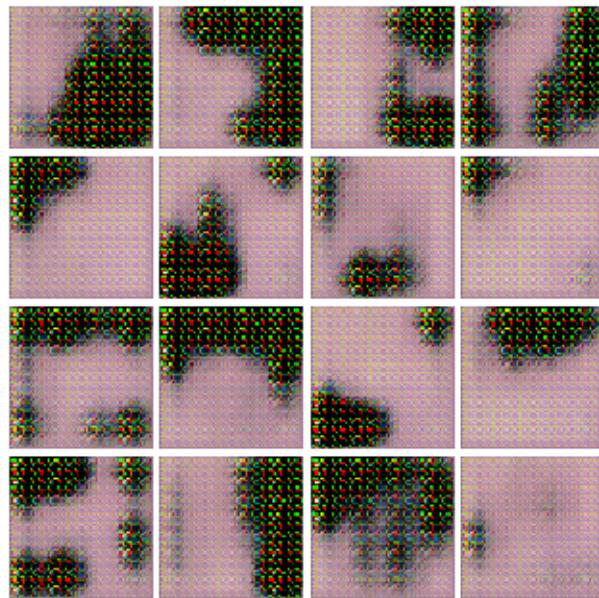
```
[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))  
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
[ ]: train(D, G, D_optimizer, G_optimizer, discriminator_loss,  
         generator_loss, num_epochs=NUM_EPOCHS, show_every=50,  
         batch_size=batch_size, train_loader=dog_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 1.455, G:8.388



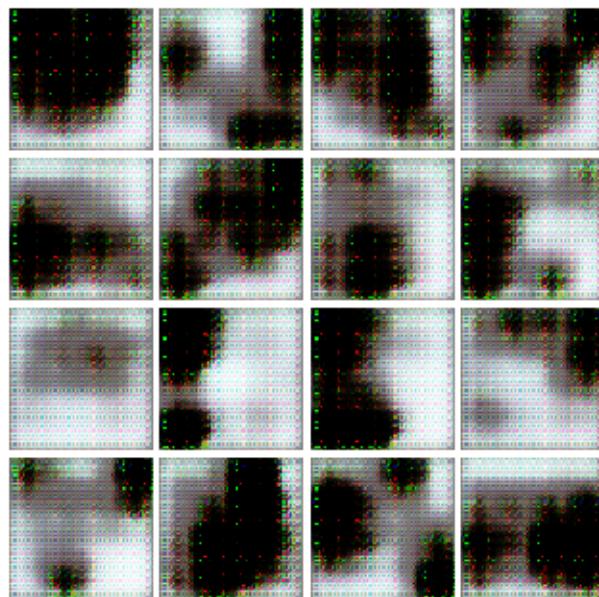
EPOCH: 2
EPOCH: 3
Iter: 50, D: 1.382, G:1.601



EPOCH: 4

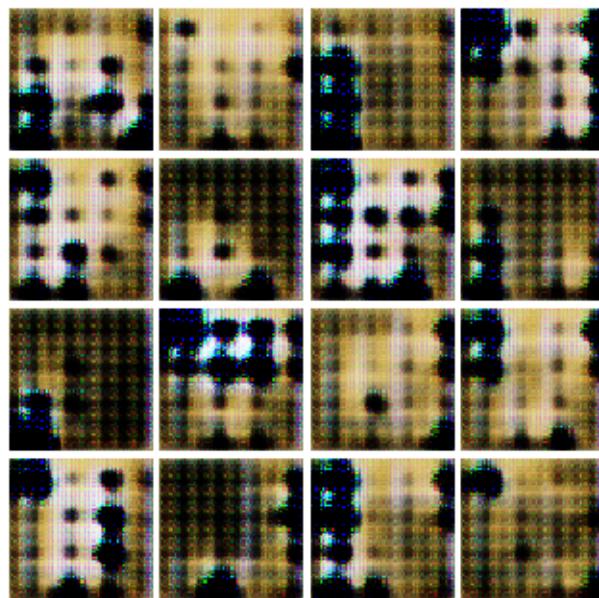
EPOCH: 5

Iter: 100, D: 0.8467, G: 5.018

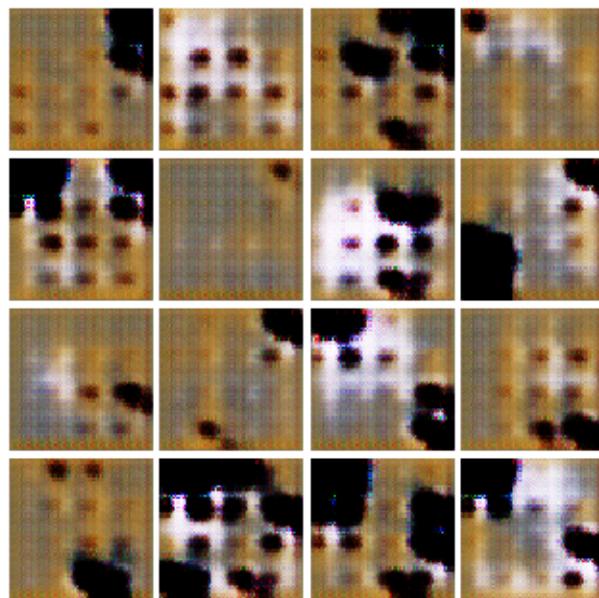


EPOCH: 6

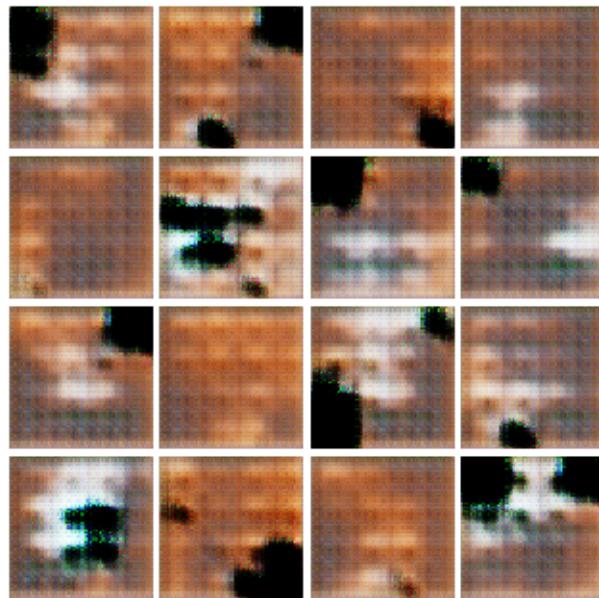
EPOCH: 7
Iter: 150, D: 0.9229, G:1.974



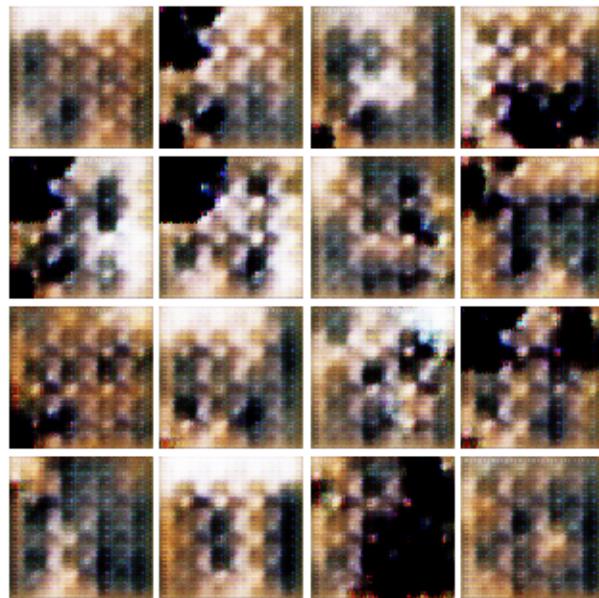
EPOCH: 8
EPOCH: 9
EPOCH: 10
Iter: 200, D: 0.9069, G:1.367



EPOCH: 11
EPOCH: 12
Iter: 250, D: 0.9697, G:2.671



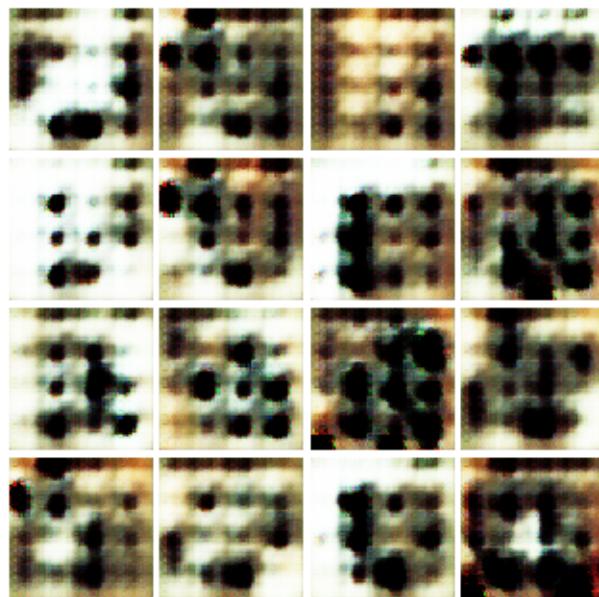
EPOCH: 13
EPOCH: 14
Iter: 300, D: 0.2901, G:2.823



EPOCH: 15

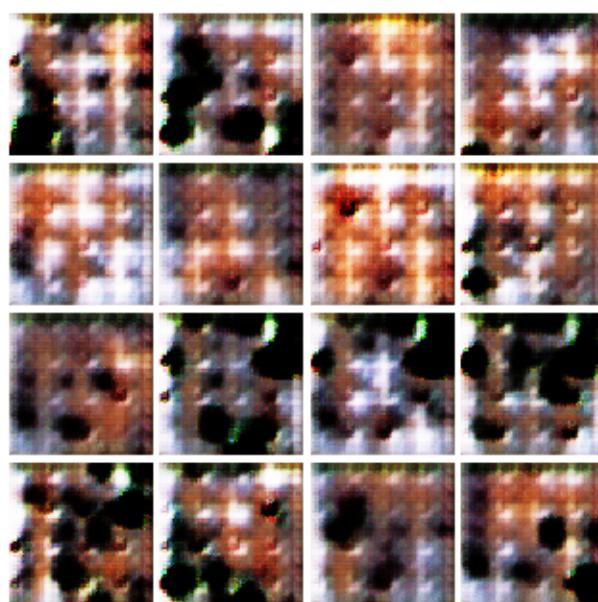
EPOCH: 16

Iter: 350, D: 0.6936, G: 3.675

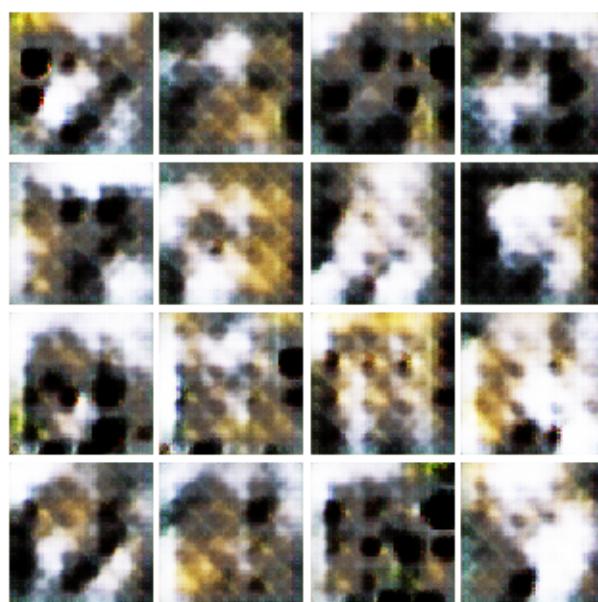


EPOCH: 17

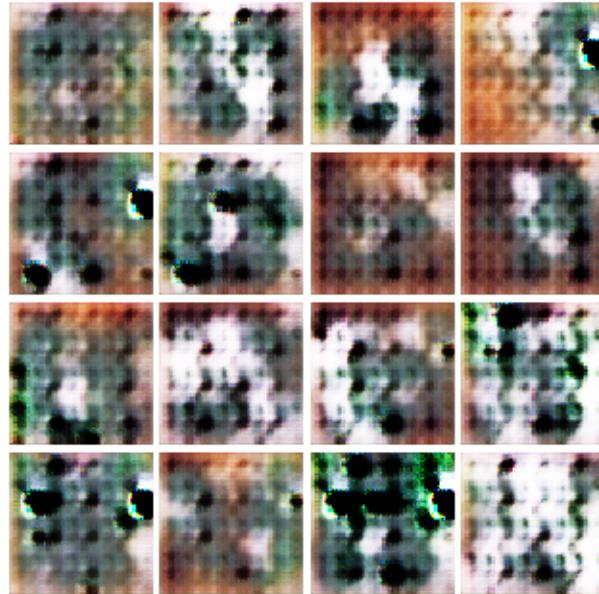
EPOCH: 18
EPOCH: 19
Iter: 400, D: 0.8892, G:4.959



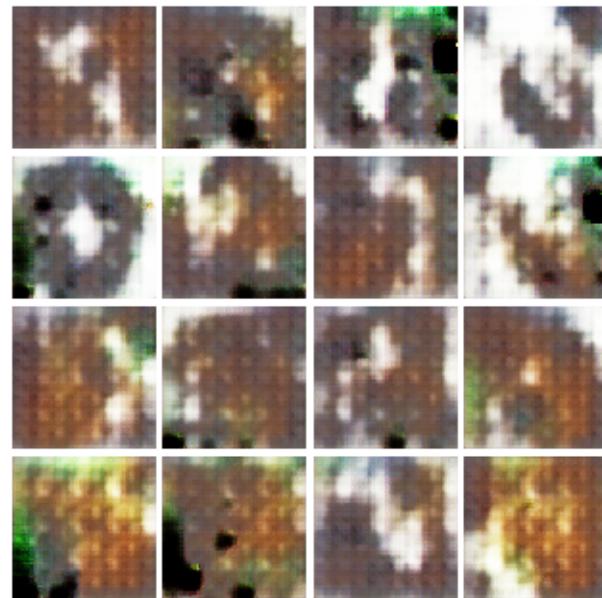
EPOCH: 20
EPOCH: 21
Iter: 450, D: 0.8315, G:1.899



EPOCH: 22
EPOCH: 23
Iter: 500, D: 0.5564, G:4.372



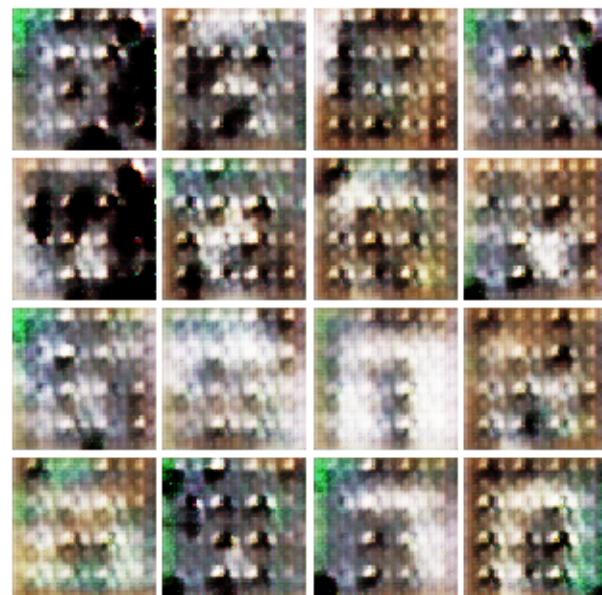
EPOCH: 24
EPOCH: 25
EPOCH: 26
Iter: 550, D: 0.8471, G:2.774



EPOCH: 27

EPOCH: 28

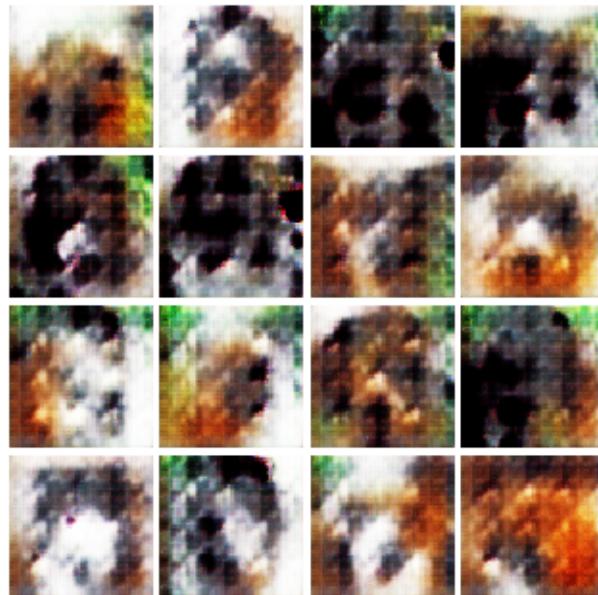
Iter: 600, D: 0.4462, G: 3.297



EPOCH: 29

EPOCH: 30

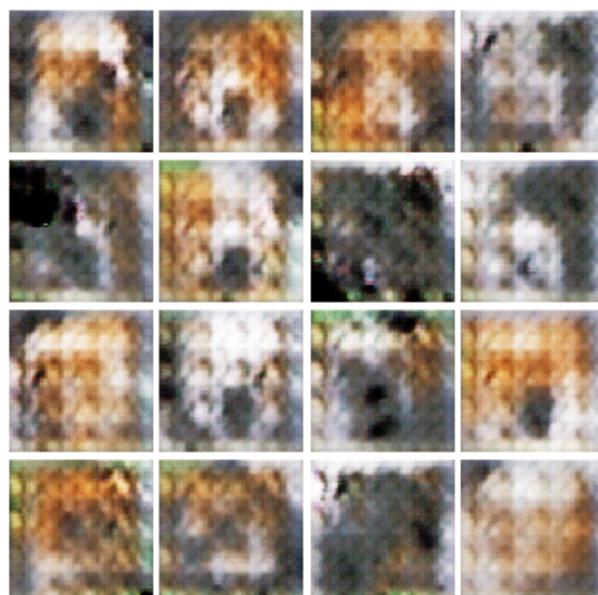
Iter: 650, D: 0.5287, G:2.792



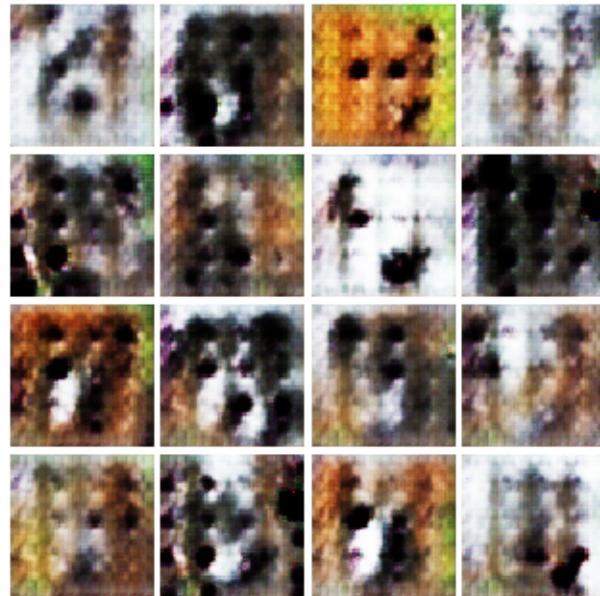
EPOCH: 31

EPOCH: 32

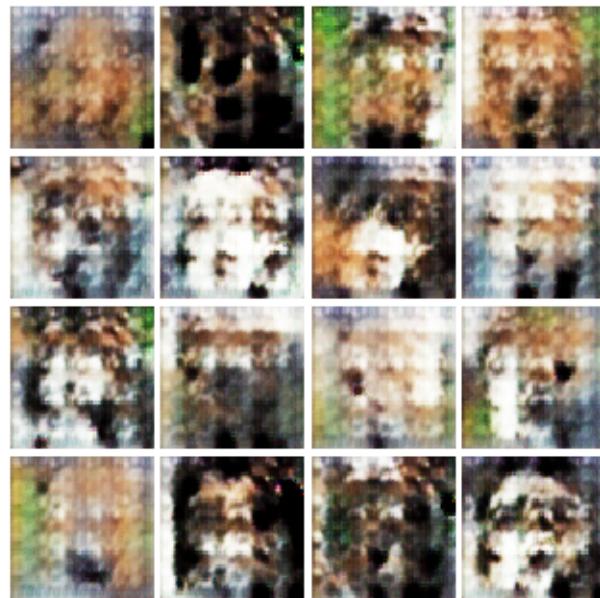
Iter: 700, D: 1.807, G:3.081



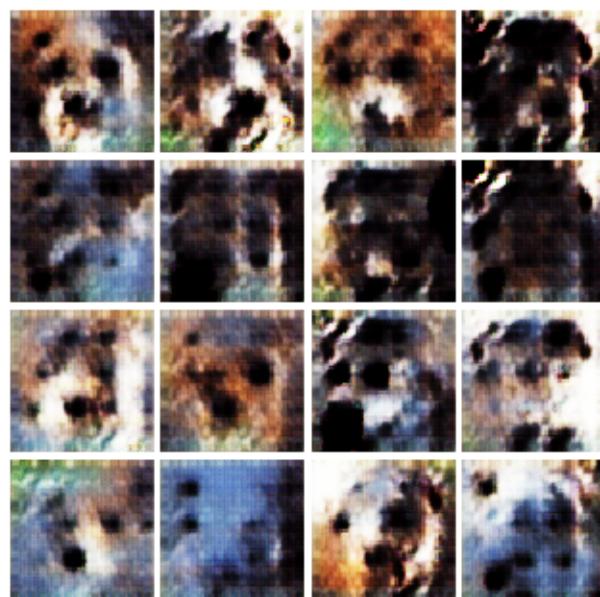
EPOCH: 33
EPOCH: 34
EPOCH: 35
Iter: 750, D: 0.7818, G:1.552



EPOCH: 36
EPOCH: 37
Iter: 800, D: 0.4931, G:4.509

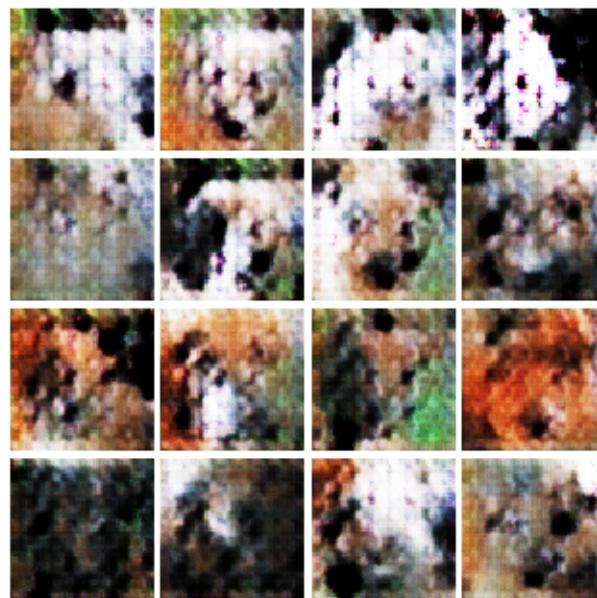


EPOCH: 38
EPOCH: 39
Iter: 850, D: 0.8428, G: 3.586

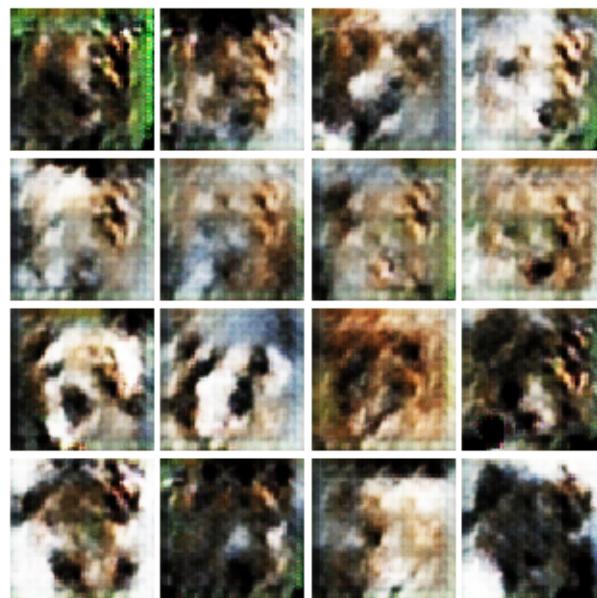


EPOCH: 40

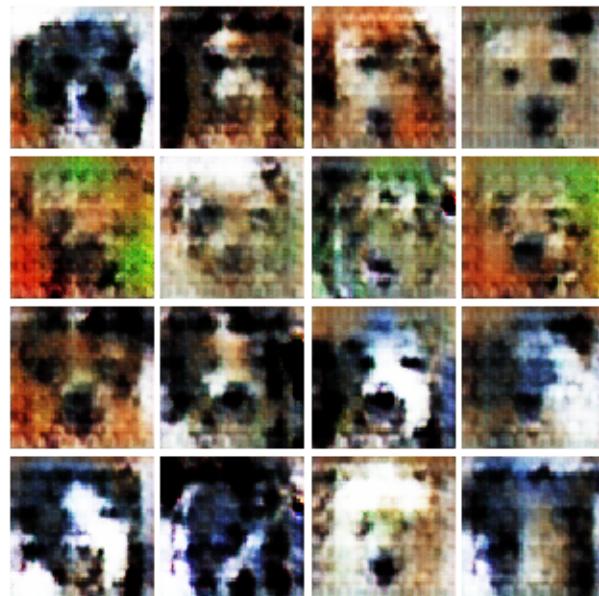
EPOCH: 41
Iter: 900, D: 1.053, G:0.9006



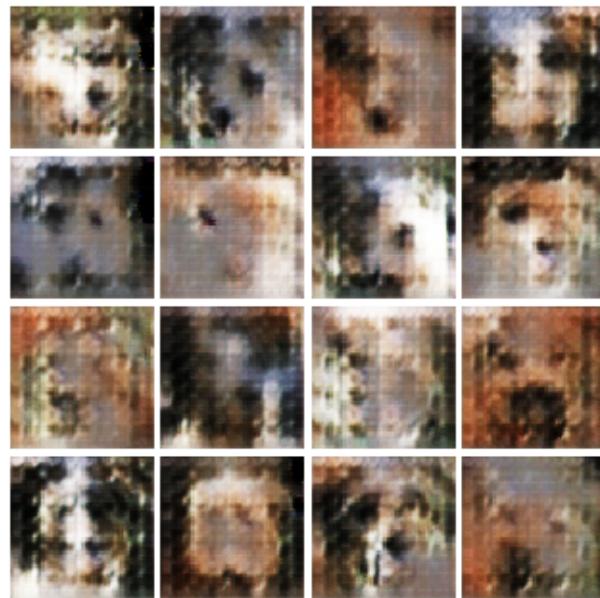
EPOCH: 42
EPOCH: 43
EPOCH: 44
Iter: 950, D: 0.4252, G:3.805



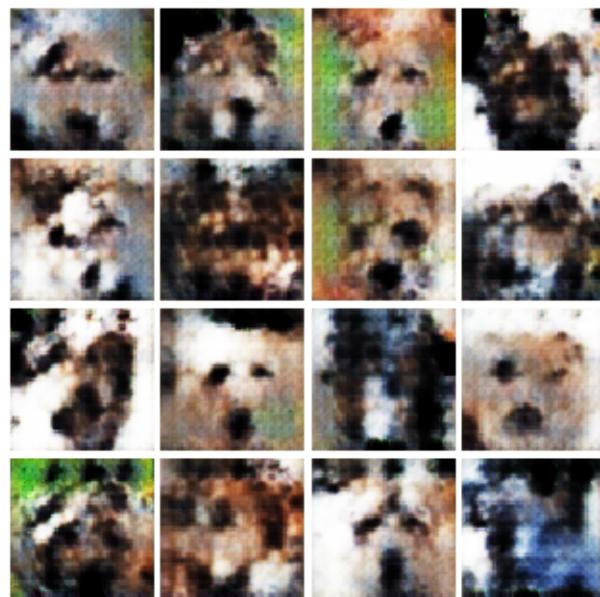
EPOCH: 45
EPOCH: 46
Iter: 1000, D: 1.081, G:2.533



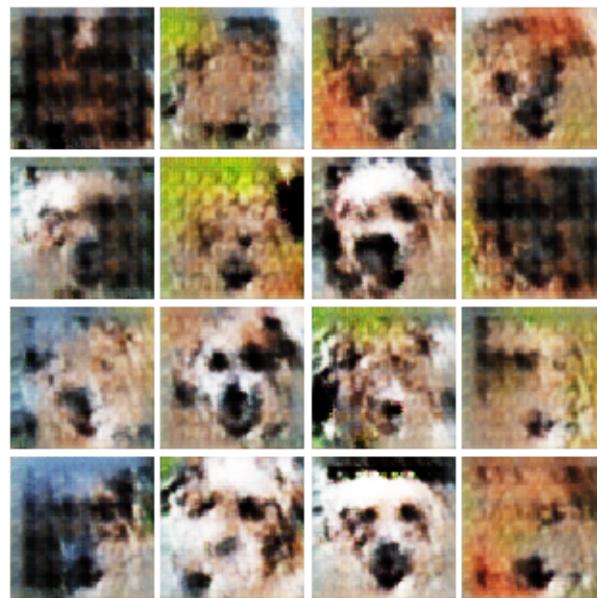
EPOCH: 47
EPOCH: 48
Iter: 1050, D: 1.886, G:4.949



EPOCH: 49
EPOCH: 50
EPOCH: 51
Iter: 1100, D: 0.6, G: 2.044



EPOCH: 52
EPOCH: 53
Iter: 1150, D: 0.4634, G:1.609



EPOCH: 54
EPOCH: 55
Iter: 1200, D: 0.6041, G:4.063



EPOCH: 56

EPOCH: 57

Iter: 1250, D: 0.3431, G:1.242



EPOCH: 58

EPOCH: 59

EPOCH: 60

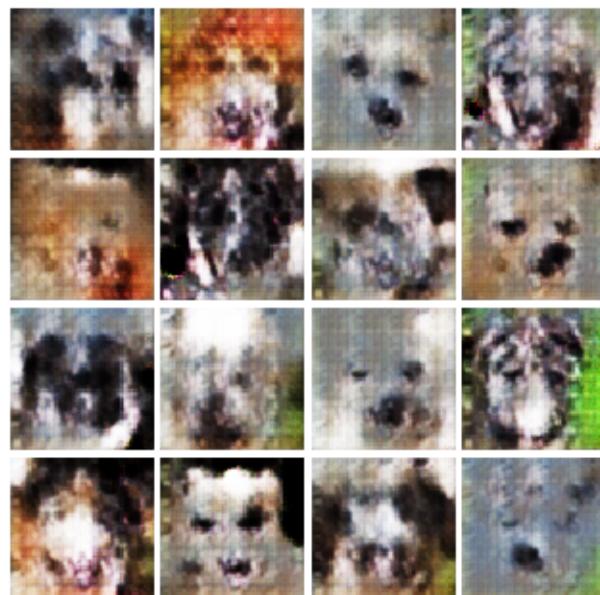
Iter: 1300, D: 0.603, G:0.9618



EPOCH: 61

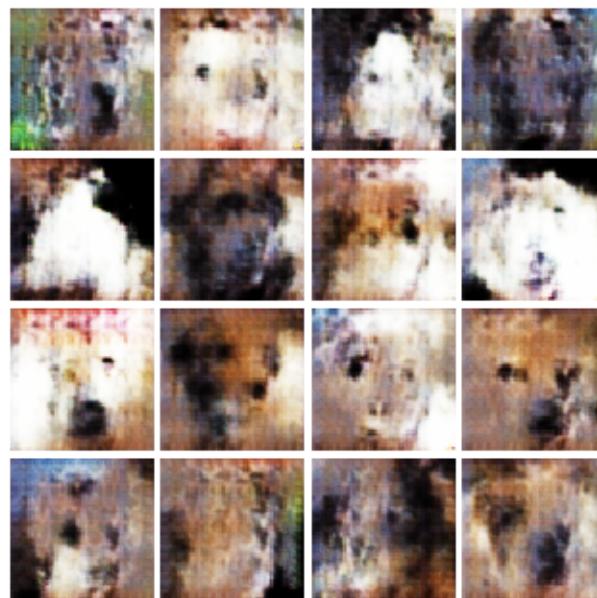
EPOCH: 62

Iter: 1350, D: 0.3746, G: 1.257

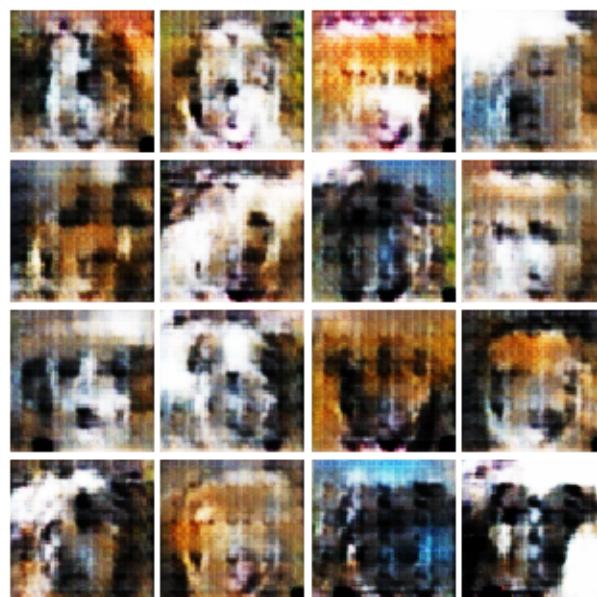


EPOCH: 63

EPOCH: 64
Iter: 1400, D: 0.8191, G:3.073



EPOCH: 65
EPOCH: 66
Iter: 1450, D: 0.6467, G:2.953

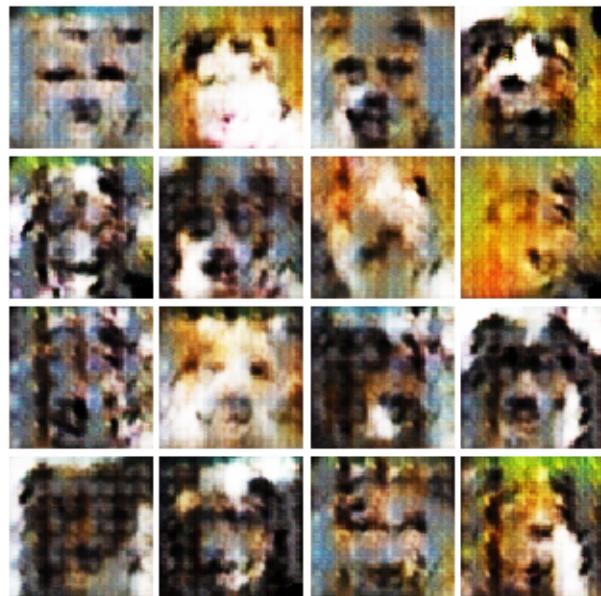


EPOCH: 67

EPOCH: 68

EPOCH: 69

Iter: 1500, D: 0.9179, G:2.315



EPOCH: 70

EPOCH: 71

Iter: 1550, D: 0.8369, G:1.526



EPOCH: 72

EPOCH: 73

Iter: 1600, D: 0.4643, G: 2.914



EPOCH: 74

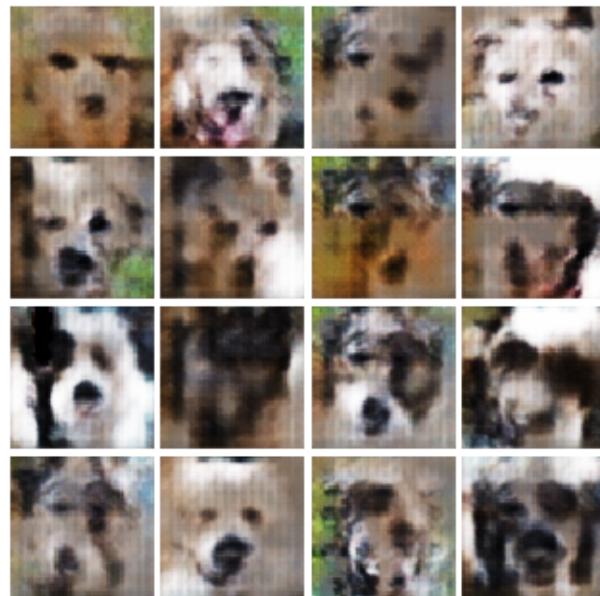
EPOCH: 75
EPOCH: 76
Iter: 1650, D: 0.6531, G:2.01



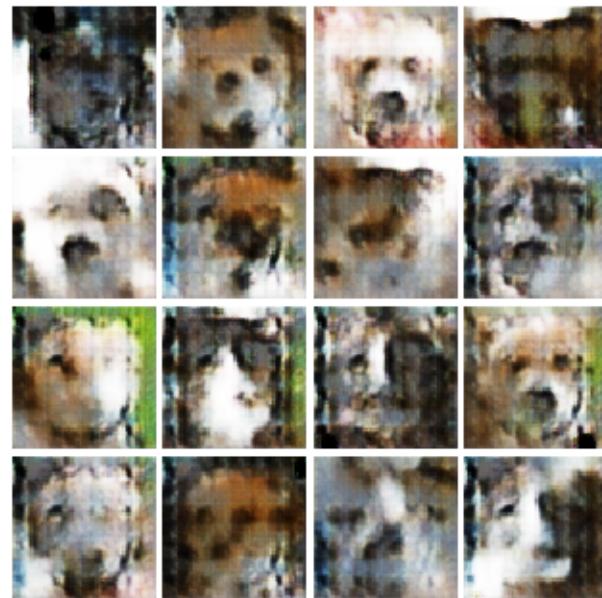
EPOCH: 77
EPOCH: 78
Iter: 1700, D: 0.5539, G:3.16



EPOCH: 79
EPOCH: 80
Iter: 1750, D: 0.3326, G:2.7



EPOCH: 81
EPOCH: 82
Iter: 1800, D: 0.408, G:3.018



EPOCH: 83

EPOCH: 84

EPOCH: 85

Iter: 1850, D: 0.7723, G: 6.317



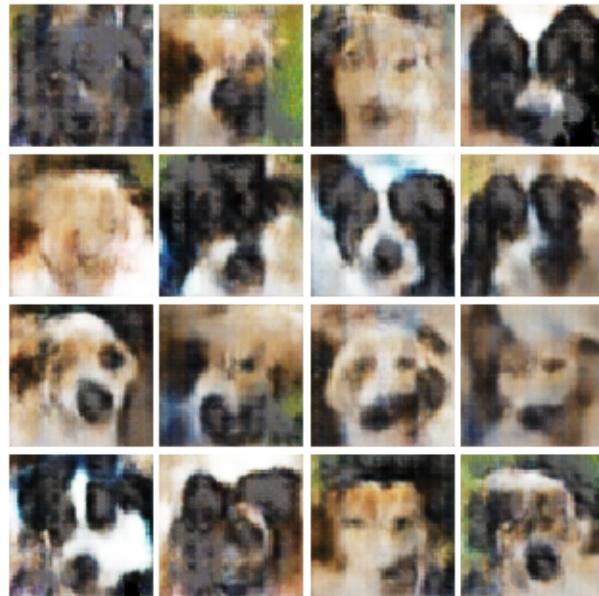
EPOCH: 86
EPOCH: 87
Iter: 1900, D: 0.5553, G:2.233



EPOCH: 88
EPOCH: 89
Iter: 1950, D: 0.7804, G:2.607



EPOCH: 90
EPOCH: 91
Iter: 2000, D: 0.6392, G:2.428



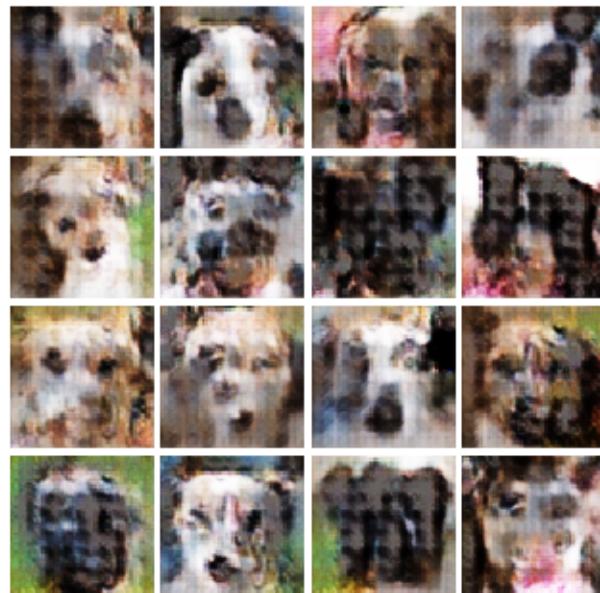
EPOCH: 92
EPOCH: 93
EPOCH: 94
Iter: 2050, D: 1.206, G:6.978



EPOCH: 95

EPOCH: 96

Iter: 2100, D: 0.6119, G:1.355



EPOCH: 97

EPOCH: 98
Iter: 2150, D: 0.646, G:2.051



EPOCH: 99
EPOCH: 100
EPOCH: 101
Iter: 2200, D: 0.3058, G:2.554



EPOCH: 102

EPOCH: 103

Iter: 2250, D: 0.3878, G:4.774



EPOCH: 104

EPOCH: 105

Iter: 2300, D: 1.085, G:3.305



EPOCH: 106

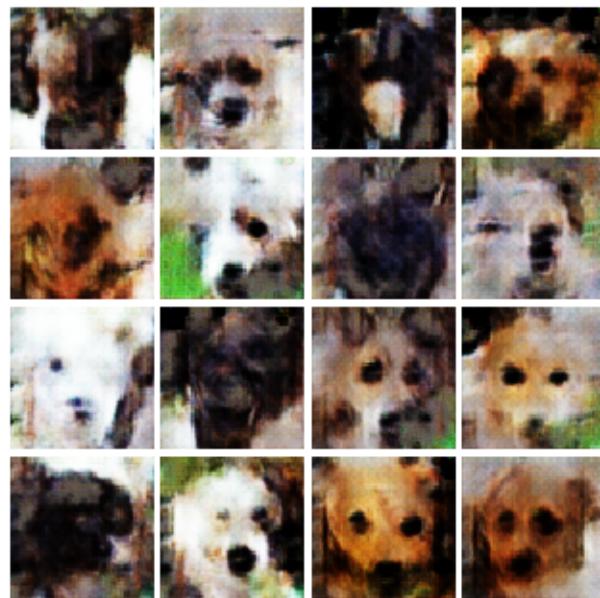
EPOCH: 107

Iter: 2350, D: 0.3186, G: 2.482



EPOCH: 108

EPOCH: 109
EPOCH: 110
Iter: 2400, D: 0.751, G:1.894



EPOCH: 111
EPOCH: 112
Iter: 2450, D: 0.7278, G:5.115



EPOCH: 113

EPOCH: 114

Iter: 2500, D: 0.4007, G:2.998



EPOCH: 115

EPOCH: 116

Iter: 2550, D: 0.5138, G:2.863



EPOCH: 117

EPOCH: 118

EPOCH: 119

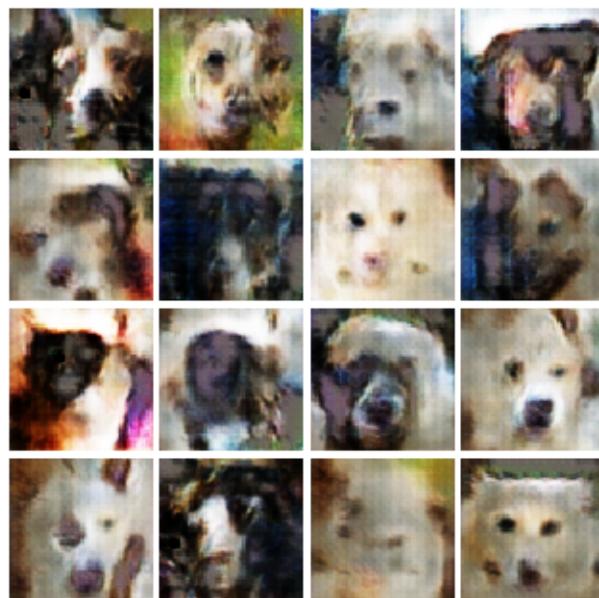
Iter: 2600, D: 0.7732, G: 4.785



EPOCH: 120
EPOCH: 121
Iter: 2650, D: 0.4183, G:2.374



EPOCH: 122
EPOCH: 123
Iter: 2700, D: 0.5076, G:4.684



EPOCH: 124
EPOCH: 125
EPOCH: 126
Iter: 2750, D: 0.2939, G:4.975



EPOCH: 127
EPOCH: 128
Iter: 2800, D: 0.477, G:4.369



EPOCH: 129

EPOCH: 130

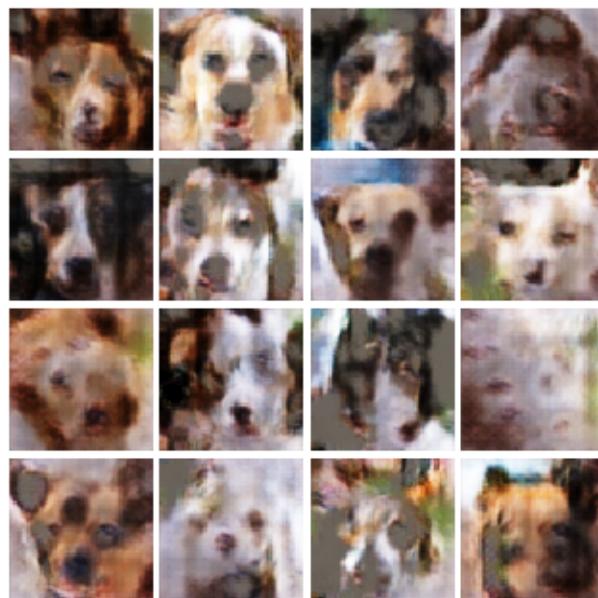
Iter: 2850, D: 0.4775, G: 2.355



EPOCH: 131

EPOCH: 132

Iter: 2900, D: 0.3916, G:4.001

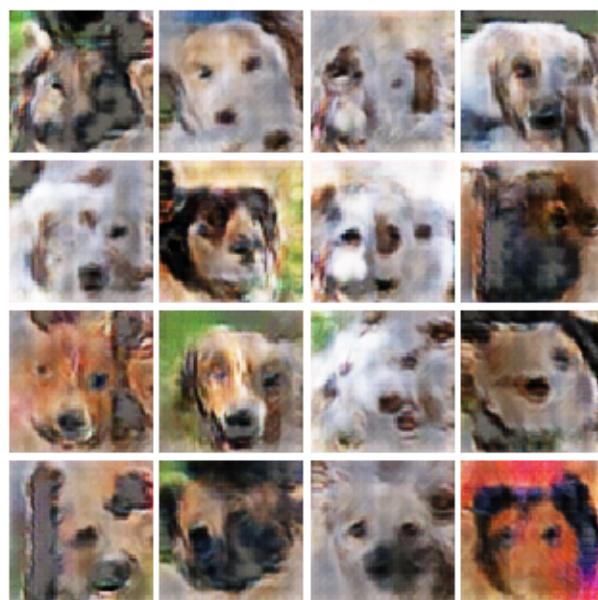


EPOCH: 133

EPOCH: 134

EPOCH: 135

Iter: 2950, D: 0.562, G:5.788



EPOCH: 136

EPOCH: 137

Iter: 3000, D: 0.3671, G:3.164



EPOCH: 138

EPOCH: 139

Iter: 3050, D: 1.058, G:4.071



EPOCH: 140

EPOCH: 141

Iter: 3100, D: 0.6665, G: 4.569



EPOCH: 142

EPOCH: 143

EPOCH: 144

Iter: 3150, D: 0.6846, G:3.193



EPOCH: 145

EPOCH: 146

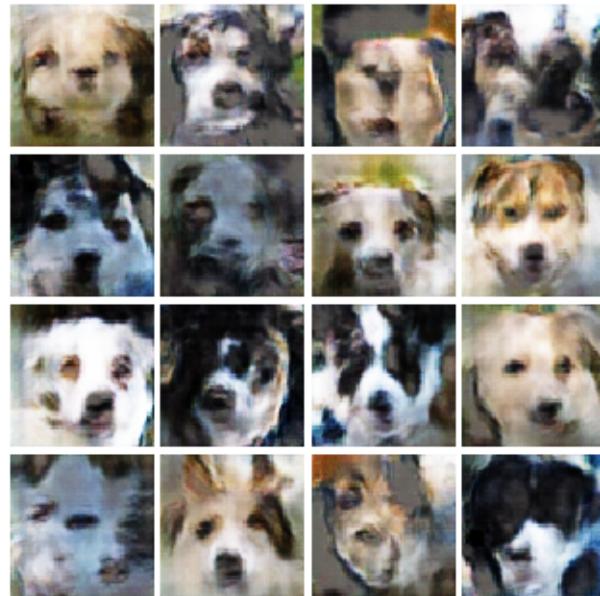
Iter: 3200, D: 0.5569, G:2.004



EPOCH: 147

EPOCH: 148

Iter: 3250, D: 0.4629, G:3.221



EPOCH: 149

EPOCH: 150

EPOCH: 151

Iter: 3300, D: 0.4111, G:3.769



EPOCH: 152

EPOCH: 153

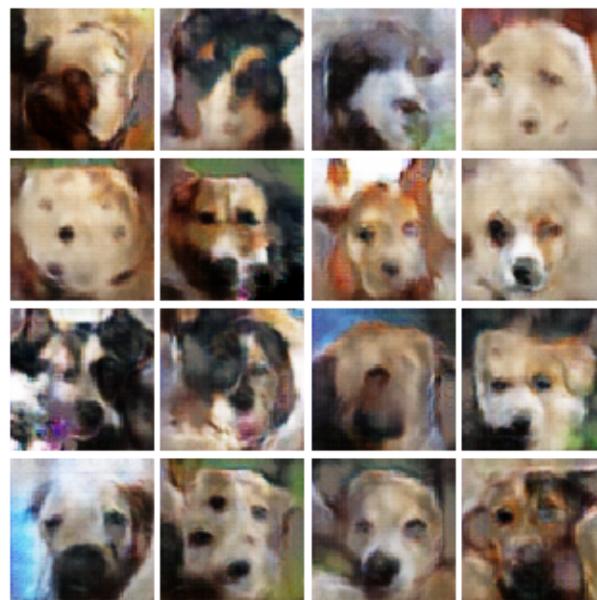
Iter: 3350, D: 0.764, G: 5.756



EPOCH: 154

EPOCH: 155

Iter: 3400, D: 0.4272, G: 3.783



EPOCH: 156

EPOCH: 157

Iter: 3450, D: 1.101, G: 1.684



EPOCH: 158

EPOCH: 159

EPOCH: 160

Iter: 3500, D: 0.5408, G:2.534



EPOCH: 161

EPOCH: 162

Iter: 3550, D: 0.605, G:3.486



EPOCH: 163

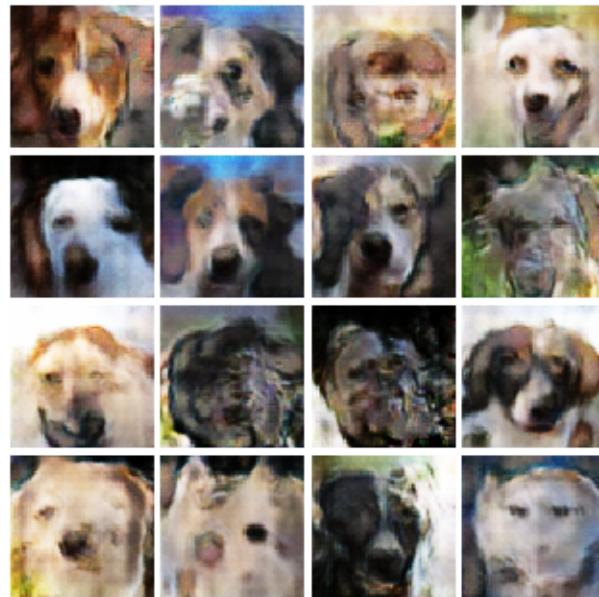
EPOCH: 164

Iter: 3600, D: 0.6403, G: 3.928



EPOCH: 165

EPOCH: 166
Iter: 3650, D: 0.52, G:2.047



EPOCH: 167
EPOCH: 168
EPOCH: 169
Iter: 3700, D: 1.784, G:5.401



EPOCH: 170

EPOCH: 171

Iter: 3750, D: 0.3819, G:3.795



EPOCH: 172

EPOCH: 173

Iter: 3800, D: 0.7515, G:2.263



EPOCH: 174

EPOCH: 175

EPOCH: 176

Iter: 3850, D: 1.114, G:2.539



EPOCH: 177
EPOCH: 178
Iter: 3900, D: 0.6965, G:4.483



EPOCH: 179
EPOCH: 180
Iter: 3950, D: 1.082, G:2.021



EPOCH: 181
EPOCH: 182
Iter: 4000, D: 2.558, G:4.945



EPOCH: 183
EPOCH: 184
EPOCH: 185
Iter: 4050, D: 0.3962, G:2.803



EPOCH: 186

EPOCH: 187

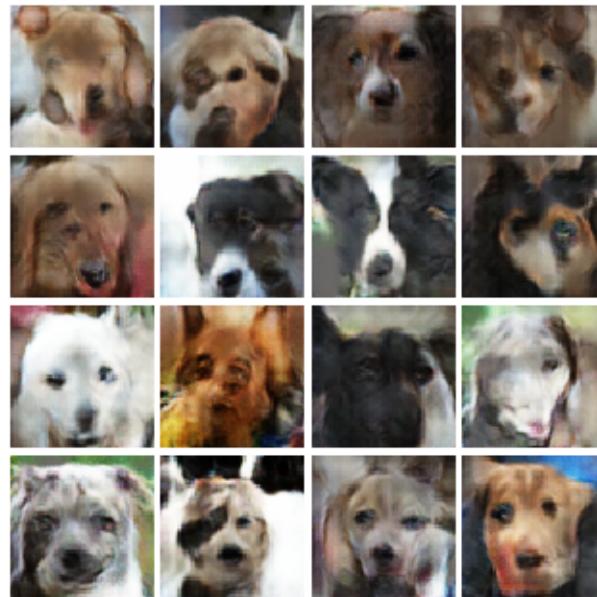
Iter: 4100, D: 0.5238, G: 2.256



EPOCH: 188

EPOCH: 189

Iter: 4150, D: 0.3937, G:3.496



EPOCH: 190

EPOCH: 191

Iter: 4200, D: 0.7048, G:3.374



EPOCH: 192
EPOCH: 193
EPOCH: 194
Iter: 4250, D: 0.4764, G:2.677



EPOCH: 195
EPOCH: 196
Iter: 4300, D: 0.3797, G:3.075



EPOCH: 197

EPOCH: 198

Iter: 4350, D: 0.4082, G: 3.572



EPOCH: 199

EPOCH: 200

```
[ ]: train(D, G, D_optimizer, G_optimizer, discriminator_loss,
           generator_loss, num_epochs=NUM_EPOCHS, show_every=50,
           batch_size=batch_size, train_loader=dog_loader_train, device=device)
```

EPOCH: 1

Iter: 0, D: 0.7443, G:4.264



EPOCH: 2

EPOCH: 3

Iter: 50, D: 0.5131, G:2.586



EPOCH: 4

EPOCH: 5

Iter: 100, D: 1.279, G: 3.748



EPOCH: 6

EPOCH: 7
Iter: 150, D: 0.5254, G:3.507



EPOCH: 8
EPOCH: 9
EPOCH: 10
Iter: 200, D: 0.4195, G:3.065



EPOCH: 11
EPOCH: 12
Iter: 250, D: 0.3574, G:2.902



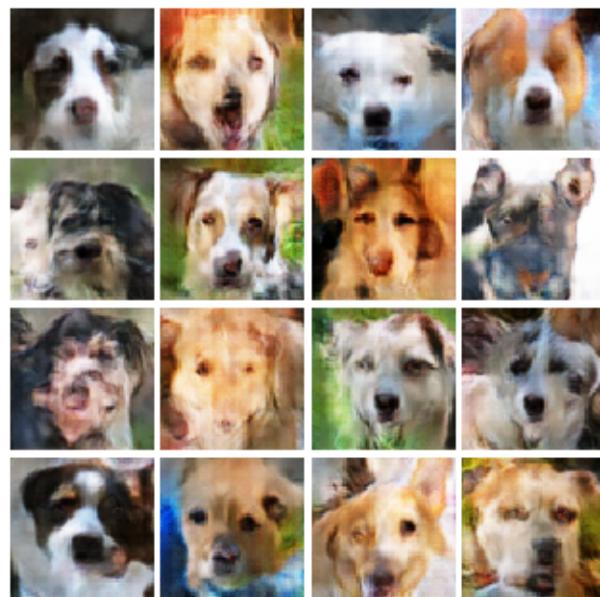
EPOCH: 13
EPOCH: 14
Iter: 300, D: 0.5297, G:1.432



EPOCH: 15

EPOCH: 16

Iter: 350, D: 0.6446, G: 3.379



EPOCH: 17

EPOCH: 18
EPOCH: 19
Iter: 400, D: 0.3581, G:3.601



EPOCH: 20
EPOCH: 21
Iter: 450, D: 0.6807, G:2.878



EPOCH: 22
EPOCH: 23
Iter: 500, D: 0.6603, G:1.655



EPOCH: 24
EPOCH: 25
EPOCH: 26
Iter: 550, D: 0.3963, G:3.366



EPOCH: 27

EPOCH: 28

Iter: 600, D: 1.226, G:1.32



EPOCH: 29

EPOCH: 30
Iter: 650, D: 0.3361, G:4.004



EPOCH: 31
EPOCH: 32
Iter: 700, D: 0.4197, G:3.824



EPOCH: 33
EPOCH: 34
EPOCH: 35
Iter: 750, D: 0.4287, G:3.131



EPOCH: 36
EPOCH: 37
Iter: 800, D: 0.6259, G:5.687



EPOCH: 38

EPOCH: 39

Iter: 850, D: 0.7528, G: 3.018



EPOCH: 40

EPOCH: 41
Iter: 900, D: 0.6188, G:2.393



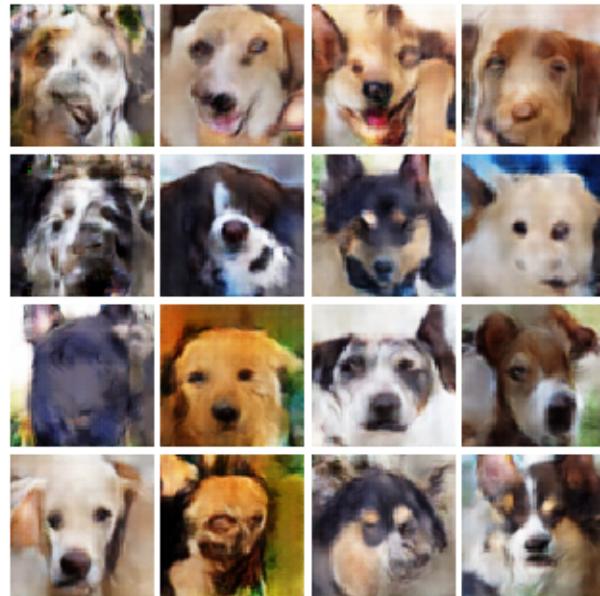
EPOCH: 42
EPOCH: 43
EPOCH: 44
Iter: 950, D: 0.6318, G:1.753



EPOCH: 45

EPOCH: 46

Iter: 1000, D: 0.6878, G:2.114



EPOCH: 47

EPOCH: 48

Iter: 1050, D: 0.3919, G:2.479



EPOCH: 49

EPOCH: 50

EPOCH: 51

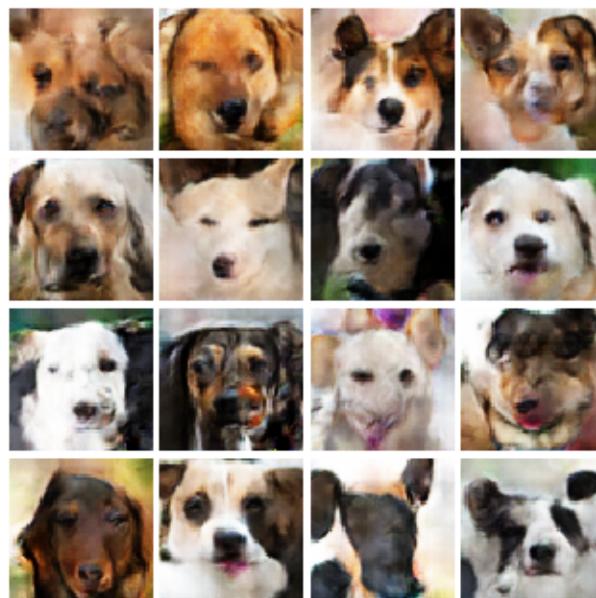
Iter: 1100, D: 0.8548, G: 3.186



EPOCH: 52
EPOCH: 53
Iter: 1150, D: 0.4915, G:3.245



EPOCH: 54
EPOCH: 55
Iter: 1200, D: 0.3002, G:3.114



EPOCH: 56

EPOCH: 57

Iter: 1250, D: 0.3677, G:3.624



EPOCH: 58

EPOCH: 59

EPOCH: 60

Iter: 1300, D: 0.3864, G:2.004



EPOCH: 61

EPOCH: 62

Iter: 1350, D: 0.3095, G: 3.226



EPOCH: 63

EPOCH: 64
Iter: 1400, D: 1.028, G:0.7783



EPOCH: 65
EPOCH: 66
Iter: 1450, D: 1.317, G:1.13



EPOCH: 67

EPOCH: 68

EPOCH: 69

Iter: 1500, D: 0.3897, G:2.675



EPOCH: 70

EPOCH: 71

Iter: 1550, D: 0.396, G:2.067



EPOCH: 72

EPOCH: 73

Iter: 1600, D: 0.3293, G: 2.819



EPOCH: 74

EPOCH: 75
EPOCH: 76
Iter: 1650, D: 0.6036, G:4.496



EPOCH: 77
EPOCH: 78
Iter: 1700, D: 0.4071, G:3.638



EPOCH: 79

EPOCH: 80

Iter: 1750, D: 0.4616, G:4.671



EPOCH: 81

EPOCH: 82

Iter: 1800, D: 0.3699, G:3.5



EPOCH: 83

EPOCH: 84

EPOCH: 85

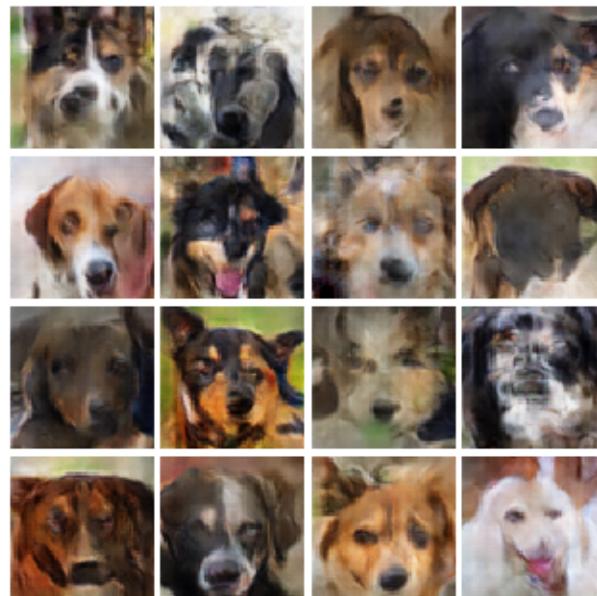
Iter: 1850, D: 0.3505, G: 1.984



EPOCH: 86
EPOCH: 87
Iter: 1900, D: 0.6654, G:1.913



EPOCH: 88
EPOCH: 89
Iter: 1950, D: 0.5576, G:3.593



EPOCH: 90

EPOCH: 91

Iter: 2000, D: 0.5075, G:2.305



EPOCH: 92

EPOCH: 93

EPOCH: 94

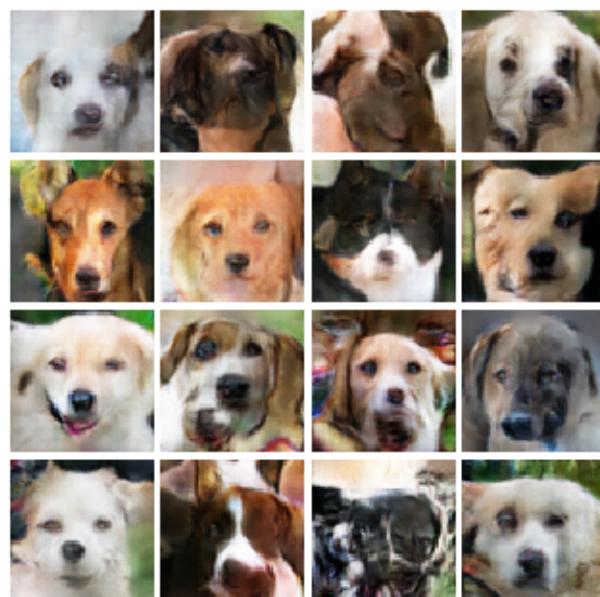
Iter: 2050, D: 0.2616, G:2.831



EPOCH: 95

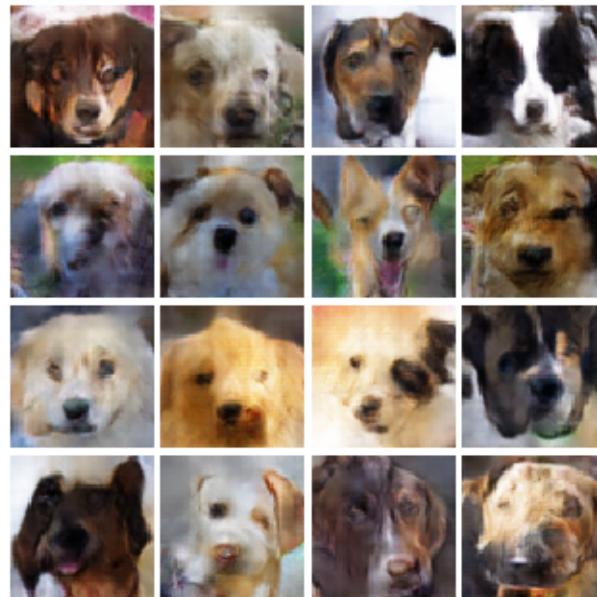
EPOCH: 96

Iter: 2100, D: 0.4052, G: 3.696



EPOCH: 97

EPOCH: 98
Iter: 2150, D: 0.7685, G:2.676



EPOCH: 99
EPOCH: 100
EPOCH: 101
Iter: 2200, D: 1.677, G:1.508



EPOCH: 102

EPOCH: 103

Iter: 2250, D: 0.7784, G:4.412



EPOCH: 104

EPOCH: 105

Iter: 2300, D: 0.2383, G:4.078



EPOCH: 106

EPOCH: 107

Iter: 2350, D: 0.2653, G: 3.742



EPOCH: 108

EPOCH: 109
EPOCH: 110
Iter: 2400, D: 0.2858, G:3.129



EPOCH: 111
EPOCH: 112
Iter: 2450, D: 4.503, G:9.163



EPOCH: 113
EPOCH: 114
Iter: 2500, D: 1.642, G:0.923



EPOCH: 115
EPOCH: 116
Iter: 2550, D: 1.464, G:1.076



EPOCH: 117

EPOCH: 118

EPOCH: 119

Iter: 2600, D: 1.294, G: 1.957



EPOCH: 120
EPOCH: 121
Iter: 2650, D: 0.7849, G:2.133



EPOCH: 122
EPOCH: 123
Iter: 2700, D: 0.4893, G:2.687



EPOCH: 124
EPOCH: 125
EPOCH: 126
Iter: 2750, D: 0.6195, G:5.036



EPOCH: 127
EPOCH: 128
Iter: 2800, D: 0.4943, G:2.019



EPOCH: 129

EPOCH: 130

Iter: 2850, D: 0.3445, G: 4.07



EPOCH: 131

EPOCH: 132
Iter: 2900, D: 0.3882, G: 4.418



EPOCH: 133
EPOCH: 134
EPOCH: 135
Iter: 2950, D: 0.5226, G: 3.061



EPOCH: 136

EPOCH: 137

Iter: 3000, D: 0.682, G:4.569



EPOCH: 138

EPOCH: 139

Iter: 3050, D: 0.4632, G:1.829



EPOCH: 140

EPOCH: 141

Iter: 3100, D: 0.3096, G: 2.905



EPOCH: 142

EPOCH: 143

EPOCH: 144

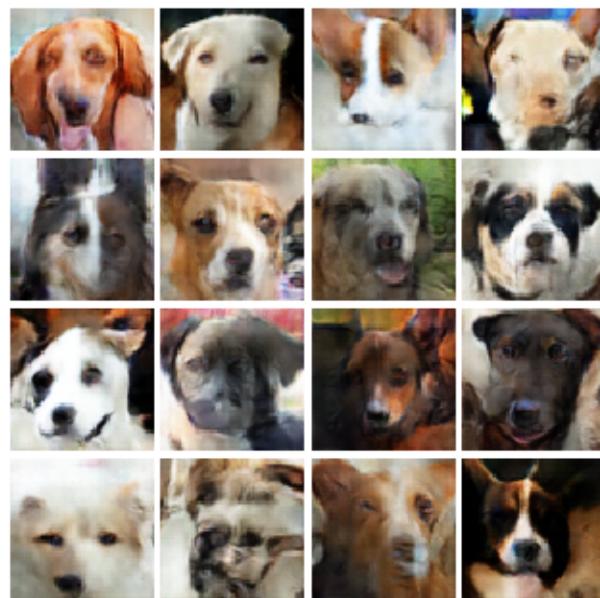
Iter: 3150, D: 0.2037, G:5.173



EPOCH: 145

EPOCH: 146

Iter: 3200, D: 0.2066, G:3.42



EPOCH: 147
EPOCH: 148
Iter: 3250, D: 1.697, G:0.955



EPOCH: 149
EPOCH: 150
EPOCH: 151
Iter: 3300, D: 1.528, G:0.7102



EPOCH: 152

EPOCH: 153

Iter: 3350, D: 1.475, G: 0.8106



EPOCH: 154

EPOCH: 155

Iter: 3400, D: 1.497, G:0.8644



EPOCH: 156

EPOCH: 157

Iter: 3450, D: 1.463, G:0.7675



EPOCH: 158
EPOCH: 159
EPOCH: 160
Iter: 3500, D: 1.461, G:0.9261



EPOCH: 161
EPOCH: 162
Iter: 3550, D: 1.432, G:0.8991



EPOCH: 163

EPOCH: 164

Iter: 3600, D: 1.402, G: 0.7002



EPOCH: 165

EPOCH: 166
Iter: 3650, D: 1.368, G:0.8707



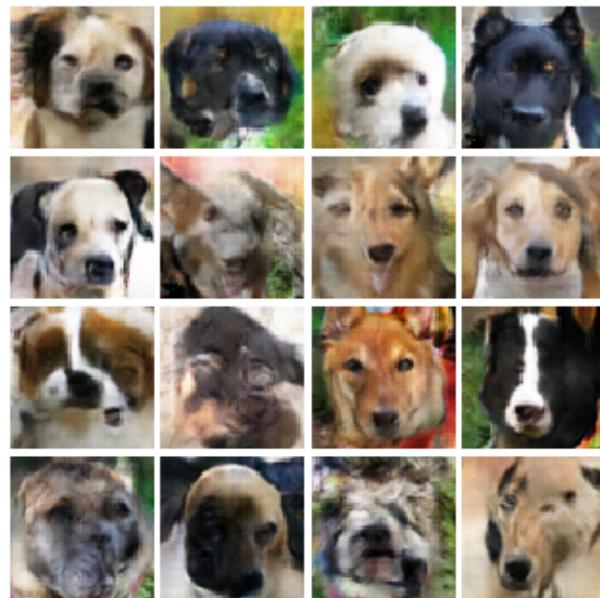
EPOCH: 167
EPOCH: 168
EPOCH: 169
Iter: 3700, D: 1.398, G:0.7297



EPOCH: 170
EPOCH: 171
Iter: 3750, D: 1.41, G:0.845



EPOCH: 172
EPOCH: 173
Iter: 3800, D: 1.384, G:0.8024



EPOCH: 174

EPOCH: 175

EPOCH: 176

Iter: 3850, D: 1.339, G: 0.756



EPOCH: 177
EPOCH: 178
Iter: 3900, D: 1.394, G:0.8869



EPOCH: 179
EPOCH: 180
Iter: 3950, D: 1.304, G:0.9702



EPOCH: 181
EPOCH: 182
Iter: 4000, D: 1.372, G:0.8883



EPOCH: 183
EPOCH: 184
EPOCH: 185
Iter: 4050, D: 1.422, G:0.9508



EPOCH: 186

EPOCH: 187

Iter: 4100, D: 1.275, G: 0.8154



EPOCH: 188

EPOCH: 189

Iter: 4150, D: 1.477, G: 0.9051



EPOCH: 190

EPOCH: 191

Iter: 4200, D: 1.374, G: 1.002



EPOCH: 192
EPOCH: 193
EPOCH: 194
Iter: 4250, D: 1.453, G:1.264



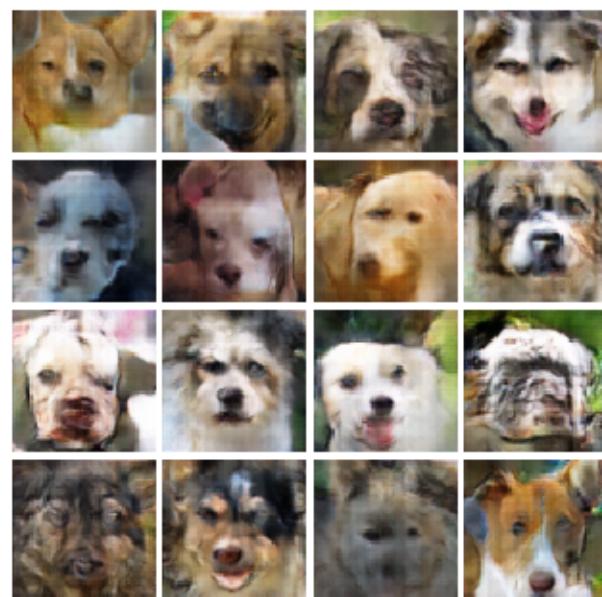
EPOCH: 195
EPOCH: 196
Iter: 4300, D: 1.16, G:2.844



EPOCH: 197

EPOCH: 198

Iter: 4350, D: 1.261, G: 4.243



EPOCH: 199

EPOCH: 200

10 Save to PDF

```
[ ]: %%capture

from google.colab import drive
drive.mount('/content/drive')
# install tex; first run may take several minutes
! apt-get install texlive-xetex
# file path and save location below are default; please change if they do not
# match yours
! jupyter nbconvert --output-dir='/content/drive/MyDrive/' '/content/drive/
MyDrive/CS444/assignment4/MP4.ipynb' --to pdf
```