

MP3_P2

April 2, 2024

```
[1]: # you will be prompted with a window asking to grant permissions
from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

```
[2]: # fill in the path in your Google Drive in the string below. Note: do not
      ↳ escape slashes or spaces
import os
datadir = "./"
if not os.path.exists(datadir):
    !ln -s "/content/drive/My Drive/Your/Assignment3/path/" $datadir # TODO: Fill
    ↳ your Assignment 3 path
os.chdir(datadir)
!pwd
```

/content

```
[ ]: # In a terminal, Run the `download_data.sh` script in the data folder of
      ↳ assignment 3 part 2
```

```
[ ]: import os
import random

import cv2
import numpy as np

import torch
from torch.utils.data import DataLoader
from torchvision import models

from src.resnet_yolo import resnet50
from yolo_loss import YoloLoss
from src.dataset import VocDetectorDataset
from src.eval_voc import evaluate
from src.predict import predict_image
from src.config import VOC_CLASSES, COLORS
from kaggle_submission import output_submission_csv
```

```
import matplotlib.pyplot as plt
import collections

%matplotlib inline
%load_ext autoreload
%autoreload 2
```

0.1 Initialization

```
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda:0

```
[ ]: # YOLO network hyperparameters
B = 2 # number of bounding box predictions per cell
S = 14 # width/height of network output grid (larger than 7x7 from paper since
↪ we use a different network)
```

To implement Yolo we will rely on a pretrained classifier as the backbone for our detection network. PyTorch offers a variety of models which are pretrained on ImageNet in the `torchvision.models` package. In particular, we will use the ResNet50 architecture as a base for our detector. This is different from the base architecture in the Yolo paper and also results in a different output grid size (14x14 instead of 7x7).

Models are typically pretrained on ImageNet since the dataset is very large (> 1 million images) and widely used. The pretrained model provides a very useful weight initialization for our detector, so that the network is able to learn quickly and effectively.

```
[ ]: load_network_path = 'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet50().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet50(pretrained=pretrained).to(device)
```

Loading saved network from checkpoints/best_detector.pth

```
[ ]: learning_rate = 0.001
num_epochs = 50
batch_size = 48
```

```
# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5
```

0.2 Reading Pascal Data

Since Pascal is a small dataset (5000 in train+val) we have combined the train and val splits to train our detector. This is not typically a good practice, but we will make an exception in this case to be able to get reasonable detection results with a comparatively small object detection dataset.

The train dataset loader also using a variety of data augmentation techniques including random shift, scaling, crop, and flips. Data augmentation is slightly more complicated for detection datasets since the bounding box annotations must be kept consistent throughout the transformations.

Since the output of the detector network we train is an $S \times S \times (B \times 5 + C)$, we use an encoder to convert the original bounding box coordinates into relative grid bounding box coordinates corresponding to the expected output. We also use a decoder which allows us to convert the opposite direction into image coordinate bounding boxes.

```
[ ]: file_root_train = r"data/VOCdevkit_2007/VOC2007/JPEGImages/"
      annotation_file_train = "data/voc2007.txt"

      train_dataset = □
        ↳VocDetectorDataset(root_img_dir=file_root_train,dataset_file=annotation_file_train,train=True,
        ↳S=S)
      train_loader = □
        ↳DataLoader(train_dataset,batch_size=batch_size,shuffle=True,num_workers=2)
      print('Loaded %d train images' % len(train_dataset))
```

Initializing dataset
Loaded 5011 train images

```
[ ]: file_root_test = "data/VOCdevkit_2007/VOC2007test/JPEGImages/"
      annotation_file_test = "data/voc2007test.txt"

      test_dataset = □
        ↳VocDetectorDataset(root_img_dir=file_root_test,dataset_file=annotation_file_test,train=False,
        ↳S=S)
      test_loader = □
        ↳DataLoader(test_dataset,batch_size=batch_size,shuffle=False,num_workers=2)
      print('Loaded %d test images' % len(test_dataset))
```

Initializing dataset
Loaded 4950 test images

```
[ ]: data = train_dataset[0]
      print(data)
```

```

(tensor([[-116., -106., -100., ..., -108., -107., -110.],
        [-115., -103., -94., ..., -104., -105., -109.],
        [-111., -112., -93., ..., -104., -106., -110.],
        ...,
        [-85., -95., -93., ..., -102., -100., -97.],
        [-78., -88., -87., ..., -103., -103., -99.],
        [-74., -84., -83., ..., -102., -100., -97.]]],

[[[-102., -96., -98., ..., -97., -95., -98.],
  [-104., -93., -90., ..., -93., -93., -96.],
  [-103., -104., -83., ..., -92., -94., -97.],
  ...,
  [-74., -80., -74., ..., -80., -77., -74.],
  [-69., -77., -70., ..., -81., -79., -75.],
  [-65., -73., -66., ..., -80., -76., -73.]]],

[[[-86., -78., -75., ..., -85., -84., -87.],
  [-88., -76., -70., ..., -82., -82., -87.],
  [-91., -90., -71., ..., -84., -84., -90.],
  ...,
  [-67., -75., -69., ..., -77., -75., -72.],
  [-63., -71., -65., ..., -79., -78., -74.],
  [-59., -67., -61., ..., -78., -75., -72.]]]), tensor([[[0.0000,
0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000]],

[[[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000]]],

```

[illegible]


```

False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, True, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False],
[False, False, False, False, False, False, False, False, False, False,
False, False, False, False]]))

```

0.3 Set up training tools

```

[ ]: criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9,
↪weight_decay=5e-4)

```

0.4 Train detector

```

[ ]: best_test_loss = np.inf
learning_rate = 1e-3
for epoch in range(num_epochs):
    net.train()

    # Update learning rate late in training
    if epoch == 30 or epoch == 40:
        learning_rate /= 10.0

    for param_group in optimizer.param_groups:
        param_group['lr'] = learning_rate

    print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
    print('Learning Rate for this epoch: {}'.format(learning_rate))

total_loss = collections.defaultdict(int)

```

```

for i, data in enumerate(train_loader):
    data = (item.to(device) for item in data)
    images, target_boxes, target_cls, has_object_map = data
    pred = net(images)
    loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
    for key in loss_dict:
        total_loss[key] += loss_dict[key].item()

    optimizer.zero_grad()
    loss_dict['total_loss'].backward()
    optimizer.step()

    if (i+1) % 50 == 0:
        outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1,
↪ num_epochs, i+1, len(train_loader)))
        outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for
↪ key, val in total_loss.items() )
        print(outstring)

# evaluate the network on the test data
if (epoch + 1) % 5 == 0:
    test_aps = evaluate(net, test_dataset_file=annotation_file_test,
↪ img_root=file_root_test)
    print(epoch, test_aps)
    with torch.no_grad():
        test_loss = 0.0
        net.eval()
        for i, data in enumerate(test_loader):
            data = (item.to(device) for item in data)
            images, target_boxes, target_cls, has_object_map = data

            pred = net(images)
            loss_dict = criterion(pred, target_boxes, target_cls,
↪ has_object_map)
            test_loss += loss_dict['total_loss'].item()
        test_loss /= len(test_loader)

    if best_test_loss > test_loss:
        best_test_loss = test_loss
        print('Updating best test loss: %.5f' % best_test_loss)
        torch.save(net.state_dict(), 'checkpoints/best_detector.pth')

    if (epoch+1) in [5, 10, 20, 30, 40]:
        torch.save(net.state_dict(), 'checkpoints/detector_epoch_%d.pth' %
↪ (epoch+1))

    torch.save(net.state_dict(), 'checkpoints/detector.pth')

```

Starting epoch 1 / 50

Learning Rate for this epoch: 0.001

Epoch [1/50], Iter [50/209], Loss: total=23.023, reg=3.690,
containing_obj=0.332, no_obj=12.304, cls=6.696

Epoch [1/50], Iter [100/209], Loss: total=14.953, reg=3.090,
containing_obj=0.452, no_obj=6.371, cls=5.040

Epoch [1/50], Iter [150/209], Loss: total=12.040, reg=2.833,
containing_obj=0.525, no_obj=4.349, cls=4.334

Epoch [1/50], Iter [200/209], Loss: total=10.444, reg=2.674,
containing_obj=0.570, no_obj=3.322, cls=3.877

Updating best test loss: 5.28874

Starting epoch 2 / 50

Learning Rate for this epoch: 0.001

Epoch [2/50], Iter [50/209], Loss: total=5.168, reg=2.059, containing_obj=0.762,
no_obj=0.197, cls=2.150

Epoch [2/50], Iter [100/209], Loss: total=5.136, reg=2.074,
containing_obj=0.779, no_obj=0.184, cls=2.099

Epoch [2/50], Iter [150/209], Loss: total=5.020, reg=2.037,
containing_obj=0.778, no_obj=0.174, cls=2.031

Epoch [2/50], Iter [200/209], Loss: total=4.918, reg=2.004,
containing_obj=0.783, no_obj=0.165, cls=1.966

Updating best test loss: 4.64379

Starting epoch 3 / 50

Learning Rate for this epoch: 0.001

Epoch [3/50], Iter [50/209], Loss: total=4.577, reg=1.916, containing_obj=0.852,
no_obj=0.124, cls=1.685

Epoch [3/50], Iter [100/209], Loss: total=4.455, reg=1.874,
containing_obj=0.838, no_obj=0.119, cls=1.623

Epoch [3/50], Iter [150/209], Loss: total=4.459, reg=1.880,
containing_obj=0.846, no_obj=0.115, cls=1.618

Epoch [3/50], Iter [200/209], Loss: total=4.385, reg=1.850,
containing_obj=0.846, no_obj=0.112, cls=1.578

Updating best test loss: 4.35423

Starting epoch 4 / 50

Learning Rate for this epoch: 0.001

Epoch [4/50], Iter [50/209], Loss: total=4.112, reg=1.788, containing_obj=0.855,
no_obj=0.095, cls=1.374

Epoch [4/50], Iter [100/209], Loss: total=4.119, reg=1.788,
containing_obj=0.869, no_obj=0.093, cls=1.369

Epoch [4/50], Iter [150/209], Loss: total=4.085, reg=1.781,

```

containing_obj=0.865, no_obj=0.092, cls=1.347
Epoch [4/50], Iter [200/209], Loss: total=4.072, reg=1.781,
containing_obj=0.874, no_obj=0.090, cls=1.327
Updating best test loss: 4.03567

Starting epoch 5 / 50
Learning Rate for this epoch: 0.001
Epoch [5/50], Iter [50/209], Loss: total=3.915, reg=1.737, containing_obj=0.889,
no_obj=0.085, cls=1.203
Epoch [5/50], Iter [100/209], Loss: total=3.842, reg=1.718,
containing_obj=0.878, no_obj=0.086, cls=1.160
Epoch [5/50], Iter [150/209], Loss: total=3.859, reg=1.733,
containing_obj=0.880, no_obj=0.086, cls=1.160
Epoch [5/50], Iter [200/209], Loss: total=3.809, reg=1.724,
containing_obj=0.868, no_obj=0.088, cls=1.128
---Evaluate model on test samples---

100%|          | 4950/4950 [01:59<00:00, 41.33it/s]

---class aeroplane ap 0.07658872725054766---
---class bicycle ap 0.05062845689428884---
---class bird ap 0.06031496703720878---
---class boat ap 0.015634869227697974---
---class bottle ap 0.0161706925891787---
---class bus ap 0.0--- (no predictions for this class)
---class car ap 0.12292450106729097---
---class cat ap 0.0111731843575419---
---class chair ap 0.03533659231025792---
---class cow ap 0.0---
---class diningtable ap 0.0--- (no predictions for this class)
---class dog ap 0.004703476482617587---
---class horse ap 0.05198014629049112---
---class motorbike ap 0.0--- (no predictions for this class)
---class person ap 0.04971351804509193---
---class pottedplant ap 0.009684292097146533---
---class sheep ap 0.019559228650137744---
---class sofa ap 0.0--- (no predictions for this class)
---class train ap 0.0--- (no predictions for this class)
---class tvmonitor ap 0.16219807005625944---
---map 0.03433053611778785---
4 [0.07658872725054766, 0.05062845689428884, 0.06031496703720878,
0.015634869227697974, 0.0161706925891787, 0.0, 0.12292450106729097,
0.0111731843575419, 0.03533659231025792, 0.0, 0.0, 0.004703476482617587,
0.05198014629049112, 0.0, 0.04971351804509193, 0.009684292097146533,
0.019559228650137744, 0.0, 0.0, 0.16219807005625944]
Updating best test loss: 3.77643

```

Starting epoch 6 / 50

Learning Rate for this epoch: 0.001

Epoch [6/50], Iter [50/209], Loss: total=3.530, reg=1.647, containing_obj=0.830, no_obj=0.098, cls=0.954

Epoch [6/50], Iter [100/209], Loss: total=3.536, reg=1.650, containing_obj=0.829, no_obj=0.102, cls=0.955

Epoch [6/50], Iter [150/209], Loss: total=3.550, reg=1.661, containing_obj=0.826, no_obj=0.107, cls=0.956

Epoch [6/50], Iter [200/209], Loss: total=3.536, reg=1.648, containing_obj=0.824, no_obj=0.110, cls=0.953

Updating best test loss: 3.60945

Starting epoch 7 / 50

Learning Rate for this epoch: 0.001

Epoch [7/50], Iter [50/209], Loss: total=3.457, reg=1.609, containing_obj=0.837, no_obj=0.124, cls=0.886

Epoch [7/50], Iter [100/209], Loss: total=3.449, reg=1.628, containing_obj=0.829, no_obj=0.122, cls=0.871

Epoch [7/50], Iter [150/209], Loss: total=3.389, reg=1.597, containing_obj=0.817, no_obj=0.122, cls=0.853

Epoch [7/50], Iter [200/209], Loss: total=3.403, reg=1.600, containing_obj=0.817, no_obj=0.124, cls=0.862

Updating best test loss: 3.55916

Starting epoch 8 / 50

Learning Rate for this epoch: 0.001

Epoch [8/50], Iter [50/209], Loss: total=3.353, reg=1.586, containing_obj=0.839, no_obj=0.124, cls=0.803

Epoch [8/50], Iter [100/209], Loss: total=3.279, reg=1.549, containing_obj=0.822, no_obj=0.127, cls=0.781

Epoch [8/50], Iter [150/209], Loss: total=3.265, reg=1.553, containing_obj=0.816, no_obj=0.128, cls=0.768

Epoch [8/50], Iter [200/209], Loss: total=3.256, reg=1.551, containing_obj=0.809, no_obj=0.128, cls=0.768

Updating best test loss: 3.41244

Starting epoch 9 / 50

Learning Rate for this epoch: 0.001

Epoch [9/50], Iter [50/209], Loss: total=3.145, reg=1.535, containing_obj=0.804, no_obj=0.130, cls=0.676

Epoch [9/50], Iter [100/209], Loss: total=3.182, reg=1.542, containing_obj=0.810, no_obj=0.133, cls=0.697

Epoch [9/50], Iter [150/209], Loss: total=3.187, reg=1.536, containing_obj=0.815, no_obj=0.131, cls=0.704

Epoch [9/50], Iter [200/209], Loss: total=3.181, reg=1.537,

containing_obj=0.811, no_obj=0.131, cls=0.702
Updating best test loss: 3.35227

Starting epoch 10 / 50

Learning Rate for this epoch: 0.001

Epoch [10/50], Iter [50/209], Loss: total=3.063, reg=1.488,
containing_obj=0.799, no_obj=0.131, cls=0.645

Epoch [10/50], Iter [100/209], Loss: total=3.017, reg=1.466,
containing_obj=0.787, no_obj=0.133, cls=0.631

Epoch [10/50], Iter [150/209], Loss: total=3.039, reg=1.481,
containing_obj=0.793, no_obj=0.133, cls=0.632

Epoch [10/50], Iter [200/209], Loss: total=3.037, reg=1.478,
containing_obj=0.794, no_obj=0.134, cls=0.632

---Evaluate model on test samples---

100%| | 4950/4950 [01:55<00:00, 42.94it/s]

---class aeroplane ap 0.29672594414934905---

---class bicycle ap 0.3422849102545965---

---class bird ap 0.26962383409735075---

---class boat ap 0.11380292495688714---

---class bottle ap 0.0465665892740073---

---class bus ap 0.25234830916967904---

---class car ap 0.4968924330943932---

---class cat ap 0.5211356780331985---

---class chair ap 0.16107418955618402---

---class cow ap 0.14420768319136462---

---class diningtable ap 0.0048543689320388345---

---class dog ap 0.44054838096319154---

---class horse ap 0.44923080267119464---

---class motorbike ap 0.3760258183215436---

---class person ap 0.2922193643238318---

---class pottedplant ap 0.06539052004378372---

---class sheep ap 0.18705394618455345---

---class sofa ap 0.16438100439114636---

---class train ap 0.45297452404630134---

---class tvmonitor ap 0.3515317055710824---

---map 0.2714436465612839---

9 [0.29672594414934905, 0.3422849102545965, 0.26962383409735075,
0.11380292495688714, 0.0465665892740073, 0.25234830916967904,
0.4968924330943932, 0.5211356780331985, 0.16107418955618402,
0.14420768319136462, 0.0048543689320388345, 0.44054838096319154,
0.44923080267119464, 0.3760258183215436, 0.2922193643238318,
0.06539052004378372, 0.18705394618455345, 0.16438100439114636,
0.45297452404630134, 0.3515317055710824]

Updating best test loss: 3.32889

Starting epoch 11 / 50
Learning Rate for this epoch: 0.001
Epoch [11/50], Iter [50/209], Loss: total=2.953, reg=1.446,
containing_obj=0.790, no_obj=0.135, cls=0.582
Epoch [11/50], Iter [100/209], Loss: total=2.943, reg=1.435,
containing_obj=0.790, no_obj=0.139, cls=0.579
Epoch [11/50], Iter [150/209], Loss: total=2.946, reg=1.444,
containing_obj=0.788, no_obj=0.140, cls=0.574
Epoch [11/50], Iter [200/209], Loss: total=2.926, reg=1.438,
containing_obj=0.780, no_obj=0.141, cls=0.567
Updating best test loss: 3.20719

Starting epoch 12 / 50
Learning Rate for this epoch: 0.001
Epoch [12/50], Iter [50/209], Loss: total=2.938, reg=1.448,
containing_obj=0.788, no_obj=0.154, cls=0.547
Epoch [12/50], Iter [100/209], Loss: total=2.832, reg=1.394,
containing_obj=0.761, no_obj=0.156, cls=0.520
Epoch [12/50], Iter [150/209], Loss: total=2.823, reg=1.391,
containing_obj=0.758, no_obj=0.156, cls=0.518
Epoch [12/50], Iter [200/209], Loss: total=2.831, reg=1.394,
containing_obj=0.759, no_obj=0.155, cls=0.524
Updating best test loss: 3.12265

Starting epoch 13 / 50
Learning Rate for this epoch: 0.001
Epoch [13/50], Iter [50/209], Loss: total=2.705, reg=1.329,
containing_obj=0.739, no_obj=0.161, cls=0.475
Epoch [13/50], Iter [100/209], Loss: total=2.816, reg=1.391,
containing_obj=0.753, no_obj=0.165, cls=0.507
Epoch [13/50], Iter [150/209], Loss: total=2.810, reg=1.392,
containing_obj=0.748, no_obj=0.167, cls=0.503
Epoch [13/50], Iter [200/209], Loss: total=2.778, reg=1.379,
containing_obj=0.737, no_obj=0.166, cls=0.496

Starting epoch 14 / 50
Learning Rate for this epoch: 0.001
Epoch [14/50], Iter [50/209], Loss: total=2.766, reg=1.355,
containing_obj=0.744, no_obj=0.166, cls=0.501
Epoch [14/50], Iter [100/209], Loss: total=2.730, reg=1.347,
containing_obj=0.730, no_obj=0.172, cls=0.482
Epoch [14/50], Iter [150/209], Loss: total=2.732, reg=1.362,
containing_obj=0.730, no_obj=0.172, cls=0.468
Epoch [14/50], Iter [200/209], Loss: total=2.706, reg=1.340,
containing_obj=0.726, no_obj=0.171, cls=0.468

Updating best test loss: 3.06067

Starting epoch 15 / 50

Learning Rate for this epoch: 0.001

Epoch [15/50], Iter [50/209], Loss: total=2.670, reg=1.348,
containing_obj=0.734, no_obj=0.171, cls=0.417

Epoch [15/50], Iter [100/209], Loss: total=2.691, reg=1.348,
containing_obj=0.731, no_obj=0.172, cls=0.440

Epoch [15/50], Iter [150/209], Loss: total=2.654, reg=1.317,
containing_obj=0.726, no_obj=0.172, cls=0.440

Epoch [15/50], Iter [200/209], Loss: total=2.636, reg=1.305,
containing_obj=0.718, no_obj=0.172, cls=0.441

---Evaluate model on test samples---

100%| | 4950/4950 [01:49<00:00, 45.00it/s]

---class aeroplane ap 0.353683881278508---

---class bicycle ap 0.5256176577417708---

---class bird ap 0.3356674510308097---

---class boat ap 0.16124285849869038---

---class bottle ap 0.08810782975750252---

---class bus ap 0.4250292824366394---

---class car ap 0.5366674313982549---

---class cat ap 0.598304841212767---

---class chair ap 0.2232592629911557---

---class cow ap 0.2789533399214591---

---class diningtable ap 0.05270457697642164---

---class dog ap 0.5249338844330504---

---class horse ap 0.5374441023029601---

---class motorbike ap 0.42264480520659303---

---class person ap 0.40555980096429206---

---class pottedplant ap 0.11200195599409636---

---class sheep ap 0.29994128351451826---

---class sofa ap 0.31837120455312473---

---class train ap 0.60919846631995---

---class tvmonitor ap 0.35808844254523087---

---map 0.3583711179538897---

14 [0.353683881278508, 0.5256176577417708, 0.3356674510308097,
0.16124285849869038, 0.08810782975750252, 0.4250292824366394,
0.5366674313982549, 0.598304841212767, 0.2232592629911557, 0.2789533399214591,
0.05270457697642164, 0.5249338844330504, 0.5374441023029601,
0.42264480520659303, 0.40555980096429206, 0.11200195599409636,
0.29994128351451826, 0.31837120455312473, 0.60919846631995, 0.35808844254523087]

Updating best test loss: 3.04224

Starting epoch 16 / 50

Learning Rate for this epoch: 0.001

Epoch [16/50], Iter [50/209], Loss: total=2.486, reg=1.218,
containing_obj=0.691, no_obj=0.172, cls=0.405
Epoch [16/50], Iter [100/209], Loss: total=2.511, reg=1.239,
containing_obj=0.683, no_obj=0.177, cls=0.413
Epoch [16/50], Iter [150/209], Loss: total=2.532, reg=1.250,
containing_obj=0.694, no_obj=0.175, cls=0.413
Epoch [16/50], Iter [200/209], Loss: total=2.553, reg=1.266,
containing_obj=0.697, no_obj=0.174, cls=0.417
Updating best test loss: 3.01288

Starting epoch 17 / 50

Learning Rate for this epoch: 0.001

Epoch [17/50], Iter [50/209], Loss: total=2.566, reg=1.276,
containing_obj=0.713, no_obj=0.180, cls=0.397
Epoch [17/50], Iter [100/209], Loss: total=2.487, reg=1.223,
containing_obj=0.695, no_obj=0.182, cls=0.387
Epoch [17/50], Iter [150/209], Loss: total=2.475, reg=1.217,
containing_obj=0.691, no_obj=0.182, cls=0.384
Epoch [17/50], Iter [200/209], Loss: total=2.489, reg=1.221,
containing_obj=0.697, no_obj=0.180, cls=0.391
Updating best test loss: 2.97498

Starting epoch 18 / 50

Learning Rate for this epoch: 0.001

Epoch [18/50], Iter [50/209], Loss: total=2.517, reg=1.262,
containing_obj=0.710, no_obj=0.187, cls=0.358
Epoch [18/50], Iter [100/209], Loss: total=2.508, reg=1.253,
containing_obj=0.705, no_obj=0.183, cls=0.367
Epoch [18/50], Iter [150/209], Loss: total=2.484, reg=1.234,
containing_obj=0.695, no_obj=0.181, cls=0.374
Epoch [18/50], Iter [200/209], Loss: total=2.507, reg=1.244,
containing_obj=0.700, no_obj=0.181, cls=0.382
Updating best test loss: 2.89958

Starting epoch 19 / 50

Learning Rate for this epoch: 0.001

Epoch [19/50], Iter [50/209], Loss: total=2.396, reg=1.171,
containing_obj=0.688, no_obj=0.175, cls=0.362
Epoch [19/50], Iter [100/209], Loss: total=2.444, reg=1.215,
containing_obj=0.692, no_obj=0.177, cls=0.360
Epoch [19/50], Iter [150/209], Loss: total=2.421, reg=1.198,
containing_obj=0.682, no_obj=0.180, cls=0.361
Epoch [19/50], Iter [200/209], Loss: total=2.426, reg=1.201,
containing_obj=0.681, no_obj=0.182, cls=0.362

```

Starting epoch 20 / 50
Learning Rate for this epoch: 0.001
Epoch [20/50], Iter [50/209], Loss: total=2.345, reg=1.163,
containing_obj=0.684, no_obj=0.180, cls=0.318
Epoch [20/50], Iter [100/209], Loss: total=2.331, reg=1.150,
containing_obj=0.673, no_obj=0.182, cls=0.326
Epoch [20/50], Iter [150/209], Loss: total=2.346, reg=1.162,
containing_obj=0.679, no_obj=0.183, cls=0.322
Epoch [20/50], Iter [200/209], Loss: total=2.356, reg=1.165,
containing_obj=0.678, no_obj=0.184, cls=0.329
---Evaluate model on test samples---

100%|          | 4950/4950 [01:50<00:00, 44.98it/s]

---class aeroplane ap 0.44728433469512224---
---class bicycle ap 0.5209844377553562---
---class bird ap 0.4001112153420066---
---class boat ap 0.1782360244094975---
---class bottle ap 0.12431433578342119---
---class bus ap 0.49089998788611183---
---class car ap 0.5543484487470753---
---class cat ap 0.5593981607331366---
---class chair ap 0.22098080775915271---
---class cow ap 0.42783670584498845---
---class diningtable ap 0.17728054829978707---
---class dog ap 0.5597769708665804---
---class horse ap 0.6254283583663549---
---class motorbike ap 0.5664770219279924---
---class person ap 0.4564772852679984---
---class pottedplant ap 0.1360019135184708---
---class sheep ap 0.3841037883120979---
---class sofa ap 0.4167732522132447---
---class train ap 0.5855054093211073---
---class tvmonitor ap 0.3582455920919174---
---map 0.40952322995707097---
19 [0.44728433469512224, 0.5209844377553562, 0.4001112153420066,
0.1782360244094975, 0.12431433578342119, 0.49089998788611183,
0.5543484487470753, 0.5593981607331366, 0.22098080775915271,
0.42783670584498845, 0.17728054829978707, 0.5597769708665804,
0.6254283583663549, 0.5664770219279924, 0.4564772852679984, 0.1360019135184708,
0.3841037883120979, 0.4167732522132447, 0.5855054093211073, 0.3582455920919174]

Starting epoch 21 / 50
Learning Rate for this epoch: 0.001
Epoch [21/50], Iter [50/209], Loss: total=2.316, reg=1.158,
containing_obj=0.658, no_obj=0.179, cls=0.321
Epoch [21/50], Iter [100/209], Loss: total=2.303, reg=1.138,

```

containing_obj=0.668, no_obj=0.184, cls=0.312
Epoch [21/50], Iter [150/209], Loss: total=2.344, reg=1.164,
containing_obj=0.676, no_obj=0.187, cls=0.317
Epoch [21/50], Iter [200/209], Loss: total=2.339, reg=1.157,
containing_obj=0.673, no_obj=0.188, cls=0.321

Starting epoch 22 / 50

Learning Rate for this epoch: 0.001

Epoch [22/50], Iter [50/209], Loss: total=2.302, reg=1.154,
containing_obj=0.672, no_obj=0.181, cls=0.295
Epoch [22/50], Iter [100/209], Loss: total=2.287, reg=1.136,
containing_obj=0.671, no_obj=0.184, cls=0.295
Epoch [22/50], Iter [150/209], Loss: total=2.297, reg=1.145,
containing_obj=0.675, no_obj=0.184, cls=0.293
Epoch [22/50], Iter [200/209], Loss: total=2.313, reg=1.146,
containing_obj=0.677, no_obj=0.186, cls=0.304
Updating best test loss: 2.86011

Starting epoch 23 / 50

Learning Rate for this epoch: 0.001

Epoch [23/50], Iter [50/209], Loss: total=2.217, reg=1.090,
containing_obj=0.643, no_obj=0.185, cls=0.299
Epoch [23/50], Iter [100/209], Loss: total=2.229, reg=1.098,
containing_obj=0.640, no_obj=0.187, cls=0.304
Epoch [23/50], Iter [150/209], Loss: total=2.212, reg=1.084,
containing_obj=0.642, no_obj=0.188, cls=0.297
Epoch [23/50], Iter [200/209], Loss: total=2.229, reg=1.088,
containing_obj=0.648, no_obj=0.190, cls=0.302

Starting epoch 24 / 50

Learning Rate for this epoch: 0.001

Epoch [24/50], Iter [50/209], Loss: total=2.119, reg=1.043,
containing_obj=0.638, no_obj=0.194, cls=0.244
Epoch [24/50], Iter [100/209], Loss: total=2.162, reg=1.061,
containing_obj=0.639, no_obj=0.193, cls=0.268
Epoch [24/50], Iter [150/209], Loss: total=2.198, reg=1.079,
containing_obj=0.647, no_obj=0.193, cls=0.280
Epoch [24/50], Iter [200/209], Loss: total=2.201, reg=1.083,
containing_obj=0.647, no_obj=0.193, cls=0.278
Updating best test loss: 2.82955

Starting epoch 25 / 50

Learning Rate for this epoch: 0.001

Epoch [25/50], Iter [50/209], Loss: total=2.209, reg=1.079,

```

containing_obj=0.650, no_obj=0.195, cls=0.286
Epoch [25/50], Iter [100/209], Loss: total=2.174, reg=1.066,
containing_obj=0.640, no_obj=0.194, cls=0.274
Epoch [25/50], Iter [150/209], Loss: total=2.181, reg=1.071,
containing_obj=0.646, no_obj=0.193, cls=0.271
Epoch [25/50], Iter [200/209], Loss: total=2.164, reg=1.061,
containing_obj=0.641, no_obj=0.194, cls=0.268
---Evaluate model on test samples---

100%|          | 4950/4950 [02:02<00:00, 40.43it/s]

---class aeroplane ap 0.5149458020731885---
---class bicycle ap 0.5106475588148521---
---class bird ap 0.46117354577234015---
---class boat ap 0.219702083376801---
---class bottle ap 0.15645203304209096---
---class bus ap 0.5717432535826358---
---class car ap 0.6270526877149574---
---class cat ap 0.6405853965777057---
---class chair ap 0.2519842755268026---
---class cow ap 0.35256324374888814---
---class diningtable ap 0.21343934006471996---
---class dog ap 0.5578261301333471---
---class horse ap 0.6606943988185616---
---class motorbike ap 0.5226537682252406---
---class person ap 0.48934373545790183---
---class pottedplant ap 0.15633199963217598---
---class sheep ap 0.4387286403033152---
---class sofa ap 0.45115321023353133---
---class train ap 0.6454942618790018---
---class tvmonitor ap 0.43682795137218233---
---map 0.44396716581751205---
24 [0.5149458020731885, 0.5106475588148521, 0.46117354577234015,
0.219702083376801, 0.15645203304209096, 0.5717432535826358, 0.6270526877149574,
0.6405853965777057, 0.2519842755268026, 0.35256324374888814,
0.21343934006471996, 0.5578261301333471, 0.6606943988185616, 0.5226537682252406,
0.48934373545790183, 0.15633199963217598, 0.4387286403033152,
0.45115321023353133, 0.6454942618790018, 0.43682795137218233]

```

Starting epoch 26 / 50

Learning Rate for this epoch: 0.001

```

Epoch [26/50], Iter [50/209], Loss: total=2.261, reg=1.096,
containing_obj=0.684, no_obj=0.192, cls=0.289
Epoch [26/50], Iter [100/209], Loss: total=2.158, reg=1.043,
containing_obj=0.649, no_obj=0.196, cls=0.269
Epoch [26/50], Iter [150/209], Loss: total=2.130, reg=1.033,
containing_obj=0.637, no_obj=0.197, cls=0.263
Epoch [26/50], Iter [200/209], Loss: total=2.133, reg=1.036,

```

containing_obj=0.638, no_obj=0.194, cls=0.264

Starting epoch 27 / 50

Learning Rate for this epoch: 0.001

Epoch [27/50], Iter [50/209], Loss: total=2.077, reg=1.036,
containing_obj=0.606, no_obj=0.193, cls=0.243

Epoch [27/50], Iter [100/209], Loss: total=2.110, reg=1.041,
containing_obj=0.617, no_obj=0.192, cls=0.259

Epoch [27/50], Iter [150/209], Loss: total=2.112, reg=1.029,
containing_obj=0.623, no_obj=0.196, cls=0.265

Epoch [27/50], Iter [200/209], Loss: total=2.096, reg=1.021,
containing_obj=0.617, no_obj=0.197, cls=0.261

Updating best test loss: 2.81818

Starting epoch 28 / 50

Learning Rate for this epoch: 0.001

Epoch [28/50], Iter [50/209], Loss: total=2.155, reg=1.037,
containing_obj=0.642, no_obj=0.188, cls=0.288

Epoch [28/50], Iter [100/209], Loss: total=2.093, reg=1.023,
containing_obj=0.620, no_obj=0.192, cls=0.258

Epoch [28/50], Iter [150/209], Loss: total=2.103, reg=1.025,
containing_obj=0.622, no_obj=0.194, cls=0.262

Epoch [28/50], Iter [200/209], Loss: total=2.100, reg=1.029,
containing_obj=0.621, no_obj=0.195, cls=0.255

Updating best test loss: 2.80208

Starting epoch 29 / 50

Learning Rate for this epoch: 0.001

Epoch [29/50], Iter [50/209], Loss: total=2.043, reg=0.984,
containing_obj=0.614, no_obj=0.198, cls=0.247

Epoch [29/50], Iter [100/209], Loss: total=2.067, reg=1.005,
containing_obj=0.620, no_obj=0.199, cls=0.243

Epoch [29/50], Iter [150/209], Loss: total=2.059, reg=1.001,
containing_obj=0.620, no_obj=0.198, cls=0.240

Epoch [29/50], Iter [200/209], Loss: total=2.047, reg=0.996,
containing_obj=0.614, no_obj=0.198, cls=0.240

Updating best test loss: 2.79635

Starting epoch 30 / 50

Learning Rate for this epoch: 0.001

Epoch [30/50], Iter [50/209], Loss: total=2.055, reg=1.001,
containing_obj=0.617, no_obj=0.200, cls=0.237

Epoch [30/50], Iter [100/209], Loss: total=2.030, reg=0.992,
containing_obj=0.613, no_obj=0.196, cls=0.228

```

Epoch [30/50], Iter [150/209], Loss: total=2.047, reg=1.001,
containing_obj=0.617, no_obj=0.195, cls=0.234
Epoch [30/50], Iter [200/209], Loss: total=2.046, reg=0.998,
containing_obj=0.617, no_obj=0.196, cls=0.235
---Evaluate model on test samples---

100%|      | 4950/4950 [02:16<00:00, 36.25it/s]

---class aeroplane ap 0.4719819809800481---
---class bicycle ap 0.5780565318334638---
---class bird ap 0.4494709722559183---
---class boat ap 0.2924411458666887---
---class bottle ap 0.17512803769885577---
---class bus ap 0.562238838562554---
---class car ap 0.6349132255562765---
---class cat ap 0.6546596536208156---
---class chair ap 0.2358103752066274---
---class cow ap 0.34905603398848234---
---class diningtable ap 0.2802260169057064---
---class dog ap 0.5733626063495525---
---class horse ap 0.6389986491725825---
---class motorbike ap 0.5315041878862564---
---class person ap 0.49390194785564623---
---class pottedplant ap 0.15866587891490264---
---class sheep ap 0.4187775294345249---
---class sofa ap 0.42169926734926033---
---class train ap 0.63268485842851---
---class tvmonitor ap 0.47652801270692446---
---map 0.45150528752867986---
29 [0.4719819809800481, 0.5780565318334638, 0.4494709722559183,
0.2924411458666887, 0.17512803769885577, 0.562238838562554, 0.6349132255562765,
0.6546596536208156, 0.2358103752066274, 0.34905603398848234, 0.2802260169057064,
0.5733626063495525, 0.6389986491725825, 0.5315041878862564, 0.49390194785564623,
0.15866587891490264, 0.4187775294345249, 0.42169926734926033, 0.63268485842851,
0.47652801270692446]
Updating best test loss: 2.76395

Starting epoch 31 / 50
Learning Rate for this epoch: 0.0001
Epoch [31/50], Iter [50/209], Loss: total=2.007, reg=0.987,
containing_obj=0.607, no_obj=0.189, cls=0.225
Epoch [31/50], Iter [100/209], Loss: total=1.956, reg=0.964,
containing_obj=0.594, no_obj=0.189, cls=0.209
Epoch [31/50], Iter [150/209], Loss: total=1.953, reg=0.956,
containing_obj=0.601, no_obj=0.189, cls=0.206
Epoch [31/50], Iter [200/209], Loss: total=1.942, reg=0.950,
containing_obj=0.599, no_obj=0.192, cls=0.202
Updating best test loss: 2.70105

```

Starting epoch 32 / 50

Learning Rate for this epoch: 0.0001

Epoch [32/50], Iter [50/209], Loss: total=1.856, reg=0.879,
containing_obj=0.588, no_obj=0.203, cls=0.187

Epoch [32/50], Iter [100/209], Loss: total=1.893, reg=0.902,
containing_obj=0.595, no_obj=0.203, cls=0.194

Epoch [32/50], Iter [150/209], Loss: total=1.906, reg=0.912,
containing_obj=0.596, no_obj=0.202, cls=0.198

Epoch [32/50], Iter [200/209], Loss: total=1.892, reg=0.903,
containing_obj=0.593, no_obj=0.200, cls=0.195

Updating best test loss: 2.69919

Starting epoch 33 / 50

Learning Rate for this epoch: 0.0001

Epoch [33/50], Iter [50/209], Loss: total=1.886, reg=0.912,
containing_obj=0.577, no_obj=0.200, cls=0.196

Epoch [33/50], Iter [100/209], Loss: total=1.835, reg=0.886,
containing_obj=0.566, no_obj=0.202, cls=0.181

Epoch [33/50], Iter [150/209], Loss: total=1.847, reg=0.890,
containing_obj=0.572, no_obj=0.203, cls=0.182

Epoch [33/50], Iter [200/209], Loss: total=1.860, reg=0.894,
containing_obj=0.578, no_obj=0.201, cls=0.187

Updating best test loss: 2.68816

Starting epoch 34 / 50

Learning Rate for this epoch: 0.0001

Epoch [34/50], Iter [50/209], Loss: total=1.796, reg=0.854,
containing_obj=0.563, no_obj=0.202, cls=0.177

Epoch [34/50], Iter [100/209], Loss: total=1.819, reg=0.869,
containing_obj=0.564, no_obj=0.200, cls=0.186

Epoch [34/50], Iter [150/209], Loss: total=1.840, reg=0.881,
containing_obj=0.576, no_obj=0.202, cls=0.181

Epoch [34/50], Iter [200/209], Loss: total=1.856, reg=0.890,
containing_obj=0.582, no_obj=0.200, cls=0.183

Starting epoch 35 / 50

Learning Rate for this epoch: 0.0001

Epoch [35/50], Iter [50/209], Loss: total=1.840, reg=0.864,
containing_obj=0.585, no_obj=0.207, cls=0.183

Epoch [35/50], Iter [100/209], Loss: total=1.877, reg=0.904,
containing_obj=0.586, no_obj=0.205, cls=0.183

Epoch [35/50], Iter [150/209], Loss: total=1.850, reg=0.881,
containing_obj=0.583, no_obj=0.205, cls=0.181


```

Epoch [35/50], Iter [200/209], Loss: total=1.835, reg=0.874,
containing_obj=0.577, no_obj=0.205, cls=0.179
---Evaluate model on test samples---
100%|      | 4950/4950 [01:49<00:00, 45.06it/s]

---class aeroplane ap 0.5182232060568526---
---class bicycle ap 0.5928895331166626---
---class bird ap 0.457731500298588---
---class boat ap 0.303100847827348---
---class bottle ap 0.20191777763229413---
---class bus ap 0.61536422439128---
---class car ap 0.6348491083974395---
---class cat ap 0.6722168600532524---
---class chair ap 0.278135548608228---
---class cow ap 0.4256744403555651---
---class diningtable ap 0.3268612595417006---
---class dog ap 0.6281144088561152---
---class horse ap 0.6612629844143209---
---class motorbike ap 0.5394669509337908---
---class person ap 0.5128583117137296---
---class pottedplant ap 0.17264680245592376---
---class sheep ap 0.4384131127375066---
---class sofa ap 0.5061388734972234---
---class train ap 0.6908505299321648---
---class tvmonitor ap 0.4982419733310435---
---map 0.48374791270755146---
34 [0.5182232060568526, 0.5928895331166626, 0.457731500298588,
0.303100847827348, 0.20191777763229413, 0.61536422439128, 0.6348491083974395,
0.6722168600532524, 0.278135548608228, 0.4256744403555651, 0.3268612595417006,
0.6281144088561152, 0.6612629844143209, 0.5394669509337908, 0.5128583117137296,
0.17264680245592376, 0.4384131127375066, 0.5061388734972234, 0.6908505299321648,
0.4982419733310435]
Updating best test loss: 2.68101

```

```

Starting epoch 36 / 50
Learning Rate for this epoch: 0.0001
Epoch [36/50], Iter [50/209], Loss: total=1.842, reg=0.894,
containing_obj=0.566, no_obj=0.200, cls=0.181
Epoch [36/50], Iter [100/209], Loss: total=1.840, reg=0.895,
containing_obj=0.571, no_obj=0.200, cls=0.173
Epoch [36/50], Iter [150/209], Loss: total=1.863, reg=0.914,
containing_obj=0.576, no_obj=0.201, cls=0.172
Epoch [36/50], Iter [200/209], Loss: total=1.841, reg=0.897,
containing_obj=0.573, no_obj=0.201, cls=0.170

```

```

Starting epoch 37 / 50

```

Learning Rate for this epoch: 0.0001
Epoch [37/50], Iter [50/209], Loss: total=1.832, reg=0.890,
containing_obj=0.573, no_obj=0.199, cls=0.170
Epoch [37/50], Iter [100/209], Loss: total=1.839, reg=0.881,
containing_obj=0.579, no_obj=0.200, cls=0.179
Epoch [37/50], Iter [150/209], Loss: total=1.833, reg=0.873,
containing_obj=0.581, no_obj=0.201, cls=0.177
Epoch [37/50], Iter [200/209], Loss: total=1.836, reg=0.880,
containing_obj=0.576, no_obj=0.202, cls=0.178
Updating best test loss: 2.67865

Starting epoch 38 / 50
Learning Rate for this epoch: 0.0001
Epoch [38/50], Iter [50/209], Loss: total=1.809, reg=0.866,
containing_obj=0.564, no_obj=0.206, cls=0.172
Epoch [38/50], Iter [100/209], Loss: total=1.783, reg=0.849,
containing_obj=0.557, no_obj=0.203, cls=0.174
Epoch [38/50], Iter [150/209], Loss: total=1.811, reg=0.873,
containing_obj=0.562, no_obj=0.200, cls=0.176
Epoch [38/50], Iter [200/209], Loss: total=1.804, reg=0.873,
containing_obj=0.563, no_obj=0.200, cls=0.168

Starting epoch 39 / 50
Learning Rate for this epoch: 0.0001
Epoch [39/50], Iter [50/209], Loss: total=1.820, reg=0.875,
containing_obj=0.568, no_obj=0.197, cls=0.180
Epoch [39/50], Iter [100/209], Loss: total=1.804, reg=0.862,
containing_obj=0.569, no_obj=0.199, cls=0.173
Epoch [39/50], Iter [150/209], Loss: total=1.794, reg=0.854,
containing_obj=0.566, no_obj=0.202, cls=0.173
Epoch [39/50], Iter [200/209], Loss: total=1.807, reg=0.859,
containing_obj=0.569, no_obj=0.201, cls=0.178

Starting epoch 40 / 50
Learning Rate for this epoch: 0.0001
Epoch [40/50], Iter [50/209], Loss: total=1.796, reg=0.854,
containing_obj=0.565, no_obj=0.203, cls=0.174
Epoch [40/50], Iter [100/209], Loss: total=1.844, reg=0.885,
containing_obj=0.579, no_obj=0.203, cls=0.176
Epoch [40/50], Iter [150/209], Loss: total=1.809, reg=0.860,
containing_obj=0.571, no_obj=0.205, cls=0.173
Epoch [40/50], Iter [200/209], Loss: total=1.811, reg=0.862,
containing_obj=0.570, no_obj=0.206, cls=0.173
---Evaluate model on test samples---

100%| | 4950/4950 [01:48<00:00, 45.50it/s]

```
---class aeroplane ap 0.4963578587792927---
---class bicycle ap 0.5756793539323094---
---class bird ap 0.4759911861761871---
---class boat ap 0.34074504406841394---
---class bottle ap 0.2197668206159354---
---class bus ap 0.6013150829459295---
---class car ap 0.6693489716075912---
---class cat ap 0.6831885360371106---
---class chair ap 0.2823448817599529---
---class cow ap 0.477146367867021---
---class diningtable ap 0.34567166923093595---
---class dog ap 0.6629475806430277---
---class horse ap 0.6844687741865834---
---class motorbike ap 0.5525202917014421---
---class person ap 0.521253363990281---
---class pottedplant ap 0.18267193056813083---
---class sheep ap 0.4471823407575993---
---class sofa ap 0.5170200961599739---
---class train ap 0.705163039335304---
---class tvmonitor ap 0.489696876633244---
---map 0.4965240033498134---
39 [0.4963578587792927, 0.5756793539323094, 0.4759911861761871,
0.34074504406841394, 0.2197668206159354, 0.6013150829459295, 0.6693489716075912,
0.6831885360371106, 0.2823448817599529, 0.477146367867021, 0.34567166923093595,
0.6629475806430277, 0.6844687741865834, 0.5525202917014421, 0.521253363990281,
0.18267193056813083, 0.4471823407575993, 0.5170200961599739, 0.705163039335304,
0.489696876633244]
```

Starting epoch 41 / 50

Learning Rate for this epoch: 1e-05

Epoch [41/50], Iter [50/209], Loss: total=1.853, reg=0.896,
containing_obj=0.565, no_obj=0.210, cls=0.181

Epoch [41/50], Iter [100/209], Loss: total=1.835, reg=0.871,
containing_obj=0.578, no_obj=0.206, cls=0.181

Epoch [41/50], Iter [150/209], Loss: total=1.830, reg=0.877,
containing_obj=0.574, no_obj=0.206, cls=0.173

Epoch [41/50], Iter [200/209], Loss: total=1.823, reg=0.876,
containing_obj=0.567, no_obj=0.206, cls=0.174

Updating best test loss: 2.67833

Starting epoch 42 / 50

Learning Rate for this epoch: 1e-05

Epoch [42/50], Iter [50/209], Loss: total=1.843, reg=0.896,
containing_obj=0.572, no_obj=0.203, cls=0.172

Epoch [42/50], Iter [100/209], Loss: total=1.780, reg=0.853,
containing_obj=0.554, no_obj=0.207, cls=0.167
Epoch [42/50], Iter [150/209], Loss: total=1.810, reg=0.868,
containing_obj=0.563, no_obj=0.204, cls=0.175
Epoch [42/50], Iter [200/209], Loss: total=1.799, reg=0.863,
containing_obj=0.563, no_obj=0.203, cls=0.170
Updating best test loss: 2.67742

Starting epoch 43 / 50

Learning Rate for this epoch: 1e-05

Epoch [43/50], Iter [50/209], Loss: total=1.832, reg=0.884,
containing_obj=0.576, no_obj=0.205, cls=0.167
Epoch [43/50], Iter [100/209], Loss: total=1.790, reg=0.849,
containing_obj=0.561, no_obj=0.210, cls=0.170
Epoch [43/50], Iter [150/209], Loss: total=1.801, reg=0.861,
containing_obj=0.567, no_obj=0.205, cls=0.167
Epoch [43/50], Iter [200/209], Loss: total=1.778, reg=0.845,
containing_obj=0.562, no_obj=0.204, cls=0.167

Starting epoch 44 / 50

Learning Rate for this epoch: 1e-05

Epoch [44/50], Iter [50/209], Loss: total=1.777, reg=0.857,
containing_obj=0.559, no_obj=0.204, cls=0.157
Epoch [44/50], Iter [100/209], Loss: total=1.758, reg=0.833,
containing_obj=0.560, no_obj=0.205, cls=0.160
Epoch [44/50], Iter [150/209], Loss: total=1.743, reg=0.823,
containing_obj=0.553, no_obj=0.203, cls=0.163
Epoch [44/50], Iter [200/209], Loss: total=1.774, reg=0.841,
containing_obj=0.563, no_obj=0.202, cls=0.168

Starting epoch 45 / 50

Learning Rate for this epoch: 1e-05

Epoch [45/50], Iter [50/209], Loss: total=1.715, reg=0.808,
containing_obj=0.533, no_obj=0.211, cls=0.162
Epoch [45/50], Iter [100/209], Loss: total=1.729, reg=0.816,
containing_obj=0.541, no_obj=0.209, cls=0.162
Epoch [45/50], Iter [150/209], Loss: total=1.779, reg=0.836,
containing_obj=0.566, no_obj=0.205, cls=0.172
Epoch [45/50], Iter [200/209], Loss: total=1.777, reg=0.839,
containing_obj=0.563, no_obj=0.206, cls=0.170

---Evaluate model on test samples---

100%| | 4950/4950 [02:05<00:00, 39.43it/s]

---class aeroplane ap 0.512483127477515---

---class bicycle ap 0.5852869781267083---

```

---class bird ap 0.4803436960479395---
---class boat ap 0.32965797702728616---
---class bottle ap 0.230540525666382---
---class bus ap 0.6147311786261999---
---class car ap 0.6599836053944655---
---class cat ap 0.6758410380008759---
---class chair ap 0.2861696088756105---
---class cow ap 0.48210031349501703---
---class diningtable ap 0.352889673707156---
---class dog ap 0.6459057021764523---
---class horse ap 0.6800722319896104---
---class motorbike ap 0.5555304880513685---
---class person ap 0.5207299831311535---
---class pottedplant ap 0.1750900289035132---
---class sheep ap 0.45206182152757535---
---class sofa ap 0.505239520641471---
---class train ap 0.7003660216670626---
---class tvmonitor ap 0.5124488150235984---
---map 0.49787361677784797---
44 [0.512483127477515, 0.5852869781267083, 0.4803436960479395,
0.32965797702728616, 0.230540525666382, 0.6147311786261999, 0.6599836053944655,
0.6758410380008759, 0.2861696088756105, 0.48210031349501703, 0.352889673707156,
0.6459057021764523, 0.6800722319896104, 0.5555304880513685, 0.5207299831311535,
0.1750900289035132, 0.45206182152757535, 0.505239520641471, 0.7003660216670626,
0.5124488150235984]
Updating best test loss: 2.67660

```

```

Starting epoch 46 / 50
Learning Rate for this epoch: 1e-05
Epoch [46/50], Iter [50/209], Loss: total=1.811, reg=0.854,
containing_obj=0.567, no_obj=0.206, cls=0.184
Epoch [46/50], Iter [100/209], Loss: total=1.745, reg=0.826,
containing_obj=0.543, no_obj=0.205, cls=0.172
Epoch [46/50], Iter [150/209], Loss: total=1.768, reg=0.840,
containing_obj=0.553, no_obj=0.205, cls=0.170
Epoch [46/50], Iter [200/209], Loss: total=1.780, reg=0.848,
containing_obj=0.555, no_obj=0.205, cls=0.172

```

```

Starting epoch 47 / 50
Learning Rate for this epoch: 1e-05
Epoch [47/50], Iter [50/209], Loss: total=1.707, reg=0.805,
containing_obj=0.537, no_obj=0.204, cls=0.162
Epoch [47/50], Iter [100/209], Loss: total=1.759, reg=0.839,
containing_obj=0.548, no_obj=0.204, cls=0.168
Epoch [47/50], Iter [150/209], Loss: total=1.766, reg=0.840,
containing_obj=0.555, no_obj=0.204, cls=0.166

```

Epoch [47/50], Iter [200/209], Loss: total=1.777, reg=0.848,
containing_obj=0.560, no_obj=0.203, cls=0.167
Updating best test loss: 2.67653

Starting epoch 48 / 50

Learning Rate for this epoch: 1e-05

Epoch [48/50], Iter [50/209], Loss: total=1.817, reg=0.891,
containing_obj=0.555, no_obj=0.207, cls=0.164
Epoch [48/50], Iter [100/209], Loss: total=1.751, reg=0.842,
containing_obj=0.545, no_obj=0.205, cls=0.158
Epoch [48/50], Iter [150/209], Loss: total=1.755, reg=0.840,
containing_obj=0.551, no_obj=0.202, cls=0.162
Epoch [48/50], Iter [200/209], Loss: total=1.771, reg=0.847,
containing_obj=0.559, no_obj=0.204, cls=0.161

Starting epoch 49 / 50

Learning Rate for this epoch: 1e-05

Epoch [49/50], Iter [50/209], Loss: total=1.755, reg=0.832,
containing_obj=0.550, no_obj=0.209, cls=0.164
Epoch [49/50], Iter [100/209], Loss: total=1.805, reg=0.878,
containing_obj=0.558, no_obj=0.207, cls=0.163
Epoch [49/50], Iter [150/209], Loss: total=1.783, reg=0.852,
containing_obj=0.558, no_obj=0.207, cls=0.166
Epoch [49/50], Iter [200/209], Loss: total=1.796, reg=0.861,
containing_obj=0.558, no_obj=0.207, cls=0.169

Starting epoch 50 / 50

Learning Rate for this epoch: 1e-05

Epoch [50/50], Iter [50/209], Loss: total=1.731, reg=0.820,
containing_obj=0.549, no_obj=0.203, cls=0.159
Epoch [50/50], Iter [100/209], Loss: total=1.793, reg=0.858,
containing_obj=0.569, no_obj=0.204, cls=0.162
Epoch [50/50], Iter [150/209], Loss: total=1.781, reg=0.851,
containing_obj=0.565, no_obj=0.203, cls=0.162
Epoch [50/50], Iter [200/209], Loss: total=1.786, reg=0.853,
containing_obj=0.558, no_obj=0.204, cls=0.170

---Evaluate model on test samples---

100%| | 4950/4950 [01:49<00:00, 45.04it/s]

---class aeroplane ap 0.5173585424561593---
---class bicycle ap 0.583714545511857---
---class bird ap 0.470908732081237---
---class boat ap 0.3220532102661491---
---class bottle ap 0.22284560678560017---
---class bus ap 0.6070306394455354---

```

---class car ap 0.6581933465351568---
---class cat ap 0.6884812056830767---
---class chair ap 0.2801820451820764---
---class cow ap 0.4824978902261956---
---class diningtable ap 0.33360376561974125---
---class dog ap 0.659913921396504---
---class horse ap 0.693596605213306---
---class motorbike ap 0.5690485472680742---
---class person ap 0.5323960906993205---
---class pottedplant ap 0.18505575579428019---
---class sheep ap 0.44878501670270876---
---class sofa ap 0.5058467567143687---
---class train ap 0.6846409861475697---
---class tvmonitor ap 0.520872493525345---
---map 0.49835128516271326---
49 [0.5173585424561593, 0.583714545511857, 0.470908732081237,
0.3220532102661491, 0.22284560678560017, 0.6070306394455354, 0.6581933465351568,
0.6884812056830767, 0.2801820451820764, 0.4824978902261956, 0.33360376561974125,
0.659913921396504, 0.693596605213306, 0.5690485472680742, 0.5323960906993205,
0.18505575579428019, 0.44878501670270876, 0.5058467567143687,
0.6846409861475697, 0.520872493525345]

```

```

[ ]: load_network_path = 'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet50().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet50(pretrained=pretrained).to(device)

```

Loading saved network from checkpoints/best_detector.pth

```

[ ]: # train for another 20 epochs
best_test_loss = np.inf
learning_rate = 1e-7
num_epochs = 20
for epoch in range(num_epochs):
    net.train()

    # Update learning rate late in training
    # if epoch == 10:
    #     learning_rate /= 10.0

```

```

for param_group in optimizer.param_groups:
    param_group['lr'] = learning_rate

print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
print('Learning Rate for this epoch: {}'.format(learning_rate))

total_loss = collections.defaultdict(int)

for i, data in enumerate(train_loader):
    data = (item.to(device) for item in data)
    images, target_boxes, target_cls, has_object_map = data
    pred = net(images)
    loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
    for key in loss_dict:
        total_loss[key] += loss_dict[key].item()

    optimizer.zero_grad()
    loss_dict['total_loss'].backward()
    optimizer.step()

    if (i+1) % 50 == 0:
        outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1,
↪num_epochs, i+1, len(train_loader)))
        outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for
↪key, val in total_loss.items() )
        print(outstring)

# evaluate the network on the test data
if (epoch + 1) % 5 == 0:
    test_aps = evaluate(net, test_dataset_file=annotation_file_test,
↪img_root=file_root_test)
    print(epoch, test_aps)
    with torch.no_grad():
        test_loss = 0.0
        net.eval()
        for i, data in enumerate(test_loader):
            data = (item.to(device) for item in data)
            images, target_boxes, target_cls, has_object_map = data

            pred = net(images)
            loss_dict = criterion(pred, target_boxes, target_cls,
↪has_object_map)
            test_loss += loss_dict['total_loss'].item()
        test_loss /= len(test_loader)

    if best_test_loss > test_loss:
        best_test_loss = test_loss

```



```

print('Updating best test loss: %.5f' % best_test_loss)
# torch.save(net.state_dict(), 'checkpoints/best_detector.pth')

if (epoch+1) in [5, 10]:
    torch.save(net.state_dict(), 'checkpoints/detector_epoch_%d.pth' %
↪(epoch+1))

torch.save(net.state_dict(), 'checkpoints/detector.pth')

```

Starting epoch 1 / 20

Learning Rate for this epoch: 1e-07

Epoch [1/20], Iter [50/105], Loss: total=1.752, reg=0.847, containing_obj=0.559,
no_obj=0.202, cls=0.144

Epoch [1/20], Iter [100/105], Loss: total=1.760, reg=0.856,
containing_obj=0.555, no_obj=0.201, cls=0.148

Updating best test loss: 2.67042

Starting epoch 2 / 20

Learning Rate for this epoch: 1e-07

Epoch [2/20], Iter [50/105], Loss: total=1.727, reg=0.813, containing_obj=0.552,
no_obj=0.205, cls=0.157

Epoch [2/20], Iter [100/105], Loss: total=1.738, reg=0.829,
containing_obj=0.554, no_obj=0.205, cls=0.150

Updating best test loss: 2.66498

Starting epoch 3 / 20

Learning Rate for this epoch: 1e-07

Epoch [3/20], Iter [50/105], Loss: total=1.729, reg=0.835, containing_obj=0.545,
no_obj=0.202, cls=0.148

Epoch [3/20], Iter [100/105], Loss: total=1.777, reg=0.857,
containing_obj=0.559, no_obj=0.202, cls=0.159

Starting epoch 4 / 20

Learning Rate for this epoch: 1e-07

Epoch [4/20], Iter [50/105], Loss: total=1.817, reg=0.869, containing_obj=0.584,
no_obj=0.201, cls=0.162

Epoch [4/20], Iter [100/105], Loss: total=1.756, reg=0.843,
containing_obj=0.557, no_obj=0.202, cls=0.153

Updating best test loss: 2.66042

Starting epoch 5 / 20

```

Learning Rate for this epoch: 1e-07
Epoch [5/20], Iter [50/105], Loss: total=1.751, reg=0.833, containing_obj=0.564,
no_obj=0.203, cls=0.149
Epoch [5/20], Iter [100/105], Loss: total=1.756, reg=0.840,
containing_obj=0.562, no_obj=0.203, cls=0.152
---Evaluate model on test samples---

100%|          | 4950/4950 [01:45<00:00, 46.90it/s]

---class aeroplane ap 0.5156360272465859---
---class bicycle ap 0.5821791933248942---
---class bird ap 0.4754608255767426---
---class boat ap 0.32956161643160997---
---class bottle ap 0.23962633438956316---
---class bus ap 0.6098380274170774---
---class car ap 0.6648527938897819---
---class cat ap 0.6839515576802225---
---class chair ap 0.2838002041505672---
---class cow ap 0.47104410313192613---
---class diningtable ap 0.3477910782984703---
---class dog ap 0.6620648664117479---
---class horse ap 0.6918500943034918---
---class motorbike ap 0.5617132807435516---
---class person ap 0.5268194852696351---
---class pottedplant ap 0.17710022135345396---
---class sheep ap 0.46297048589591366---
---class sofa ap 0.5044173320168852---
---class train ap 0.7046449198400089---
---class tvmonitor ap 0.507086273890087---
---map 0.5001204360631107---
4 [0.5156360272465859, 0.5821791933248942, 0.4754608255767426,
0.32956161643160997, 0.23962633438956316, 0.6098380274170774,
0.6648527938897819, 0.6839515576802225, 0.2838002041505672, 0.47104410313192613,
0.3477910782984703, 0.6620648664117479, 0.6918500943034918, 0.5617132807435516,
0.5268194852696351, 0.17710022135345396, 0.46297048589591366,
0.5044173320168852, 0.7046449198400089, 0.507086273890087]

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[16], line 45
      43 net.eval()
      44 for i, data in enumerate(test_loader):
----> 45     data = (item.to(device) for item in data)
      46     images, target_boxes, target_cls, has_object_map = data
      48     pred = net(images)

KeyboardInterrupt:

```

1 View example predictions

```
[ ]: net.eval()

# select random image from test set
image_name = random.choice(test_dataset.fnames)
image = cv2.imread(os.path.join(file_root_test, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_test)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.4, 1)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
    text_size[0], p1[1] + text_size[1]),
                color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.
    FONT_HERSHEY_SIMPLEX, 0.4, (255, 255, 255), 1, 8)

plt.figure(figsize = (15,15))
plt.imshow(image)
```

predicting...

```
[ ]: <matplotlib.image.AxesImage at 0x14dc30942f50>
```



1.1 Evaluate on Test

To evaluate detection results we use mAP (mean of average precision over each class)

```
[ ]: test_aps = evaluate(net, test_dataset_file=annotation_file_test,
    ↪img_root=file_root_test)
```

---Evaluate model on test samples---

100%| | 4950/4950 [01:44<00:00, 47.40it/s]

---class aeroplane ap 0.5156360272465859---

---class bicycle ap 0.5821791933248942---

---class bird ap 0.4754608255767426---

---class boat ap 0.32956161643160997---

---class bottle ap 0.23962633438956316---

---class bus ap 0.6098380274170774---

---class car ap 0.6648527938897819---

---class cat ap 0.6839515576802225---

---class chair ap 0.2838002041505672---

---class cow ap 0.47104410313192613---

---class diningtable ap 0.3477910782984703---

```

---class dog ap 0.6620648664117479---
---class horse ap 0.6918500943034918---
---class motorbike ap 0.5617132807435516---
---class person ap 0.5268194852696351---
---class pottedplant ap 0.17710022135345396---
---class sheep ap 0.46297048589591366---
---class sofa ap 0.5044173320168852---
---class train ap 0.7046449198400089---
---class tvmonitor ap 0.507086273890087---
---map 0.5001204360631107---

```

1.1.1 Cell added to get intermediate mAP values for students

```

[ ]: network_paths = ['./checkpoints/detector_epoch_%d.pth' % epoch for epoch in [5,
↳10, 20, 30, 40]]+ ['./checkpoints/detector.pth']
for load_network_path in network_paths:
    print('Loading saved network from {}'.format(load_network_path))
    net_loaded = resnet50().to(device)
    net_loaded.load_state_dict(torch.load(load_network_path))
    evaluate(net_loaded, test_dataset_file=annotation_file_test,
↳img_root=file_root_test)

```

Loading saved network from ./checkpoints/detector_epoch_5.pth

---Evaluate model on test samples---

100%| | 4950/4950 [01:33<00:00, 53.02it/s]

```

---class aeroplane ap 0.07658872725054766---
---class bicycle ap 0.05062845689428884---
---class bird ap 0.06031496703720878---
---class boat ap 0.015634869227697974---
---class bottle ap 0.0161706925891787---
---class bus ap 0.0--- (no predictions for this class)
---class car ap 0.12292450106729097---
---class cat ap 0.0111731843575419---
---class chair ap 0.03533659231025792---
---class cow ap 0.0---
---class diningtable ap 0.0--- (no predictions for this class)
---class dog ap 0.004703476482617587---
---class horse ap 0.05198014629049112---
---class motorbike ap 0.0--- (no predictions for this class)
---class person ap 0.04971351804509193---
---class pottedplant ap 0.009684292097146533---
---class sheep ap 0.019559228650137744---
---class sofa ap 0.0--- (no predictions for this class)
---class train ap 0.0--- (no predictions for this class)
---class tvmonitor ap 0.16219807005625944---
---map 0.03433053611778785---

```

Loading saved network from ./checkpoints/detector_epoch_10.pth

---Evaluate model on test samples---

100%| | 4950/4950 [01:41<00:00, 48.83it/s]

---class aeroplane ap 0.29672594414934905---
---class bicycle ap 0.3422849102545965---
---class bird ap 0.26962383409735075---
---class boat ap 0.11380292495688714---
---class bottle ap 0.0465665892740073---
---class bus ap 0.25234830916967904---
---class car ap 0.4968924330943932---
---class cat ap 0.5211356780331985---
---class chair ap 0.16107418955618402---
---class cow ap 0.14420768319136462---
---class diningtable ap 0.0048543689320388345---
---class dog ap 0.44054838096319154---
---class horse ap 0.44923080267119464---
---class motorbike ap 0.3760258183215436---
---class person ap 0.2922193643238318---
---class pottedplant ap 0.06539052004378372---
---class sheep ap 0.18705394618455345---
---class sofa ap 0.16438100439114636---
---class train ap 0.45297452404630134---
---class tvmonitor ap 0.3515317055710824---
---map 0.2714436465612839---

Loading saved network from ./checkpoints/detector_epoch_20.pth

---Evaluate model on test samples---

100%| | 4950/4950 [01:51<00:00, 44.27it/s]

---class aeroplane ap 0.44728433469512224---
---class bicycle ap 0.5209844377553562---
---class bird ap 0.4001112153420066---
---class boat ap 0.1782360244094975---
---class bottle ap 0.12431433578342119---
---class bus ap 0.49089998788611183---
---class car ap 0.5543484487470753---
---class cat ap 0.5593981607331366---
---class chair ap 0.22098080775915271---
---class cow ap 0.42783670584498845---
---class diningtable ap 0.17728054829978707---
---class dog ap 0.5597769708665804---
---class horse ap 0.6254283583663549---
---class motorbike ap 0.5664770219279924---
---class person ap 0.4564772852679984---
---class pottedplant ap 0.1360019135184708---
---class sheep ap 0.3841037883120979---
---class sofa ap 0.4167732522132447---
---class train ap 0.5855054093211073---
---class tvmonitor ap 0.3582455920919174---

```

---map 0.40952322995707097---
Loading saved network from ./checkpoints/detector_epoch_30.pth
---Evaluate model on test samples---

100%|      | 4950/4950 [01:51<00:00, 44.30it/s]

---class aeroplane ap 0.4719819809800481---
---class bicycle ap 0.5780565318334638---
---class bird ap 0.4494709722559183---
---class boat ap 0.2924411458666887---
---class bottle ap 0.17512803769885577---
---class bus ap 0.562238838562554---
---class car ap 0.6349132255562765---
---class cat ap 0.6546596536208156---
---class chair ap 0.2358103752066274---
---class cow ap 0.34905603398848234---
---class diningtable ap 0.2802260169057064---
---class dog ap 0.5733626063495525---
---class horse ap 0.6389986491725825---
---class motorbike ap 0.5315041878862564---
---class person ap 0.49390194785564623---
---class pottedplant ap 0.15866587891490264---
---class sheep ap 0.4187775294345249---
---class sofa ap 0.42169926734926033---
---class train ap 0.63268485842851---
---class tvmonitor ap 0.47652801270692446---
---map 0.45150528752867986---
Loading saved network from ./checkpoints/detector_epoch_40.pth
---Evaluate model on test samples---

100%|      | 4950/4950 [01:48<00:00, 45.72it/s]

---class aeroplane ap 0.4963578587792927---
---class bicycle ap 0.5756793539323094---
---class bird ap 0.4759911861761871---
---class boat ap 0.34074504406841394---
---class bottle ap 0.2197668206159354---
---class bus ap 0.6013150829459295---
---class car ap 0.6693489716075912---
---class cat ap 0.6831885360371106---
---class chair ap 0.2823448817599529---
---class cow ap 0.477146367867021---
---class diningtable ap 0.34567166923093595---
---class dog ap 0.6629475806430277---
---class horse ap 0.6844687741865834---
---class motorbike ap 0.5525202917014421---
---class person ap 0.521253363990281---
---class pottedplant ap 0.18267193056813083---
---class sheep ap 0.4471823407575993---
---class sofa ap 0.5170200961599739---

```

```

---class train ap 0.705163039335304---
---class tvmonitor ap 0.489696876633244---
---map 0.4965240033498134---
Loading saved network from ./checkpoints/detector.pth
---Evaluate model on test samples---

100%|      | 4950/4950 [02:00<00:00, 40.91it/s]

---class aeroplane ap 0.5173585424561593---
---class bicycle ap 0.583714545511857---
---class bird ap 0.470908732081237---
---class boat ap 0.3220532102661491---
---class bottle ap 0.22284560678560017---
---class bus ap 0.6070306394455354---
---class car ap 0.6581933465351568---
---class cat ap 0.6884812056830767---
---class chair ap 0.2801820451820764---
---class cow ap 0.4824978902261956---
---class diningtable ap 0.33360376561974125---
---class dog ap 0.659913921396504---
---class horse ap 0.693596605213306---
---class motorbike ap 0.5690485472680742---
---class person ap 0.5323960906993205---
---class pottedplant ap 0.18505575579428019---
---class sheep ap 0.44878501670270876---
---class sofa ap 0.5058467567143687---
---class train ap 0.6846409861475697---
---class tvmonitor ap 0.520872493525345---
---map 0.49835128516271326---

```

```
[ ]: output_submission_csv('my_new_solution.csv', test_aps)
```

2 Extra Credit 1: Video Object Detection

```
[ ]: import numpy as np
import cv2
import pafy
import matplotlib.pyplot as plt
from matplotlib import cm
from PIL import Image
from tqdm import trange

import torch
from torch import nn
from torchvision import transforms

url = "https://www.youtube.com/watch?v=xZGahvrep3o"
```



```
[ ]: !pip install pytube
```

Collecting pytube

Downloading pytube-15.0.0-py3-none-any.whl.metadata (5.0 kB)

Downloading pytube-15.0.0-py3-none-any.whl (57 kB)

57.6/57.6 kB

297.8 kB/s eta 0:00:00:00:01

Installing collected packages: pytube

Successfully installed pytube-15.0.0

```
[ ]: ### Download YouTube Video ###
from pytube import YouTube
url = "https://www.youtube.com/watch?v=xZGahvrep3o"
path = "./video"
video = YouTube(url)
video.streams.filter(progressive=True)
stream = video.streams.get_by_itag(22)
stream.download(path)
```

```
[ ]: '/u/qilong/mistral/mistral2/./video/SNL Digital Short YOLO - SNL.mp4'
```

```
[ ]: ### Video to Frame ###

def video2frame(videos_path, frames_save_path, time_interval):

    '''
    :param videos_path: path to video
    :param frames_save_path: path to frames
    :param time_interval: time interval
    :return:
    '''

    vidcap = cv2.VideoCapture(videos_path)
    success, image = vidcap.read()
    count = 0
    while success:
        success, image = vidcap.read()
        count += 1
        if count % time_interval == 0:
            try:
                cv2.imencode('.jpg', image)[1].tofile(frames_save_path + "%d.
↪jpg" % count)
            except:
                break
        print(count)

    return count
```

```
[ ]: videos_path = r'./video/SNL_Digital_Short_YOLO_SNL.mp4'
frames_save_path = r'./video/frame2'
time_interval = 2
count = video2frame(videos_path, frames_save_path, time_interval)
```

5265

```
[ ]: ### Load Detector ###
load_network_path = "checkpoints/best_detector.pth"
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet50().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet50(pretrained=pretrained).to(device)
```

Loading saved network from checkpoints/best_detector.pth

```
[ ]: ### Test on one frame ###
net.eval()
detected_save_path = "./video/detected2"

# load image from frame
idx = 100
image = cv2.imread(frames_save_path + "{}/.jpg".format(idx * 10))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

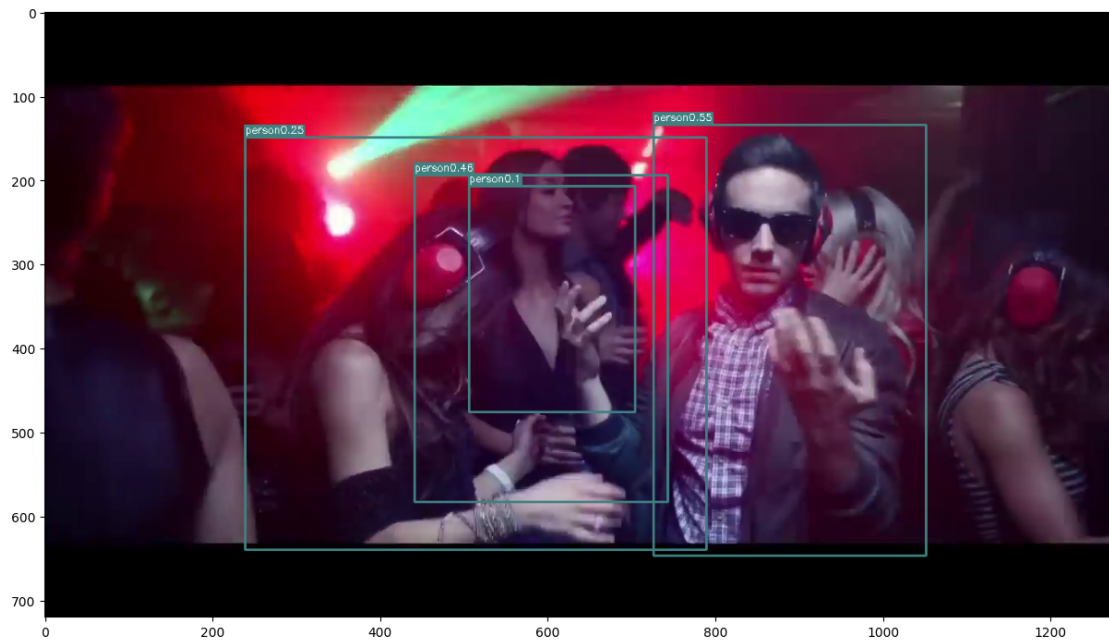
result = predict_image(net, "{}/.jpg".format(idx * 10),
    ↪root_img_directory=frames_save_path + "/"
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.4,
    ↪1)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
    ↪text_size[0], p1[1] + text_size[1]),
        color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.
    ↪FONT_HERSHEY_SIMPLEX, 0.4, (255, 255, 255), 1, 8)

# show image
```

```
plt.figure(figsize = (15,15))
plt.imshow(image)

# save image
cv2.imwrite(detected_save_path + "{:}.jpg".format(idx * 10), image)
```

[]: True



```
[ ]: ### Detection ###
net.eval()
detected_save_path = "./video/detected2"
if count is None:
    count = 5265

# load image from frame
for i in trange(count // time_interval):
    idx = i + 1
    try:
        image = cv2.imread("./video/frame2/{:}.jpg".format(idx * time_interval))
    except:
        print("Can not load frame {}".format(idx * time_interval))
        continue
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    result = predict_image(net, "{:}.jpg".format(idx * time_interval),
        root_img_directory="./video/frame2/")
```

```

    for left_up, right_bottom, class_name, _, prob in result:
        color = COLORS[VOC_CLASSES.index(class_name)]
        cv2.rectangle(image, left_up, right_bottom, color, 2)
        label = class_name + str(round(prob, 2))
        text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX,
↪0.4, 1)
        p1 = (left_up[0], left_up[1] - text_size[1])
        cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
↪text_size[0], p1[1] + text_size[1]), color, -1)
        cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.
↪FONT_HERSHEY_SIMPLEX, 0.4, (255, 255, 255), 1, 8)

    # save image
    cv2.imwrite(detected_save_path + "{}/{}.jpg".format(idx * time_interval),
↪image)

```

100%| | 2632/2632 [02:17<00:00, 19.16it/s]

[]: *### Frame to Video ###*

```

def frame2video(im_dir, video_dir, fps):

    im_list = []
    for idx in range(1, count // time_interval + 1):
        im_list.append("{}{}.jpg".format(idx * time_interval))
    img = Image.open(im_dir + im_list[100])
    img_size = img.size
    print(img_size)

    fourcc = cv2.VideoWriter_fourcc(*'mp4v')
    videoWriter = cv2.VideoWriter(video_dir, fourcc, fps, img_size)

    print("Start producing video...")
    for i in range(len(im_list)):
        im_name = im_dir + im_list[i]
        frame = cv2.imread(im_name)
        videoWriter.write(frame)
    videoWriter.release()
    print('finish')

```

```

[ ]: im_dir = detected_save_path + "/"
    video_dir = r'./video/detected2_video.mp4'
    fps = 24
    frame2video(im_dir, video_dir, fps)

```

(1280, 720)

Start producing video...

```
100%|          | 2632/2632 [00:32<00:00, 81.15it/s]
finish
```

3 Extra Credit 2: Better Pretrained Model

```
[ ]: from src.resnet_yolo import resnet101

load_network_path = None #'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet101().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet101(pretrained=pretrained).to(device)
```

Load pre-trained model

```
[ ]: learning_rate = 0.001
num_epochs = 50
batch_size = 24

# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5

criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9,
    ↪weight_decay=5e-4)
```

```
[ ]: best_test_loss = np.inf
for epoch in range(num_epochs):
    net.train()

    # Update learning rate late in training
    if epoch == 30 or epoch == 40:
        learning_rate /= 10.0

    for param_group in optimizer.param_groups:
        param_group['lr'] = learning_rate

    print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
```

```

print('Learning Rate for this epoch: {}'.format(learning_rate))

total_loss = collections.defaultdict(int)

for i, data in enumerate(train_loader):
    data = (item.to(device) for item in data)
    images, target_boxes, target_cls, has_object_map = data
    pred = net(images)
    loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
    for key in loss_dict:
        total_loss[key] += loss_dict[key].item()

    optimizer.zero_grad()
    loss_dict['total_loss'].backward()
    optimizer.step()

    if (i+1) % 50 == 0:
        outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1,
↪ num_epochs, i+1, len(train_loader)))
        outstring += ', '.join( "%s=%.3f" % (key[:-5], val / (i+1)) for
↪ key, val in total_loss.items() )
        print(outstring)

# evaluate the network on the test data
if (epoch + 1) % 5 == 0:
    test_aps = evaluate(net, test_dataset_file=annotation_file_test,
↪ img_root=file_root_test)
    print(epoch, test_aps)
    with torch.no_grad():
        test_loss = 0.0
        net.eval()
        for i, data in enumerate(test_loader):
            data = (item.to(device) for item in data)
            images, target_boxes, target_cls, has_object_map = data

            pred = net(images)
            loss_dict = criterion(pred, target_boxes, target_cls,
↪ has_object_map)
            test_loss += loss_dict['total_loss'].item()
        test_loss /= len(test_loader)

    if best_test_loss > test_loss:
        best_test_loss = test_loss
        print('Updating best test loss: %.5f' % best_test_loss)
        torch.save(net.state_dict(), 'checkpoints/res101_best_detector.pth')

if (epoch+1) in [5, 10, 20, 30, 40]:

```

```
torch.save(net.state_dict(), 'checkpoints/res101_detector_epoch_%d.pth' %  
↪ (epoch+1))  
  
torch.save(net.state_dict(), 'checkpoints/res101_detector.pth')
```

Starting epoch 1 / 50

Learning Rate for this epoch: 0.001

Epoch [1/50], Iter [50/209], Loss: total=23.906, reg=3.820,
containing_obj=0.347, no_obj=12.476, cls=7.263

Epoch [1/50], Iter [100/209], Loss: total=15.517, reg=3.252,
containing_obj=0.453, no_obj=6.467, cls=5.345

Epoch [1/50], Iter [150/209], Loss: total=12.311, reg=2.919,
containing_obj=0.516, no_obj=4.418, cls=4.458

Epoch [1/50], Iter [200/209], Loss: total=10.481, reg=2.697,
containing_obj=0.555, no_obj=3.377, cls=3.852

Updating best test loss: 5.02656

Starting epoch 2 / 50

Learning Rate for this epoch: 0.001

Epoch [2/50], Iter [50/209], Loss: total=4.768, reg=1.978, containing_obj=0.740,
no_obj=0.204, cls=1.846

Epoch [2/50], Iter [100/209], Loss: total=4.751, reg=1.962,
containing_obj=0.764, no_obj=0.190, cls=1.835

Epoch [2/50], Iter [150/209], Loss: total=4.743, reg=1.963,
containing_obj=0.778, no_obj=0.179, cls=1.824

Epoch [2/50], Iter [200/209], Loss: total=4.652, reg=1.935,
containing_obj=0.786, no_obj=0.169, cls=1.761

Updating best test loss: 4.39795

Starting epoch 3 / 50

Learning Rate for this epoch: 0.001

Epoch [3/50], Iter [50/209], Loss: total=4.397, reg=1.890, containing_obj=0.864,
no_obj=0.126, cls=1.517

Epoch [3/50], Iter [100/209], Loss: total=4.289, reg=1.848,
containing_obj=0.862, no_obj=0.121, cls=1.457

Epoch [3/50], Iter [150/209], Loss: total=4.225, reg=1.826,
containing_obj=0.870, no_obj=0.117, cls=1.413

Epoch [3/50], Iter [200/209], Loss: total=4.187, reg=1.821,
containing_obj=0.872, no_obj=0.113, cls=1.381

Updating best test loss: 4.08262

Starting epoch 4 / 50

Learning Rate for this epoch: 0.001
Epoch [4/50], Iter [50/209], Loss: total=3.852, reg=1.713, containing_obj=0.898,
no_obj=0.094, cls=1.146
Epoch [4/50], Iter [100/209], Loss: total=3.823, reg=1.693,
containing_obj=0.903, no_obj=0.092, cls=1.135
Epoch [4/50], Iter [150/209], Loss: total=3.813, reg=1.702,
containing_obj=0.908, no_obj=0.089, cls=1.114
Epoch [4/50], Iter [200/209], Loss: total=3.791, reg=1.705,
containing_obj=0.912, no_obj=0.087, cls=1.087
Updating best test loss: 3.77741

Starting epoch 5 / 50

Learning Rate for this epoch: 0.001
Epoch [5/50], Iter [50/209], Loss: total=3.573, reg=1.631, containing_obj=0.915,
no_obj=0.077, cls=0.951
Epoch [5/50], Iter [100/209], Loss: total=3.568, reg=1.636,
containing_obj=0.925, no_obj=0.075, cls=0.931
Epoch [5/50], Iter [150/209], Loss: total=3.542, reg=1.627,
containing_obj=0.928, no_obj=0.074, cls=0.913
Epoch [5/50], Iter [200/209], Loss: total=3.558, reg=1.639,
containing_obj=0.937, no_obj=0.073, cls=0.909
---Evaluate model on test samples---

100%| | 4950/4950 [02:15<00:00, 36.62it/s]

---class aeroplane ap 0.0---
---class bicycle ap 0.0--- (no predictions for this class)
---class bird ap 0.0--- (no predictions for this class)
---class boat ap 0.0--- (no predictions for this class)
---class bottle ap 0.0--- (no predictions for this class)
---class bus ap 0.0--- (no predictions for this class)
---class car ap 0.0--- (no predictions for this class)
---class cat ap 0.0--- (no predictions for this class)
---class chair ap 0.0--- (no predictions for this class)
---class cow ap 0.0--- (no predictions for this class)
---class diningtable ap 0.0--- (no predictions for this class)
---class dog ap 0.0--- (no predictions for this class)
---class horse ap 0.0--- (no predictions for this class)
---class motorbike ap 0.0--- (no predictions for this class)
---class person ap 0.0--- (no predictions for this class)
---class pottedplant ap 0.0--- (no predictions for this class)
---class sheep ap 0.0--- (no predictions for this class)
---class sofa ap 0.0--- (no predictions for this class)
---class train ap 0.0--- (no predictions for this class)
---class tvmonitor ap 0.0--- (no predictions for this class)
---map 0.0---
4 [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
0.0, 0.0, 0.0, 0.0, 0.0]

Updating best test loss: 3.66728

Starting epoch 6 / 50

Learning Rate for this epoch: 0.001

Epoch [6/50], Iter [50/209], Loss: total=3.335, reg=1.534, containing_obj=0.952, no_obj=0.066, cls=0.784

Epoch [6/50], Iter [100/209], Loss: total=3.348, reg=1.557, containing_obj=0.945, no_obj=0.065, cls=0.780

Epoch [6/50], Iter [150/209], Loss: total=3.405, reg=1.583, containing_obj=0.971, no_obj=0.064, cls=0.787

Epoch [6/50], Iter [200/209], Loss: total=3.371, reg=1.571, containing_obj=0.966, no_obj=0.063, cls=0.771

Updating best test loss: 3.58269

Starting epoch 7 / 50

Learning Rate for this epoch: 0.001

Epoch [7/50], Iter [50/209], Loss: total=3.310, reg=1.555, containing_obj=0.997, no_obj=0.059, cls=0.699

Epoch [7/50], Iter [100/209], Loss: total=3.281, reg=1.546, containing_obj=0.994, no_obj=0.059, cls=0.682

Epoch [7/50], Iter [150/209], Loss: total=3.210, reg=1.513, containing_obj=0.979, no_obj=0.058, cls=0.660

Epoch [7/50], Iter [200/209], Loss: total=3.220, reg=1.514, containing_obj=0.985, no_obj=0.058, cls=0.663

Updating best test loss: 3.50733

Starting epoch 8 / 50

Learning Rate for this epoch: 0.001

Epoch [8/50], Iter [50/209], Loss: total=3.192, reg=1.520, containing_obj=1.012, no_obj=0.056, cls=0.604

Epoch [8/50], Iter [100/209], Loss: total=3.126, reg=1.492, containing_obj=0.992, no_obj=0.056, cls=0.587

Epoch [8/50], Iter [150/209], Loss: total=3.119, reg=1.486, containing_obj=0.992, no_obj=0.056, cls=0.585

Epoch [8/50], Iter [200/209], Loss: total=3.110, reg=1.481, containing_obj=0.989, no_obj=0.056, cls=0.584

Updating best test loss: 3.36429

Starting epoch 9 / 50

Learning Rate for this epoch: 0.001

Epoch [9/50], Iter [50/209], Loss: total=3.065, reg=1.489, containing_obj=0.973, no_obj=0.060, cls=0.543

Epoch [9/50], Iter [100/209], Loss: total=3.048, reg=1.460,
containing_obj=0.972, no_obj=0.062, cls=0.553
Epoch [9/50], Iter [150/209], Loss: total=3.024, reg=1.453,
containing_obj=0.956, no_obj=0.067, cls=0.548
Epoch [9/50], Iter [200/209], Loss: total=2.995, reg=1.442,
containing_obj=0.942, no_obj=0.074, cls=0.537
Updating best test loss: 3.21619

Starting epoch 10 / 50

Learning Rate for this epoch: 0.001

Epoch [10/50], Iter [50/209], Loss: total=2.815, reg=1.380,
containing_obj=0.843, no_obj=0.127, cls=0.465
Epoch [10/50], Iter [100/209], Loss: total=2.853, reg=1.405,
containing_obj=0.843, no_obj=0.132, cls=0.473
Epoch [10/50], Iter [150/209], Loss: total=2.828, reg=1.380,
containing_obj=0.845, no_obj=0.134, cls=0.469
Epoch [10/50], Iter [200/209], Loss: total=2.833, reg=1.388,
containing_obj=0.837, no_obj=0.135, cls=0.473

---Evaluate model on test samples---

100%| | 4950/4950 [02:16<00:00, 36.20it/s]

---class aeroplane ap 0.37014882268865834---
---class bicycle ap 0.33202984858019907---
---class bird ap 0.32980370752002586---
---class boat ap 0.17872398942594217---
---class bottle ap 0.03817529058860585---
---class bus ap 0.39699784808851935---
---class car ap 0.5011312021898204---
---class cat ap 0.48249772972528926---
---class chair ap 0.166253515159359---
---class cow ap 0.3905101878308053---
---class diningtable ap 0.0048543689320388345---
---class dog ap 0.5036399047043938---
---class horse ap 0.5081073248781507---
---class motorbike ap 0.20567201037109883---
---class person ap 0.37773878026190766---
---class pottedplant ap 0.0719055271847999---
---class sheep ap 0.21640730222603838---
---class sofa ap 0.14106169740623525---
---class train ap 0.49811441302951087---
---class tvmonitor ap 0.40449863737274727---
---map 0.3059136054082073---

9 [0.37014882268865834, 0.33202984858019907, 0.32980370752002586,
0.17872398942594217, 0.03817529058860585, 0.39699784808851935,
0.5011312021898204, 0.48249772972528926, 0.166253515159359, 0.3905101878308053,
0.0048543689320388345, 0.5036399047043938, 0.5081073248781507,
0.20567201037109883, 0.37773878026190766, 0.0719055271847999,

0.21640730222603838, 0.14106169740623525, 0.49811441302951087,
0.40449863737274727]
Updating best test loss: 3.14993

Starting epoch 11 / 50

Learning Rate for this epoch: 0.001

Epoch [11/50], Iter [50/209], Loss: total=2.708, reg=1.310,
containing_obj=0.816, no_obj=0.144, cls=0.438
Epoch [11/50], Iter [100/209], Loss: total=2.716, reg=1.325,
containing_obj=0.829, no_obj=0.141, cls=0.422
Epoch [11/50], Iter [150/209], Loss: total=2.742, reg=1.334,
containing_obj=0.831, no_obj=0.140, cls=0.438
Epoch [11/50], Iter [200/209], Loss: total=2.757, reg=1.354,
containing_obj=0.825, no_obj=0.141, cls=0.438

Starting epoch 12 / 50

Learning Rate for this epoch: 0.001

Epoch [12/50], Iter [50/209], Loss: total=2.700, reg=1.327,
containing_obj=0.821, no_obj=0.142, cls=0.410
Epoch [12/50], Iter [100/209], Loss: total=2.685, reg=1.328,
containing_obj=0.803, no_obj=0.141, cls=0.413
Epoch [12/50], Iter [150/209], Loss: total=2.671, reg=1.319,
containing_obj=0.805, no_obj=0.140, cls=0.406
Epoch [12/50], Iter [200/209], Loss: total=2.682, reg=1.311,
containing_obj=0.819, no_obj=0.139, cls=0.413

Starting epoch 13 / 50

Learning Rate for this epoch: 0.001

Epoch [13/50], Iter [50/209], Loss: total=2.524, reg=1.224,
containing_obj=0.795, no_obj=0.138, cls=0.368
Epoch [13/50], Iter [100/209], Loss: total=2.596, reg=1.255,
containing_obj=0.820, no_obj=0.135, cls=0.385
Epoch [13/50], Iter [150/209], Loss: total=2.643, reg=1.286,
containing_obj=0.839, no_obj=0.134, cls=0.384
Epoch [13/50], Iter [200/209], Loss: total=2.648, reg=1.295,
containing_obj=0.831, no_obj=0.134, cls=0.389
Updating best test loss: 3.03751

Starting epoch 14 / 50

Learning Rate for this epoch: 0.001

Epoch [14/50], Iter [50/209], Loss: total=2.533, reg=1.245,
containing_obj=0.777, no_obj=0.141, cls=0.370
Epoch [14/50], Iter [100/209], Loss: total=2.480, reg=1.214,
containing_obj=0.777, no_obj=0.142, cls=0.347

Epoch [14/50], Iter [150/209], Loss: total=2.520, reg=1.227,
containing_obj=0.797, no_obj=0.143, cls=0.354
Epoch [14/50], Iter [200/209], Loss: total=2.537, reg=1.231,
containing_obj=0.802, no_obj=0.142, cls=0.362

Starting epoch 15 / 50

Learning Rate for this epoch: 0.001

Epoch [15/50], Iter [50/209], Loss: total=2.447, reg=1.170,
containing_obj=0.820, no_obj=0.134, cls=0.323
Epoch [15/50], Iter [100/209], Loss: total=2.434, reg=1.165,
containing_obj=0.806, no_obj=0.140, cls=0.324
Epoch [15/50], Iter [150/209], Loss: total=2.458, reg=1.177,
containing_obj=0.813, no_obj=0.137, cls=0.330
Epoch [15/50], Iter [200/209], Loss: total=2.520, reg=1.217,
containing_obj=0.831, no_obj=0.136, cls=0.337

---Evaluate model on test samples---

100%| | 4950/4950 [02:24<00:00, 34.36it/s]

---class aeroplane ap 0.4356732454051556---
---class bicycle ap 0.44810200781632786---
---class bird ap 0.3801434682582944---
---class boat ap 0.29092424903022696---
---class bottle ap 0.12055072478432052---
---class bus ap 0.5441086483550813---
---class car ap 0.5993128759021323---
---class cat ap 0.5496272609274507---
---class chair ap 0.20023555237683116---
---class cow ap 0.49366906046884235---
---class diningtable ap 0.14186838954549316---
---class dog ap 0.5555851985361089---
---class horse ap 0.5968196230312743---
---class motorbike ap 0.4222586849264899---
---class person ap 0.37512585571712664---
---class pottedplant ap 0.1738742044375341---
---class sheep ap 0.416266436233785---
---class sofa ap 0.30789329560014533---
---class train ap 0.5761414438038914---
---class tvmonitor ap 0.36003730538745715---
---map 0.3994108765271985---

14 [0.4356732454051556, 0.44810200781632786, 0.3801434682582944,
0.29092424903022696, 0.12055072478432052, 0.5441086483550813,
0.5993128759021323, 0.5496272609274507, 0.20023555237683116,
0.49366906046884235, 0.14186838954549316, 0.5555851985361089,
0.5968196230312743, 0.4222586849264899, 0.37512585571712664, 0.1738742044375341,
0.416266436233785, 0.30789329560014533, 0.5761414438038914, 0.36003730538745715]

Starting epoch 16 / 50
Learning Rate for this epoch: 0.001
Epoch [16/50], Iter [50/209], Loss: total=2.463, reg=1.200,
containing_obj=0.814, no_obj=0.141, cls=0.308
Epoch [16/50], Iter [100/209], Loss: total=2.486, reg=1.211,
containing_obj=0.834, no_obj=0.135, cls=0.306
Epoch [16/50], Iter [150/209], Loss: total=2.447, reg=1.183,
containing_obj=0.818, no_obj=0.137, cls=0.309
Epoch [16/50], Iter [200/209], Loss: total=2.456, reg=1.179,
containing_obj=0.826, no_obj=0.135, cls=0.315
Updating best test loss: 2.94462

Starting epoch 17 / 50
Learning Rate for this epoch: 0.001
Epoch [17/50], Iter [50/209], Loss: total=2.406, reg=1.157,
containing_obj=0.822, no_obj=0.140, cls=0.287
Epoch [17/50], Iter [100/209], Loss: total=2.414, reg=1.163,
containing_obj=0.819, no_obj=0.144, cls=0.288
Epoch [17/50], Iter [150/209], Loss: total=2.398, reg=1.142,
containing_obj=0.821, no_obj=0.142, cls=0.293
Epoch [17/50], Iter [200/209], Loss: total=2.408, reg=1.147,
containing_obj=0.823, no_obj=0.140, cls=0.298

Starting epoch 18 / 50
Learning Rate for this epoch: 0.001
Epoch [18/50], Iter [50/209], Loss: total=2.487, reg=1.203,
containing_obj=0.847, no_obj=0.137, cls=0.300
Epoch [18/50], Iter [100/209], Loss: total=2.434, reg=1.156,
containing_obj=0.836, no_obj=0.139, cls=0.302
Epoch [18/50], Iter [150/209], Loss: total=2.404, reg=1.140,
containing_obj=0.833, no_obj=0.137, cls=0.293
Epoch [18/50], Iter [200/209], Loss: total=2.401, reg=1.137,
containing_obj=0.833, no_obj=0.137, cls=0.293

Starting epoch 19 / 50
Learning Rate for this epoch: 0.001
Epoch [19/50], Iter [50/209], Loss: total=2.349, reg=1.104,
containing_obj=0.829, no_obj=0.132, cls=0.284
Epoch [19/50], Iter [100/209], Loss: total=2.315, reg=1.100,
containing_obj=0.808, no_obj=0.136, cls=0.271
Epoch [19/50], Iter [150/209], Loss: total=2.338, reg=1.111,
containing_obj=0.814, no_obj=0.140, cls=0.272
Epoch [19/50], Iter [200/209], Loss: total=2.334, reg=1.105,
containing_obj=0.815, no_obj=0.141, cls=0.272
Updating best test loss: 2.92473

Starting epoch 20 / 50

Learning Rate for this epoch: 0.001

Epoch [20/50], Iter [50/209], Loss: total=2.296, reg=1.087,
containing_obj=0.826, no_obj=0.132, cls=0.251

Epoch [20/50], Iter [100/209], Loss: total=2.329, reg=1.105,
containing_obj=0.832, no_obj=0.136, cls=0.256

Epoch [20/50], Iter [150/209], Loss: total=2.308, reg=1.094,
containing_obj=0.824, no_obj=0.137, cls=0.252

Epoch [20/50], Iter [200/209], Loss: total=2.291, reg=1.086,
containing_obj=0.812, no_obj=0.138, cls=0.254

---Evaluate model on test samples---

100%| | 4950/4950 [02:23<00:00, 34.58it/s]

---class aeroplane ap 0.4500290385934387---

---class bicycle ap 0.6119677548336837---

---class bird ap 0.47045234354917054---

---class boat ap 0.33807746224643087---

---class bottle ap 0.17938049648987325---

---class bus ap 0.5857919548744188---

---class car ap 0.6546783146730449---

---class cat ap 0.6193989618047662---

---class chair ap 0.2984238286421016---

---class cow ap 0.5704814432674307---

---class diningtable ap 0.1726416483340955---

---class dog ap 0.6256145581359478---

---class horse ap 0.6486127212155299---

---class motorbike ap 0.4805464038752212---

---class person ap 0.49900208745793134---

---class pottedplant ap 0.2053964101368331---

---class sheep ap 0.4259484914784117---

---class sofa ap 0.4440582509298917---

---class train ap 0.6016929814373102---

---class tvmonitor ap 0.4807186224829285---

---map 0.46814568872292306---

19 [0.4500290385934387, 0.6119677548336837, 0.47045234354917054,
0.33807746224643087, 0.17938049648987325, 0.5857919548744188,
0.6546783146730449, 0.6193989618047662, 0.2984238286421016, 0.5704814432674307,
0.1726416483340955, 0.6256145581359478, 0.6486127212155299, 0.4805464038752212,
0.49900208745793134, 0.2053964101368331, 0.4259484914784117, 0.4440582509298917,
0.6016929814373102, 0.4807186224829285]

Updating best test loss: 2.91945

Starting epoch 21 / 50

Learning Rate for this epoch: 0.001

Epoch [21/50], Iter [50/209], Loss: total=2.316, reg=1.075,

containing_obj=0.848, no_obj=0.137, cls=0.256
Epoch [21/50], Iter [100/209], Loss: total=2.253, reg=1.031,
containing_obj=0.825, no_obj=0.139, cls=0.258
Epoch [21/50], Iter [150/209], Loss: total=2.287, reg=1.053,
containing_obj=0.832, no_obj=0.139, cls=0.264
Epoch [21/50], Iter [200/209], Loss: total=2.274, reg=1.062,
containing_obj=0.821, no_obj=0.137, cls=0.255

Starting epoch 22 / 50

Learning Rate for this epoch: 0.001
Epoch [22/50], Iter [50/209], Loss: total=2.252, reg=1.053,
containing_obj=0.814, no_obj=0.143, cls=0.242
Epoch [22/50], Iter [100/209], Loss: total=2.232, reg=1.035,
containing_obj=0.816, no_obj=0.143, cls=0.238
Epoch [22/50], Iter [150/209], Loss: total=2.239, reg=1.040,
containing_obj=0.821, no_obj=0.140, cls=0.238
Epoch [22/50], Iter [200/209], Loss: total=2.231, reg=1.039,
containing_obj=0.814, no_obj=0.142, cls=0.236
Updating best test loss: 2.87325

Starting epoch 23 / 50

Learning Rate for this epoch: 0.001
Epoch [23/50], Iter [50/209], Loss: total=2.135, reg=0.992,
containing_obj=0.770, no_obj=0.151, cls=0.221
Epoch [23/50], Iter [100/209], Loss: total=2.119, reg=1.002,
containing_obj=0.749, no_obj=0.154, cls=0.215
Epoch [23/50], Iter [150/209], Loss: total=2.148, reg=1.006,
containing_obj=0.750, no_obj=0.161, cls=0.230
Epoch [23/50], Iter [200/209], Loss: total=2.157, reg=1.021,
containing_obj=0.738, no_obj=0.167, cls=0.230
Updating best test loss: 2.85476

Starting epoch 24 / 50

Learning Rate for this epoch: 0.001
Epoch [24/50], Iter [50/209], Loss: total=2.188, reg=1.058,
containing_obj=0.676, no_obj=0.205, cls=0.248
Epoch [24/50], Iter [100/209], Loss: total=2.099, reg=1.018,
containing_obj=0.646, no_obj=0.207, cls=0.227
Epoch [24/50], Iter [150/209], Loss: total=2.079, reg=1.005,
containing_obj=0.639, no_obj=0.208, cls=0.226
Epoch [24/50], Iter [200/209], Loss: total=2.075, reg=1.004,
containing_obj=0.641, no_obj=0.207, cls=0.222
Updating best test loss: 2.76584

```

Starting epoch 25 / 50
Learning Rate for this epoch: 0.001
Epoch [25/50], Iter [50/209], Loss: total=2.002, reg=0.976,
containing_obj=0.618, no_obj=0.209, cls=0.199
Epoch [25/50], Iter [100/209], Loss: total=2.063, reg=1.007,
containing_obj=0.640, no_obj=0.209, cls=0.206
Epoch [25/50], Iter [150/209], Loss: total=2.033, reg=0.988,
containing_obj=0.627, no_obj=0.207, cls=0.211
Epoch [25/50], Iter [200/209], Loss: total=2.028, reg=0.985,
containing_obj=0.624, no_obj=0.209, cls=0.210
---Evaluate model on test samples---

100%|      | 4950/4950 [02:09<00:00, 38.27it/s]

---class aeroplane ap 0.4493943018533752---
---class bicycle ap 0.5815049295793329---
---class bird ap 0.40438008537311365---
---class boat ap 0.3059144934344722---
---class bottle ap 0.1880503664061213---
---class bus ap 0.5294113446256266---
---class car ap 0.6277563886035822---
---class cat ap 0.6563460194727262---
---class chair ap 0.2818475300113292---
---class cow ap 0.46979663971899654---
---class diningtable ap 0.2724369353427165---
---class dog ap 0.6138955865319873---
---class horse ap 0.6640598729179764---
---class motorbike ap 0.48992167955986476---
---class person ap 0.5380842520165102---
---class pottedplant ap 0.15371697792721345---
---class sheep ap 0.4520906661645119---
---class sofa ap 0.4427459560395489---
---class train ap 0.5235328183834421---
---class tvmonitor ap 0.46897552682954124---
---map 0.45569311853959943---
24 [0.4493943018533752, 0.5815049295793329, 0.40438008537311365,
0.3059144934344722, 0.1880503664061213, 0.5294113446256266, 0.6277563886035822,
0.6563460194727262, 0.2818475300113292, 0.46979663971899654, 0.2724369353427165,
0.6138955865319873, 0.6640598729179764, 0.48992167955986476, 0.5380842520165102,
0.15371697792721345, 0.4520906661645119, 0.4427459560395489, 0.5235328183834421,
0.46897552682954124]
Updating best test loss: 2.73554

```

```

Starting epoch 26 / 50
Learning Rate for this epoch: 0.001
Epoch [26/50], Iter [50/209], Loss: total=1.988, reg=0.970,
containing_obj=0.620, no_obj=0.206, cls=0.192
Epoch [26/50], Iter [100/209], Loss: total=2.002, reg=0.961,

```


containing_obj=0.625, no_obj=0.207, cls=0.209
Epoch [26/50], Iter [150/209], Loss: total=2.005, reg=0.967,
containing_obj=0.618, no_obj=0.209, cls=0.211
Epoch [26/50], Iter [200/209], Loss: total=1.999, reg=0.962,
containing_obj=0.612, no_obj=0.208, cls=0.217

Starting epoch 27 / 50

Learning Rate for this epoch: 0.001

Epoch [27/50], Iter [50/209], Loss: total=1.905, reg=0.920,
containing_obj=0.591, no_obj=0.201, cls=0.194
Epoch [27/50], Iter [100/209], Loss: total=1.956, reg=0.936,
containing_obj=0.614, no_obj=0.207, cls=0.200
Epoch [27/50], Iter [150/209], Loss: total=1.966, reg=0.946,
containing_obj=0.610, no_obj=0.207, cls=0.203
Epoch [27/50], Iter [200/209], Loss: total=1.973, reg=0.952,
containing_obj=0.607, no_obj=0.208, cls=0.206

Starting epoch 28 / 50

Learning Rate for this epoch: 0.001

Epoch [28/50], Iter [50/209], Loss: total=1.973, reg=0.961,
containing_obj=0.606, no_obj=0.203, cls=0.203
Epoch [28/50], Iter [100/209], Loss: total=1.901, reg=0.931,
containing_obj=0.580, no_obj=0.201, cls=0.189
Epoch [28/50], Iter [150/209], Loss: total=1.906, reg=0.939,
containing_obj=0.581, no_obj=0.201, cls=0.185
Epoch [28/50], Iter [200/209], Loss: total=1.932, reg=0.954,
containing_obj=0.591, no_obj=0.200, cls=0.187

Starting epoch 29 / 50

Learning Rate for this epoch: 0.001

Epoch [29/50], Iter [50/209], Loss: total=1.883, reg=0.899,
containing_obj=0.585, no_obj=0.205, cls=0.194
Epoch [29/50], Iter [100/209], Loss: total=1.848, reg=0.889,
containing_obj=0.576, no_obj=0.202, cls=0.180
Epoch [29/50], Iter [150/209], Loss: total=1.889, reg=0.917,
containing_obj=0.583, no_obj=0.203, cls=0.186
Epoch [29/50], Iter [200/209], Loss: total=1.884, reg=0.916,
containing_obj=0.580, no_obj=0.204, cls=0.185

Starting epoch 30 / 50

Learning Rate for this epoch: 0.001

Epoch [30/50], Iter [50/209], Loss: total=1.783, reg=0.856,
containing_obj=0.565, no_obj=0.201, cls=0.161
Epoch [30/50], Iter [100/209], Loss: total=1.828, reg=0.871,

```

containing_obj=0.573, no_obj=0.205, cls=0.180
Epoch [30/50], Iter [150/209], Loss: total=1.835, reg=0.875,
containing_obj=0.571, no_obj=0.209, cls=0.180
Epoch [30/50], Iter [200/209], Loss: total=1.849, reg=0.882,
containing_obj=0.575, no_obj=0.208, cls=0.184
---Evaluate model on test samples---
100%|      | 4950/4950 [02:11<00:00, 37.68it/s]
---class aeroplane ap 0.4370347377431284---
---class bicycle ap 0.6063564896750449---
---class bird ap 0.4513877136383035---
---class boat ap 0.3177819913946993---
---class bottle ap 0.22718827519784496---
---class bus ap 0.6232368912901032---
---class car ap 0.6722871253118491---
---class cat ap 0.5940839534050067---
---class chair ap 0.32175522703517934---
---class cow ap 0.5520750850568417---
---class diningtable ap 0.2596434021904313---
---class dog ap 0.5652328641056429---
---class horse ap 0.6683811709784176---
---class motorbike ap 0.5780252501379156---
---class person ap 0.5511284046251623---
---class pottedplant ap 0.20948417858732127---
---class sheep ap 0.44027965034567046---
---class sofa ap 0.4079782852830158---
---class train ap 0.5780049653813829---
---class tvmonitor ap 0.5218785526769412---
---map 0.47916121070299517---
29 [0.4370347377431284, 0.6063564896750449, 0.4513877136383035,
0.3177819913946993, 0.22718827519784496, 0.6232368912901032, 0.6722871253118491,
0.5940839534050067, 0.32175522703517934, 0.5520750850568417, 0.2596434021904313,
0.5652328641056429, 0.6683811709784176, 0.5780252501379156, 0.5511284046251623,
0.20948417858732127, 0.44027965034567046, 0.4079782852830158,
0.5780049653813829, 0.5218785526769412]

```

Starting epoch 31 / 50

Learning Rate for this epoch: 0.0001

```

Epoch [31/50], Iter [50/209], Loss: total=1.789, reg=0.856,
containing_obj=0.585, no_obj=0.200, cls=0.148
Epoch [31/50], Iter [100/209], Loss: total=1.808, reg=0.866,
containing_obj=0.591, no_obj=0.200, cls=0.152
Epoch [31/50], Iter [150/209], Loss: total=1.788, reg=0.854,
containing_obj=0.571, no_obj=0.203, cls=0.160
Epoch [31/50], Iter [200/209], Loss: total=1.784, reg=0.856,
containing_obj=0.564, no_obj=0.203, cls=0.161

```

Updating best test loss: 2.66404

Starting epoch 32 / 50

Learning Rate for this epoch: 0.0001

Epoch [32/50], Iter [50/209], Loss: total=1.718, reg=0.824,
containing_obj=0.540, no_obj=0.207, cls=0.147
Epoch [32/50], Iter [100/209], Loss: total=1.700, reg=0.817,
containing_obj=0.533, no_obj=0.206, cls=0.143
Epoch [32/50], Iter [150/209], Loss: total=1.739, reg=0.833,
containing_obj=0.549, no_obj=0.203, cls=0.153
Epoch [32/50], Iter [200/209], Loss: total=1.726, reg=0.825,
containing_obj=0.546, no_obj=0.204, cls=0.152

Starting epoch 33 / 50

Learning Rate for this epoch: 0.0001

Epoch [33/50], Iter [50/209], Loss: total=1.736, reg=0.837,
containing_obj=0.546, no_obj=0.194, cls=0.158
Epoch [33/50], Iter [100/209], Loss: total=1.690, reg=0.812,
containing_obj=0.529, no_obj=0.201, cls=0.147
Epoch [33/50], Iter [150/209], Loss: total=1.694, reg=0.809,
containing_obj=0.536, no_obj=0.203, cls=0.146
Epoch [33/50], Iter [200/209], Loss: total=1.689, reg=0.800,
containing_obj=0.541, no_obj=0.205, cls=0.143
Updating best test loss: 2.66111

Starting epoch 34 / 50

Learning Rate for this epoch: 0.0001

Epoch [34/50], Iter [50/209], Loss: total=1.700, reg=0.798,
containing_obj=0.551, no_obj=0.204, cls=0.146
Epoch [34/50], Iter [100/209], Loss: total=1.686, reg=0.789,
containing_obj=0.544, no_obj=0.207, cls=0.146
Epoch [34/50], Iter [150/209], Loss: total=1.667, reg=0.781,
containing_obj=0.539, no_obj=0.205, cls=0.142
Epoch [34/50], Iter [200/209], Loss: total=1.667, reg=0.785,
containing_obj=0.535, no_obj=0.208, cls=0.139
Updating best test loss: 2.65856

Starting epoch 35 / 50

Learning Rate for this epoch: 0.0001

Epoch [35/50], Iter [50/209], Loss: total=1.629, reg=0.779,
containing_obj=0.529, no_obj=0.202, cls=0.119
Epoch [35/50], Iter [100/209], Loss: total=1.646, reg=0.785,
containing_obj=0.530, no_obj=0.202, cls=0.128
Epoch [35/50], Iter [150/209], Loss: total=1.638, reg=0.776,
containing_obj=0.529, no_obj=0.204, cls=0.129

Epoch [35/50], Iter [200/209], Loss: total=1.652, reg=0.782,
containing_obj=0.532, no_obj=0.205, cls=0.133
---Evaluate model on test samples---

100%| | 4950/4950 [02:26<00:00, 33.73it/s]

---class aeroplane ap 0.5124813531391292---
---class bicycle ap 0.6406666728820745---
---class bird ap 0.5267294913528213---
---class boat ap 0.3520827537300927---
---class bottle ap 0.23707125023635497---
---class bus ap 0.6381050036675867---
---class car ap 0.6898492314049264---
---class cat ap 0.6869430381256758---
---class chair ap 0.32649655488905494---
---class cow ap 0.5229444279305109---
---class diningtable ap 0.32013941508431654---
---class dog ap 0.6362180706662386---
---class horse ap 0.6834630980947245---
---class motorbike ap 0.6141312524877574---
---class person ap 0.5667527822081906---
---class pottedplant ap 0.22793007147545105---
---class sheep ap 0.46117512972559727---
---class sofa ap 0.4738739595996908---
---class train ap 0.6174561207624698---
---class tvmonitor ap 0.5261043183866694---
---map 0.5130306997924666---

34 [0.5124813531391292, 0.6406666728820745, 0.5267294913528213,
0.3520827537300927, 0.23707125023635497, 0.6381050036675867, 0.6898492314049264,
0.6869430381256758, 0.32649655488905494, 0.5229444279305109,
0.32013941508431654, 0.6362180706662386, 0.6834630980947245, 0.6141312524877574,
0.5667527822081906, 0.22793007147545105, 0.46117512972559727,
0.4738739595996908, 0.6174561207624698, 0.5261043183866694]

Starting epoch 36 / 50

Learning Rate for this epoch: 0.0001

Epoch [36/50], Iter [50/209], Loss: total=1.661, reg=0.783,
containing_obj=0.522, no_obj=0.214, cls=0.141

Epoch [36/50], Iter [100/209], Loss: total=1.698, reg=0.798,
containing_obj=0.542, no_obj=0.212, cls=0.145

Epoch [36/50], Iter [150/209], Loss: total=1.685, reg=0.790,
containing_obj=0.537, no_obj=0.210, cls=0.147

Epoch [36/50], Iter [200/209], Loss: total=1.663, reg=0.782,
containing_obj=0.532, no_obj=0.208, cls=0.140

Starting epoch 37 / 50

Learning Rate for this epoch: 0.0001

Epoch [37/50], Iter [50/209], Loss: total=1.622, reg=0.756,
containing_obj=0.531, no_obj=0.206, cls=0.130
Epoch [37/50], Iter [100/209], Loss: total=1.586, reg=0.736,
containing_obj=0.517, no_obj=0.208, cls=0.125
Epoch [37/50], Iter [150/209], Loss: total=1.590, reg=0.738,
containing_obj=0.517, no_obj=0.206, cls=0.128
Epoch [37/50], Iter [200/209], Loss: total=1.609, reg=0.755,
containing_obj=0.520, no_obj=0.206, cls=0.128
Updating best test loss: 2.65273

Starting epoch 38 / 50

Learning Rate for this epoch: 0.0001

Epoch [38/50], Iter [50/209], Loss: total=1.565, reg=0.728,
containing_obj=0.507, no_obj=0.204, cls=0.125
Epoch [38/50], Iter [100/209], Loss: total=1.584, reg=0.736,
containing_obj=0.519, no_obj=0.201, cls=0.129
Epoch [38/50], Iter [150/209], Loss: total=1.594, reg=0.744,
containing_obj=0.517, no_obj=0.205, cls=0.127
Epoch [38/50], Iter [200/209], Loss: total=1.604, reg=0.752,
containing_obj=0.517, no_obj=0.206, cls=0.129

Starting epoch 39 / 50

Learning Rate for this epoch: 0.0001

Epoch [39/50], Iter [50/209], Loss: total=1.599, reg=0.757,
containing_obj=0.521, no_obj=0.199, cls=0.122
Epoch [39/50], Iter [100/209], Loss: total=1.596, reg=0.756,
containing_obj=0.520, no_obj=0.200, cls=0.122
Epoch [39/50], Iter [150/209], Loss: total=1.592, reg=0.752,
containing_obj=0.515, no_obj=0.205, cls=0.121
Epoch [39/50], Iter [200/209], Loss: total=1.599, reg=0.752,
containing_obj=0.516, no_obj=0.209, cls=0.121

Starting epoch 40 / 50

Learning Rate for this epoch: 0.0001

Epoch [40/50], Iter [50/209], Loss: total=1.646, reg=0.776,
containing_obj=0.539, no_obj=0.200, cls=0.131
Epoch [40/50], Iter [100/209], Loss: total=1.586, reg=0.747,
containing_obj=0.514, no_obj=0.204, cls=0.122
Epoch [40/50], Iter [150/209], Loss: total=1.589, reg=0.748,
containing_obj=0.516, no_obj=0.204, cls=0.121
Epoch [40/50], Iter [200/209], Loss: total=1.598, reg=0.751,
containing_obj=0.516, no_obj=0.206, cls=0.125
---Evaluate model on test samples---

100%| | 4950/4950 [02:19<00:00, 35.40it/s]

```

---class aeroplane ap 0.5041299236726144---
---class bicycle ap 0.6702644041146242---
---class bird ap 0.5443064755457228---
---class boat ap 0.34595750439014633---
---class bottle ap 0.24733893947292582---
---class bus ap 0.6477820421617064---
---class car ap 0.6934408739149409---
---class cat ap 0.6793073088527579---
---class chair ap 0.3329980944348795---
---class cow ap 0.5591092977325537---
---class diningtable ap 0.3303599087870781---
---class dog ap 0.6339171482342213---
---class horse ap 0.7086375921419621---
---class motorbike ap 0.6158403105188133---
---class person ap 0.565730483438492---
---class pottedplant ap 0.22363083087776076---
---class sheep ap 0.5013237735445348---
---class sofa ap 0.4762274064285803---
---class train ap 0.6200301629226648---
---class tvmonitor ap 0.5077442503051058---
---map 0.5204038365746043---
39 [0.5041299236726144, 0.6702644041146242, 0.5443064755457228,
0.34595750439014633, 0.24733893947292582, 0.6477820421617064,
0.6934408739149409, 0.6793073088527579, 0.3329980944348795, 0.5591092977325537,
0.3303599087870781, 0.6339171482342213, 0.7086375921419621, 0.6158403105188133,
0.565730483438492, 0.22363083087776076, 0.5013237735445348, 0.4762274064285803,
0.6200301629226648, 0.5077442503051058]
Updating best test loss: 2.63933

```

```

Starting epoch 41 / 50
Learning Rate for this epoch: 1e-05
Epoch [41/50], Iter [50/209], Loss: total=1.688, reg=0.812,
containing_obj=0.558, no_obj=0.200, cls=0.119
Epoch [41/50], Iter [100/209], Loss: total=1.604, reg=0.769,
containing_obj=0.523, no_obj=0.204, cls=0.109
Epoch [41/50], Iter [150/209], Loss: total=1.588, reg=0.756,
containing_obj=0.519, no_obj=0.204, cls=0.108
Epoch [41/50], Iter [200/209], Loss: total=1.588, reg=0.753,
containing_obj=0.517, no_obj=0.204, cls=0.114

```

```

Starting epoch 42 / 50
Learning Rate for this epoch: 1e-05
Epoch [42/50], Iter [50/209], Loss: total=1.491, reg=0.693,
containing_obj=0.490, no_obj=0.201, cls=0.108
Epoch [42/50], Iter [100/209], Loss: total=1.529, reg=0.716,
containing_obj=0.501, no_obj=0.196, cls=0.115

```

Epoch [42/50], Iter [150/209], Loss: total=1.536, reg=0.722,
containing_obj=0.502, no_obj=0.196, cls=0.117
Epoch [42/50], Iter [200/209], Loss: total=1.545, reg=0.722,
containing_obj=0.506, no_obj=0.199, cls=0.118

Starting epoch 43 / 50

Learning Rate for this epoch: 1e-05

Epoch [43/50], Iter [50/209], Loss: total=1.499, reg=0.708,
containing_obj=0.483, no_obj=0.207, cls=0.101
Epoch [43/50], Iter [100/209], Loss: total=1.561, reg=0.740,
containing_obj=0.503, no_obj=0.206, cls=0.112
Epoch [43/50], Iter [150/209], Loss: total=1.577, reg=0.750,
containing_obj=0.506, no_obj=0.201, cls=0.119
Epoch [43/50], Iter [200/209], Loss: total=1.573, reg=0.745,
containing_obj=0.505, no_obj=0.203, cls=0.120

Starting epoch 44 / 50

Learning Rate for this epoch: 1e-05

Epoch [44/50], Iter [50/209], Loss: total=1.470, reg=0.689,
containing_obj=0.467, no_obj=0.210, cls=0.106
Epoch [44/50], Iter [100/209], Loss: total=1.548, reg=0.732,
containing_obj=0.493, no_obj=0.205, cls=0.118
Epoch [44/50], Iter [150/209], Loss: total=1.556, reg=0.728,
containing_obj=0.504, no_obj=0.206, cls=0.118
Epoch [44/50], Iter [200/209], Loss: total=1.572, reg=0.737,
containing_obj=0.510, no_obj=0.205, cls=0.121

Starting epoch 45 / 50

Learning Rate for this epoch: 1e-05

Epoch [45/50], Iter [50/209], Loss: total=1.631, reg=0.771,
containing_obj=0.538, no_obj=0.202, cls=0.120
Epoch [45/50], Iter [100/209], Loss: total=1.594, reg=0.747,
containing_obj=0.527, no_obj=0.202, cls=0.118
Epoch [45/50], Iter [150/209], Loss: total=1.563, reg=0.728,
containing_obj=0.514, no_obj=0.206, cls=0.116
Epoch [45/50], Iter [200/209], Loss: total=1.556, reg=0.722,
containing_obj=0.514, no_obj=0.204, cls=0.116

---Evaluate model on test samples---

100%| | 4950/4950 [02:18<00:00, 35.72it/s]

---class aeroplane ap 0.5281957531118149---

---class bicycle ap 0.6787577370933039---

---class bird ap 0.5471821781273337---

---class boat ap 0.371997979324474---

---class bottle ap 0.24145822286658025---

```

---class bus ap 0.6387569872644105---
---class car ap 0.6910533909333798---
---class cat ap 0.6815432789489099---
---class chair ap 0.3370886467266352---
---class cow ap 0.556894605082036---
---class diningtable ap 0.336280279594243---
---class dog ap 0.6324824447097587---
---class horse ap 0.691316028434873---
---class motorbike ap 0.6006904397326223---
---class person ap 0.57728022472045---
---class pottedplant ap 0.22064712850499352---
---class sheep ap 0.49075543672158295---
---class sofa ap 0.4660795964177883---
---class train ap 0.6189256732133483---
---class tvmonitor ap 0.5111932286065323---
---map 0.5209289630067535---
44 [0.5281957531118149, 0.6787577370933039, 0.5471821781273337,
0.371997979324474, 0.24145822286658025, 0.6387569872644105, 0.6910533909333798,
0.6815432789489099, 0.3370886467266352, 0.556894605082036, 0.336280279594243,
0.6324824447097587, 0.691316028434873, 0.6006904397326223, 0.57728022472045,
0.22064712850499352, 0.49075543672158295, 0.4660795964177883,
0.6189256732133483, 0.5111932286065323]

```

Starting epoch 46 / 50

Learning Rate for this epoch: 1e-05

```

Epoch [46/50], Iter [50/209], Loss: total=1.555, reg=0.733,
containing_obj=0.505, no_obj=0.197, cls=0.119
Epoch [46/50], Iter [100/209], Loss: total=1.548, reg=0.723,
containing_obj=0.503, no_obj=0.204, cls=0.118
Epoch [46/50], Iter [150/209], Loss: total=1.555, reg=0.726,
containing_obj=0.507, no_obj=0.204, cls=0.118
Epoch [46/50], Iter [200/209], Loss: total=1.562, reg=0.729,
containing_obj=0.513, no_obj=0.204, cls=0.116

```

Starting epoch 47 / 50

Learning Rate for this epoch: 1e-05

```

Epoch [47/50], Iter [50/209], Loss: total=1.532, reg=0.733,
containing_obj=0.483, no_obj=0.203, cls=0.112
Epoch [47/50], Iter [100/209], Loss: total=1.550, reg=0.728,
containing_obj=0.505, no_obj=0.200, cls=0.117
Epoch [47/50], Iter [150/209], Loss: total=1.554, reg=0.728,
containing_obj=0.507, no_obj=0.202, cls=0.117
Epoch [47/50], Iter [200/209], Loss: total=1.548, reg=0.721,
containing_obj=0.506, no_obj=0.203, cls=0.118

```


Starting epoch 48 / 50
Learning Rate for this epoch: 1e-05
Epoch [48/50], Iter [50/209], Loss: total=1.509, reg=0.699,
containing_obj=0.491, no_obj=0.209, cls=0.110
Epoch [48/50], Iter [100/209], Loss: total=1.557, reg=0.728,
containing_obj=0.504, no_obj=0.206, cls=0.119
Epoch [48/50], Iter [150/209], Loss: total=1.565, reg=0.731,
containing_obj=0.508, no_obj=0.206, cls=0.119
Epoch [48/50], Iter [200/209], Loss: total=1.566, reg=0.728,
containing_obj=0.510, no_obj=0.208, cls=0.120

Starting epoch 49 / 50
Learning Rate for this epoch: 1e-05
Epoch [49/50], Iter [50/209], Loss: total=1.544, reg=0.711,
containing_obj=0.511, no_obj=0.204, cls=0.118
Epoch [49/50], Iter [100/209], Loss: total=1.550, reg=0.722,
containing_obj=0.506, no_obj=0.200, cls=0.122
Epoch [49/50], Iter [150/209], Loss: total=1.558, reg=0.734,
containing_obj=0.504, no_obj=0.200, cls=0.120
Epoch [49/50], Iter [200/209], Loss: total=1.575, reg=0.745,
containing_obj=0.511, no_obj=0.200, cls=0.119

Starting epoch 50 / 50
Learning Rate for this epoch: 1e-05
Epoch [50/50], Iter [50/209], Loss: total=1.563, reg=0.728,
containing_obj=0.521, no_obj=0.202, cls=0.112
Epoch [50/50], Iter [100/209], Loss: total=1.527, reg=0.714,
containing_obj=0.498, no_obj=0.207, cls=0.108
Epoch [50/50], Iter [150/209], Loss: total=1.543, reg=0.718,
containing_obj=0.505, no_obj=0.205, cls=0.115
Epoch [50/50], Iter [200/209], Loss: total=1.538, reg=0.715,
containing_obj=0.502, no_obj=0.206, cls=0.115
---Evaluate model on test samples---

100%| | 4950/4950 [02:19<00:00, 35.59it/s]

---class aeroplane ap 0.5313191239370854---
---class bicycle ap 0.6576073499532791---
---class bird ap 0.5366736391270497---
---class boat ap 0.37716124846787535---
---class bottle ap 0.24085737575341248---
---class bus ap 0.6485617792947089---
---class car ap 0.687869477840831---
---class cat ap 0.6731492633349193---
---class chair ap 0.3445235890792374---
---class cow ap 0.5548337618464732---
---class diningtable ap 0.350204457200313---

```

---class dog ap 0.6516651484153018---
---class horse ap 0.6954082953991244---
---class motorbike ap 0.6141076107499881---
---class person ap 0.5783747980777809---
---class pottedplant ap 0.22892791858007538---
---class sheep ap 0.4697638239596463---
---class sofa ap 0.4842278084485112---
---class train ap 0.6416915438594533---
---class tvmonitor ap 0.5286599705774611---
---map 0.5247793991951263---
49 [0.5313191239370854, 0.6576073499532791, 0.5366736391270497,
0.37716124846787535, 0.24085737575341248, 0.6485617792947089, 0.687869477840831,
0.6731492633349193, 0.3445235890792374, 0.5548337618464732, 0.350204457200313,
0.6516651484153018, 0.6954082953991244, 0.6141076107499881, 0.5783747980777809,
0.22892791858007538, 0.4697638239596463, 0.4842278084485112, 0.6416915438594533,
0.5286599705774611]

```

```

[ ]: ### Print Prediction ###
net.eval()

# select random image from test set
image_name = random.choice(test_dataset.fnames)
image = cv2.imread(os.path.join(file_root_test, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_test)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.4, 1)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
text_size[0], p1[1] + text_size[1]),
color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.
FONT_HERSHEY_SIMPLEX, 0.4, (255, 255, 255), 1, 8)

plt.figure(figsize = (15,15))
plt.imshow(image)

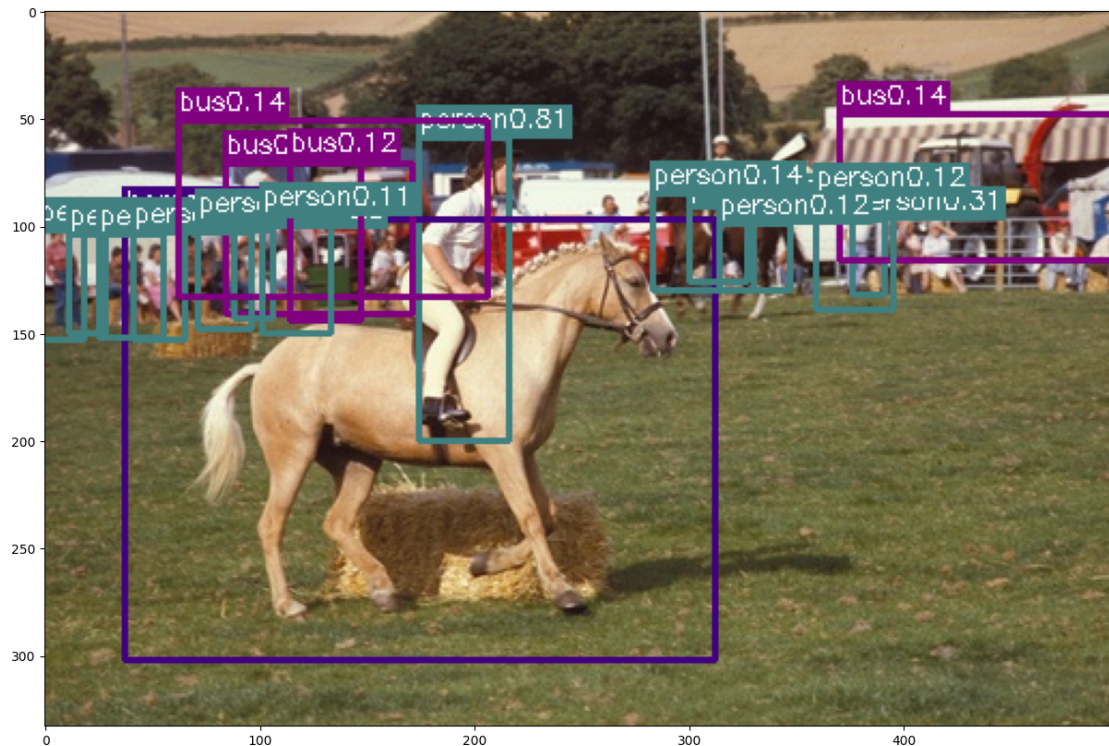
```

predicting...

```

[ ]: <matplotlib.image.AxesImage at 0x151c66ea7a90>

```



```
[ ]: ### Evaluate on Test ###
test_aps = evaluate(net, test_dataset_file=annotation_file_test,
img_root=file_root_test)
```

---Evaluate model on test samples---

100%| | 4950/4950 [02:24<00:00, 34.37it/s]

```
---class aeroplane ap 0.5313191239370854---
---class bicycle ap 0.6576073499532791---
---class bird ap 0.5366736391270497---
---class boat ap 0.37716124846787535---
---class bottle ap 0.24085737575341248---
---class bus ap 0.6485617792947089---
---class car ap 0.687869477840831---
---class cat ap 0.6731492633349193---
---class chair ap 0.3445235890792374---
---class cow ap 0.5548337618464732---
---class diningtable ap 0.350204457200313---
---class dog ap 0.6516651484153018---
---class horse ap 0.6954082953991244---
---class motorbike ap 0.6141076107499881---
---class person ap 0.5783747980777809---
---class pottedplant ap 0.22892791858007538---
```

```
---class sheep ap 0.4697638239596463---  
---class sofa ap 0.4842278084485112---  
---class train ap 0.6416915438594533---  
---class tvmonitor ap 0.5286599705774611---  
---map 0.5247793991951263---
```

4 Save as PDF

```
[ ]: %%capture  
  
from google.colab import drive  
drive.mount('/content/drive')  
# install tex; first run may take several minutes  
[!] apt-get install texlive-xetex  
# file path and save location below are default; please change if they do not  
↪match yours  
[!] jupyter nbconvert --output-dir='/content/drive/MyDrive/' '/content/drive/  
↪MyDrive/CS444/assignment3_starter_sp24/assignment3_part2/MP3_P2.ipynb' --to_  
↪pdf
```